Intro

Practical Python for Data Science

This material is focused on writing complete, command-line programs that you can understand and use to get something done. Most of the programs included as examples and my implementations of the exercises are 100 lines or fewer. By leveraging standard modules and data types in Python, I will show you how to write readable, documented, and testable code. The chapters show many examples of complete programs you can run and alter, and the exercises expect you to write complete programs that satisfy a given test suite. After completing this material, you should be able to:

- Write, test, and document programs in bash and Python
- Use the source code management system Git to version, share, and distribute code
- Use parallelization techniques and hardware (HPC) to run programs faster
- Package and distribute software to create reproducible workflows

Using Git

I recommend you fork this repository and then add this repo as an upstream source:

```
$ git clone <your_fork_of_the_repo> ppds
$ cd ppds
$ REPO=https://github.com/kyclark/practical_python_for_data_science.git
$ git remote add upstream $REPO
```

That way you can `git pull upstream master` to get new content. I write the exercises in such a way that you will create new content that should not conflict with content I make.

Programming Environment

The material begins with the Unix command line. If you are working on Windows, I highly recommend you install Windows Subsystem for Linux and probably GitBash. If you are on an Apple computer, you have a full Unix system available through your Terminal app. The author uses a Mac with iTerm and vim editor to write, debug, and run programs. You may wish to use an editor like Sublime, TextWrangler, or Atom or an integrated develoment environment (IDE) like VSCode or PyCharm. However you choose to *write* code, this material assume you will *run* it from the command line. For many reasons, I have chosen not to

use Jupyter Notebooks. Some chapters may include a Notebook, but I would prefer to have students write command-line programs and use a testing framework like PyTest to ensure that code runs correctly, top to bottom.

Python

I personally prefer statically typed and "functional" languages like Rust, Elm, and Haskell, but I concede the dominance of dynamically typed languages such as Perl, Python, and Ruby. This material is intended to steer the student towards best practices when working in Python to avoid what I consider to be dangerous tendencies of the language. I will be using Python version 3.

Author

> "Computer programming has always been a self-taught, maverick occupation." - Ellen Ullman

> "Every great developer you know got there by solving problems they were unqualified to solve until they actually did it." - Patrick McKenzie

My name is Ken Youens-Clark (https://orcid.org/0000-0001-9961-144X), and I'm a Senior Scientific Programmer at the University of Arizona. While I have a Masters of Science in Biosystems Engineering, my undergraduate trianing was a BA in English Lit with a minor in music. As a kid, I had a RadioShack computer (probably a TRS-80, but I don't really remember) on which I wrote maybe two programs in BASIC. I never got into computing until after completing my bachelor's degree in 1995 when I started playing with computers and databases at my first job. The next year I landed a position where I learned Visual Basic on Windows 3.1, and that was when I actually got hooked on programming and problem solving. Since then, I've worked in several languages on various operating systems, and I've learned by doing – always having to tackle problems I was never trained to solve. I spent the longest part of my career using Perl in a bioinformatics settting, but Python has definitely taken over the data processing and machine learning space, so that's what I'll cover.

Organization

> "The only way to learn a new programming language is by writing programs in it." - Dennis Ritchie

Each chapter is written in Markdown format which you can read directly on GitHub where I host the material. I also use `pandoc` to convert this document to PDF via Latex, but sometimes things don't work perfectly smoothly. Still, there are PDFs you can download and read offline (or print, gasp!) if you like. I mostly like to write lots of programs (there are over 180 in this repo!) and comment on them. I embed the program directly into the Markdown, but they are all available in the `examples` directory for you to alter and run.

Perhaps owing to my musical training, I approach my teaching with the idea that you need to practice – quite a bit. You cannot learn Python or any language just by reading a text; therefore, most chapters have an `exercises` directory that contain code you should write and tests to help you know when they work sufficiently well. Exercises ending in `_a` are considered easier than those ending in `_b`. Some of the tests are minimal, e.g., I expect something like a "usage" statement under certain circumstances but don't verify anything more than the word "usage" at the beginning of your output. I trust you care enough about your user and your craft to actually create a proper help page. I include my own version of a solution that you can use to compare. I spent many years in the Perl community where "There Is More Than One Way To Do It" (TIMTOWTDI) is something of a mantra whereas the Python community espouses "There should be one – and preferably only one – obvious way to do it" (https://www.python.org/dev/peps/pep-0020/). I disagree with this notion and believe you can find many creative and beautiful solutions. More than anything, the solution that you figured out, that you understand, and that satisfies the test suite is the "right" one for you. Your style will change as you grow more knowledgeable and confident in your programming skills.

Acknowledgements

Copyright

Chapter 1

The Command Line

Whatever operating system you are on, I assume you have access to some Unix-like command line. That is, you are looking at a blank screen (probably) with a prompt like `$` waiting for you to type something. You could try typing "hello" and then the Enter key, and you'll probably see something like this:

```
$ hello
-bash: hello: command not found
```

NB: When you see a `$` given in the example prompts, it is a metacharacter indicating that this is the prompt for a normal (not super-user) account. You should type/copy/paste all the stuff *after* the `$`. If you ever see a prompt with "#" in a tutorial, it's indicating a command that should be run as the super-user/root account, e.g., installing some software into a system-wide directory so it can be shared by all users.

The first word on the command line needs to be a *command*. On my system, there's no program or built-in command called `hello` anywhere to be found (specifically in the `$PATH`, but we'll get to that in a bit), so it tells me `command not found`. Try `help` instead and you'll likely see quite a bit of output if you are on a system that is running GNU's `bash` shell. The command `hostname` should work pretty evreywhere, telling you the name of your machine. Certainly `ls` will work to show you a directory listing. Try it!

Common Unix Commands

> "The good news about computers is that they do what you tell them to do. The bad news is that they do what you tell them to do." - Ted Nelson

Let's look at some more commands you can do. This is by no means an exhaustive list, just a few to get you going. Try running each of them. To learn more about the tools, try both `man cmd` or `cmd -h` or `cmd --help`.

- **whoami**: reports your username
- **w**: shows who is currently on a system
- **man**: show the manual page for a command
- **echo**: say something
- **cowsay**: have a cow say something
- **env**: print your environment
- **printenv**: print some/all of your environment
- **which/type**: tells you the location of a program

- **touch**: create an empty regular file
- **file**: briefly describe a file or directory
- **pwd**: print working directory, where you are right now
- **ls**: list files in current directory
- **find**: find files or directories matching some conditions
- **locate**: find files using a database, requires daemon to run
- **cd**: change directory (with no arguments, `cd $HOME`)
- **cp**: copy a file (or "cp -r" to copy a directory)
- **mv**: move a file or diretory
- **mkdir**: create a directory
- **rmdir**: remove a directory
- **rm**: remove a file (or "rm -r" to remove a directory)
- **cat**: concatenate files (cf. http://porkmail.org/era/unix/award.html)
- **column**: arrange text into columns
- **paste**: merge lines of files
- **sort**: sort text or numbers
- **uniq**: remove duplicates *from a sorted list*
- **sed**: stream editor for altering text
- **awk/gawk**: pattern scanning and processing language
- **grep**: global regular expression program (maybe?), cf. https://en.wikipedia.org/wiki/Regular_expression
- **history**: look at past commands, cf. CTRL-R for searching your history directly from the command line
- **head**: view the first few (10) lines of a file
- **tail**: view the last (10) lines of a file
- **comm**: find lines in common/unique in two sorted files
- **top**: view the programs taking the most system resources (memory, I/O, time, CPUs, etc.), cf. "htop"
- **ps**: view the currently running processes
- **cut**: select columns from the output of a program
- **wc**: (character) word (and line) count
- **more/less**: pager programs that show you a "page" of text at a time; cf. https://unix.stackexchange.com/questions/81129/what-are-the-differences-between-most-more-and-less/81131
- **bc**: calculator
- **df**: report file system disk space usage; useful to find a place to land your data
- **du**: report disk usage; recommend "du -shc"; useful to identify large directories that need to be removed
- **ssh**: secure shell, like telnet only with encryption
- **scp**: secure copy a file from/to remote systems using ssh
- **rsync**: remote sync; uses scp but only copies data that is different
- **ftp**: use "file transfer protocol" to retrieve large data sets (better than HTTP/browsers, esp for getting data onto remote systems)
- **ncftp**: more modern FTP client that automatically handles anonymous logins

- **wget**: web get a file from an HTTP location, cf. "wget is not a crime" and Aaron Schwartz
- **|**: pipe the output of a command into another command
- **>, >>**: redirect the output of a command into a file; the file will be created if it does not exist; the single arrow indicates that you wish to overwrite any existing file, while the double-arrows say to append to an existing file
- **<**: redirect contents of a file into a command
- **nano**: a very simple text editor; until you're ready to commit to vim or emacs, start here
- **md5sum**: calculate the MD5 checksum of a file
- **diff**: find the differences between two files
- **xargs**: take a list from one command, concatenate and pass as the arguments to another command

The Unix filesystem hierarchy

The Unix filesystem can thought of as a graph or a tree. The root is "/" and is called the "root directory." We can "list" the contents of a directory with `ls`. Without any arguments, this prints the contents of the current working directory which you can print with `pwd` (print working directory). You can `ls /` to see the contents of the root directory, or `ls $HOME` to see your own home directory.

Moving around the filesystem

You can print your current working directory either with `pwd` or `echo $PWD`.

The `cd` command is used to "change directory," e.g., `cd /rsgrps/bh_class/`. If you wish to return to the previous working directory, use `cd -`.

If you provide no argument to `cd`, it will change to your `$HOME` directory which is also known in bash by the `~` (tilde or twiddle). So these three commands are equivalent:

- `cd`
- `cd ~`
- `cd $HOME`

Once you are in a directory, use `ls` to inspect the contents. If you do not provide an argument to `ls`, it assumes the current directory which has the alias `.` The parent directory is `..`.

You can use both absolute and relative paths with `cd`. An absolute path starts from the root directory, e.g., "/usr/local/bin/". A relative path does not start with the leading `/` and assumes a path relative to your current working

directory. If you were in the "/usr/local" directory and wanted to change to "/usr/local/bin", you could either `cd /usr/local/bin` (absolute) or `cd bin` (relative to "/usr/local").

Once you are in "/usr/local/bin", what would `pwd` show after you did `cd ../..`?

Chaining commands

> "Programming is breaking of one big impossible task into several very small possible tasks." - Jazzwant

Most Unix commands use STDIN ("standard in"), STDOUT ("standard out"), and STDERR ("standard error"). For instance, the `env` program will show you key/value pairs that describe. your environment – things like your user name (`$USER`), your shell (`$SHELL`), your current working directory (`$PWD`). It can be quite a long list, so you could send the output (STDOUT) of `env` to `head` to see just the first few lines:

```
$ env | head
TERM_PROGRAM=iTerm.app
TERM=xterm-256color
SHELL=/bin/bash
TMPDIR=/var/folders/0h/vjzkyy052qx4p70trn2p2h400000gn/T/
PERL5LIB=/Users/kyclark/work/imicrobe/lib
Apple_PubSub_Socket_Render=/private/tmp/com.apple.launchd.MoIOCra0uS/Render
TERM_PROGRAM_VERSION=3.2.6
OLDPWD=/Users/kyclark/work/biosys-analytics/lectures
TERM_SESSION_ID=w2t0p0:34E89330-9CEE-4FA6-9039-3674CBDE4655
USER=kyclark
```

If you provide a file name as an argument to `head`, it will work on that:

```
$ head /usr/share/dict/words
A
a
aa
aal
aalii
aam
Aani
aardvark
aardwolf
Aaron
```

Without any arguments, `head` assumes you must want it to read from STDIN. Many other programs will assume STDIN if not provided an argument. For

instance, you could pipe `env` into `grep` to look for lines with the word "TERM" in them:

```
$ env | grep TERM
TERM_PROGRAM=iTerm.app
TERM=xterm-256color
TERM_PROGRAM_VERSION=3.2.6
TERM_SESSION_ID=w2t0p0:34E89330-9CEE-4FA6-9039-3674CBDE4655
ITERM_PROFILE=Default
ITERM_SESSION_ID=w2t0p0:34E89330-9CEE-4FA6-9039-3674CBDE4655
COLORTERM=truecolor
```

If you are fortunate enough to have the `fortune` and `cowsay` programs on your system, you can do this:

```
$ fortune | cowsay
 _____
/ Somebody ought to cross ball point pens \
| with coat hangers so that the pens will |
\ multiply instead of disappear.          /
 -----------------------------------------
        \   ^__^
         \  (oo)_____
            (__)\       )\/\
                ||----w |
                ||     ||
```

Manual pages

> "Programming isn't about what you know; it's about what you can figure out." - Chris Pine

The `man` program will show you the manual page for a program, if it exists. Just type `man <program>`, e.g., `man wget`. Inside a manpage, you can use the `/` to search for a string. Use `q` to "quit" `man`. Most programs will also show you a help/usage document if you run them with `-h`, `--help`, or `-help`. I often find it useful to `grep` the help, e.g.:

```
$ wget --help | grep clobber
  -nc, --no-clobber                skip downloads that would download to
       --unlink                    remove file before clobber
```

Pronunciation

- `/`: "slash"; the thing leaning the other way is a "backslash"

- **sh**: "shuh" or "ess-ach"
- **etc**: "et-see"
- **usr**: "user"
- **src**: "source"
- **#**: "hash" (NOT "hashtag") or "pound"
- **$**: "dollar"
- **!**: "bang"
- **#!**: "shebang"
- **ˆ**: "caret"
- **PID**: "pid" (not pee-eye-dee)
- **~**: "twiddle" or "tilde"; shortcut to your home directory when alone, shortcut to another user's home directory when used like "~bhurwitz"

Variables

You will see things like `$USER` and `$HOME` that start with the `$` sign. These are variables because they can change from person to person, system to system. On most systems, my username is "kyclark" but I might be "kclark" or "kyclark1" on others, but on all systems `$USER` refers to whatever value is defined for my username. Similarly, my `$HOME` directory might be "/Users/kyclark," "/home1/03137/kyclark," or "/home/u20/kyclark," but I can always refer to the idea of my home directory with the variable `$HOME`.

When you are assigning a variable, you do not use the `$`.

```
[hpc:login3@~]$ SECRET=ilikecake
[hpc:login3@~]$ echo $SECRET
ilikecake
[hpc:login3@~]$ echo SECRET
SECRET
```

To remove a variable from your environment, use `unset`:

```
[hpc:login3@~]$ unset SECRET
[hpc:login3@~]$ echo $SECRET
```

Notice that there is no error when referencing a variable that does not exist or has not been set.

Control sequences

If you launch a program that won't stop, you can use CTRL-C (where "CTRL" is the "control" key sometime written "ˆC" or "ˆ-C") to send an "interrupt" signal to the program. If it is well-behaved, it should stop, but it may not. For example, perhaps I've tried to use a text editor to open a 10G FASTA file and

now my terminal is unresponsive because the editor is using all available memory. I could open another terminal on the machine and run `ps -fu $USER` to find all the programs I am running:

```
$ ps -fu $USER
UID          PID   PPID  C STIME TTY          TIME CMD
kyclark    31718  31692  0 12:16 ?        00:00:00 sshd: kyclark@pts/75
kyclark    31723  31718  0 12:16 pts/75   00:00:00 -bash
kyclark    33265  33247  0 12:16 ?        00:00:00 sshd: kyclark@pts/86
kyclark    33277  33265  1 12:16 pts/86   00:00:00 -bash
kyclark    33792  33277  9 12:17 pts/86   00:00:00 vim maize_genome.fasta
kyclark    33806  31723  0 12:17 pts/75   00:00:00 ps -fu kyclark
```

The PID is the "process ID" and the PPID is the "parent process ID." In the above table, let's assume I want to kill `vim`, so I type `kill 33792`. If in a reasonable amount of time (a minute or so) that doesn't work, I could use `kill -9` (but it's considered a bit uncouth).

CTRL-Z is used to put a process into the background. This could be handy if, say, you were in an editor, you make a change to your code, you CTRL-Z to background the editor, you run the script to see if it worked, then you `fg` to bring it back to the foreground or `bg` it to have it resume running in the background. I would consider this a sub-optimal work environment, but it's fine if you were for some reason limited to a single terminal window.

Generally if you want a program to run in the background, you would launch it from the command line with an ampersand ("&") at the end:

```
$ my-background-prog.sh &
```

Lastly, most Unix programs interpret CTRL-D as the end-of-input signal. You can use this to send the "exit" command to most any interactive program, even your shell. Here's a way to enter some text into a file directly from the command line without using a text editor. After typing the last line (i.e., type "chickens.<Enter>"), type CTRL-D:

```
$ cat > wheelbarrow
so much depends
upon

a red wheel
barrow

glazed with rain
water

beside the white
chickens.
<CTRL-D>
```

```
$ cat wheelbarrow
so much depends
upon

a red wheel
barrow

glazed with rain
water

beside the white
chickens.
```

Handy command line shortcuts

- Tab: hit the Tab key for command completion; hit it early and often!
- **!!**: (bang-bang) execute the last command again
- **!$**: (bang-dollar) the last argument from your previous command line (think of the $ as the right anchor in a regex)
- **!^**: (band-caret) the first argument from your previous command line (think of the ^ as the left anchor in a regex)
- CTRL-R: reverse search of your history
- Up/down cursor keys: go backwards/forwards in your history
- CTRL-A, CTRL-E: jump to the start, end of the command line when in emacs mode (default)

NB: If you are on a Mac, it's easy to remap your (useless) CAPSLOCK key to CTRL. Much less strain on your hand as you will find you need CTRL quite a bit, even more so if you choose emacs for your $EDITOR.

Altering your $PATH

> I feel like there should seriously be some sort of first-year level class where people learn how to add something to their path and what that means. – Kristopher Micinski

Your `$PATH` setting is an ordered, colon-delimited list of directories that will be searched to find programs. Run `echo $PATH` to see yours. It probably looks something like this:

```
$ echo $PATH
/usr/local/sbin:/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin
```

Here's my `$PATH` on my laptop:

```
$ echo $PATH | gsed "s/:/\n/g"
/Users/kyclark/.cargo/bin
/Users/kyclark/bin
/anaconda3/bin
/Users/kyclark/.local/bin
/Users/kyclark/work/cyverse-cli/bin
/usr/local/sbin
/usr/local/bin
/usr/bin
/bin
/usr/sbin
/sbin
/Library/TeX/texbin
```

And here is my `$PATH` on the HPC at the University of Ariziona:

```
[hpc:login3@~]$ echo $PATH | sed "s/:/\n/g"
/rsgrps/bh_class/bin
/home/u20/kyclark/.cargo/bin
/home/u20/kyclark/.local/bin
/cm/local/apps/gcc/6.1.0/bin
/cm/shared/uaapps/pbspro/18.2.1/sbin
/cm/shared/uaapps/pbspro/18.2.1/bin
/opt/TurboVNC/bin
/cm/shared/uabin
/usr/lib64/qt-3.3/bin
/cm/local/apps/environment-modules/4.0.0//bin
/usr/local/bin
/bin
/usr/bin
/usr/local/sbin
/usr/sbin
/sbin
/sbin
/usr/sbin
/cm/local/apps/environment-modules/4.0.0/bin
```

I've used "sed" (or "gsed" which is GNU sed) to add a newline after each colon so you can more easily see that the directories are separated by colons. If you use the `which` command to see where a program lives, you can see that it is found in a directory that is included in your `$PATH`:

```
$ which sed
/usr/bin/sed
$ which python
/anaconda3/bin/python
```

By definition, if a program's location is not in your `$PATH`, then `which` cannot

find it, and that means you cannot run it without giving a full path to the program, e.g., `/usr/sbin/foobar/baz`.

In this repository I have included a `bin` directory that has some useful Python programs like `new_py.py` which we will use later to stub out new Python programs. In order to use them, you have three options:

1. Use the complete path to the programs. E.g., if you have installed this into `$HOME/ppds`, then execute `$HOME/ppds/bin/new_py.py`
2. Copy the contents of the `bin` directory to one of the other directories that are already in your `$PATH`, e.g., `cp $HOME/ppds/bin/* /usr/local/bin`, but that might require root privilege that you don't have.
3. Add `$HOME/ppds/bin` to your `$PATH`

```
export PATH="$HOME/ppds/bin:$PATH"
```

You just told your shell (bash) to set the `$PATH` variable to `$HOME/ppds/bin` plus whatever is was before. Since we want this to happen each time we log in, so we can add this command to `$HOME/.bashrc`:

```
echo "export PATH=$HOME/ppds/bin:$PATH" >> ~/.bashrc
```

As you find or create useful programs that you would like to have available globally on your system (i.e., not just in the current working directory), you can create a location like `$HOME/bin` (or my preferred `$HOME/.local/bin`) and add this to your `$PATH` as well. You can add as many directories as you like (within reason).

Dotfiles

"Dotfiles" are files with names that begin with a dot. They are normally hidden from view unless you use `ls -a` to list "all" files. A single dot `.` means the current directory, and two dots `..` mean the parent directory. Your ".bashrc" (or maybe ".profile" or maybe ".bash_profile" depending on your system) file is read every time you login to your system, so you can remember your customizations. "Rc" may mean "resource configuration," but who really knows?

After a while, you may wish to collect your dotfiles into a Github repo, e.g., https://github.com/kyclark/dotfiles.

Aliases

Sometimes you'll find you're using a particular command quite often and want to create a shortcut. You can assign any command to a single "alias" like so:

```
alias cx='chmod +x'
```

```
alias up2='cd ../../'
alias up3='cd ../../../'
```

If you execute this on the command line, the alias will be saved until you log out. Adding these lines to your `.bashrc` will make it available every time you log in. When you make a change and want the shell to bring those into the current environment, you need to `source` the file. The command `.` is an alias for `source`:

```
$ source ~/.bashrc
$ . ~/.bashrc
```

Permissions

When you execute `ls -l`, you'll see the "long" listing of the contents of a directory similar to this:

```
-rwxr-xr-x    1 kyclark   staff      174 Aug  9 20:21 abs.py*
drwxr-xr-x  14 kyclark   staff      476 Aug  3 12:14 anaconda3/
```

The first column of data contains a wealth of information represent in 10 bits. The first bit is:

- "-" for a regular file
- "d" for a directory
- "l" for a symlink (like a shortcut)

The other nine bits are broken into sets of three bits that represent the permissions for the "user," "group," and "other." The "abs.py" is a regular file we can tell from the first dash. The next three bits show "rwx" which means that the user ("kyclark") has read, write, and execute permissions for this file. The next three bits show "r-x" meaning that the group ("staff") can read and execute the file only. The same is true for all others.

When you create a file, the normal default is that it is not executable. You must specifically tell Unix to change the mode of the file using the "chmod" command. Often it's enough to say:

```
$ chmod +x myprog.sh
```

To turn on the execute bits for everyone, but it's possible to have much finer control of the permissions. If you only want user and group to have execute, then do:

```
$ chmod ug+x myprog.sh
```

Removing is done with a "-," so any combination of `[ugo][+-][rwx]` will usually get you what you want.

Sometimes you may see instructions to `chmod 775` a file. This is using octal notation where the three bits "rwx" correspond to the digits "421," so the first "7" is "4+2+1" which equals "rwx" whereas the "5" = "4+1" so only "rw":

```
 user    group    other
r w x    r w x    r w x
4 2 1 | 4 2 1 | 4 2 1
+ + +    + + +    + - +
 = 7      = 7      = 5
```

Therefore "chmod 775" is the same as:

```
$ chmod -rwx myfile
$ chmod ug+rwx myfile
$ chmod o+rw myfile
```

When you create ssh keys or config files, you are instructed to `chmod 600`:

```
 user    group    other
r w x    r w x    r w x
4 2 1 | 4 2 1 | 4 2 1
+ + -    - - -    - - -
 = 6      = 0      = 0
```

Which means that only you can read or write the file, and no one else can do anything with it. So you can see that it can be much faster to use the octal notation.

When you are trying to share data with your colleagues who are on the same system, you may put something into a shared location but they complain that they cannot read it or nothing is there. The problem is most likely permissions. The "uask" setting on a system determines the default permissions, and it may be that the directory and/or files are readable only by you. It may also be that you are not in a common group that you can use to grant permission, in which case you can either:

- politely ask your sysadmin to create a new group OR
- `chmod 777` the directory, which is probably the worst option as it makes the directory completely accessible to anyone to do anything. In short, don't do this unless you really don't care if someone accidentally or maliciously wipes out your data.

File system layout

The top level of a Unix file system is "/" which is called "root." Confusingly, there is also an account named "root" which is basically the super-user/sysadmin (systems administrator). Unix has always been a multi-tenant system . . .

Installing software

Much of the time, "bioinformatics" seems like little more than installing software and chaining them together with scripts. Sometimes you may be lucky enough to have a "sysadmin" (systems administrator) who can assist you, but most of the time you'll find yourself needing to take care of business yourself.

My suggestions for installing software are (in order):

Sysadmin

Go introduce yourself to your sysadmins. Take them to lunch or order them some pizza or drop off some good beer or whiskey. Whatever it takes to be on good terms because a good sysadmin who is responsive to your needs is an enormous help. If they are willing to install software for you, that is the way to go. Often, though, this is a task far beneath them, and they would expect you to be able to fend for yourself. They may provide `sudo` (https://xkcd.com/149/) privileges to allow you to install software into shared locations (e.g., `/usr/local`), but it's more likely they would expect you to install into your `$HOME`.

Package managers

There are several package management systems for Linux and OSX including apt-get, yum, homebrew, macports, and more. These usually relieve the problems of software compatibility and shared libraries. Unless you have `sudo` to install globally, you can configure to install into your `$HOME`.

Binary installations

Quite often you'll be happy to find that the maintainers of the software you need have gone to the trouble to build binary distributions for your system, which is likely to be a generic 64-bit Linux platform. Often you can just download the binaries and put them into your `$PATH`. There is usually a "README" or "INSTALL" file that will explain exactly what to do. To use the binaries, you can:

1) Always refer to the full path to the binary
2) Place them into a directory in your `$PATH` like `$HOME/.local/bin`
3) Add the new directory to your `$PATH`

Source installations

Installing from source usually means downloading a "tarball" ("tar" = "tape archive," a container of files, that is then compressed with a program like "gzip" to create a ".tar.gz" or ".tgz" file extension), running `./configure` to figure out how it can build on your system, and then `make` to build the binaries. Usually you will run `make install` to put the binaries into their proper directory, but sometimes you just `make` and copy the files yourself.

The basic steps for installing into your `$HOME` are usually:

```
$ tar xvf package.tgz
$ ./configure --prefix=$HOME/.local
$ make && make install
```

When I'm in an environment with a directory I can share with my team (like the UA HPC), I'll configure the package to install into that shared space so that others can use the program. When I'm on a system like "stampede" where I cannot share with others, I'll usually install into my `$HOME/.local` or some sort of "work" directory.

The Data and Software Carpentries

The Software Carpentry project aims to teach basic command-line usage. You should definitely work through `https://swcarpentry.github.io/shell-novice/`.

Chapter 2

Unix exercises

Here are some tasks that will introduce how the commands from the previous chapter can be combined to get things done. Sometimes you can get exactly what you need with command-line tools and never need to write a program. Here I show you output on my system, and I encourage you to type (do not copy and paste!) the commands on your system to compare. Note that we are likely to be using different versions of a Unix-like OS, so the implementations of commands like `ls` or `sed` might differ by output or arguments. Be sure to consult your manpage and versions to understand any differences you see.

Find the number of unique users on a shared system

We know that `w` will tell us the users logged in. Try it now on a system that has many users (i.e., not your laptop) and see the output. Likely there are dozens of users, so we'll connect the output of `w` to `head` using a pipe `|` so that we only see the first five lines:

```
[hpc:login2@~]$ w | head -5
 09:39:27 up 65 days, 20:05, 10 users,  load average: 0.72, 0.75, 0.78
USER     TTY      FROM             LOGIN@   IDLE   JCPU   PCPU WHAT
kyclark  pts/2    gatekeeper.hpc.a 09:38    0.00s  0.05s  0.02s w
emsenhub pts/0    gatekeeper.hpc.a 04:05   14.00s  0.87s  0.87s -bash
joneska  pts/3    gatekeeper.hpc.a 08:25    1:12m  0.16s  0.12s vim results_x2r
```

Really we want to see the first five *users*, not the first five *lines* of output. To skip the first two lines of headers from `w`, we can pipe `w` into `awk` and tell it we only want to see output when the Number of Records (`NR`) is greater than 2:

```
[hpc:login2@~]$ w | awk 'NR>2' | head -5
kyclark  pts/2    gatekeeper.hpc.a 09:38    0.00s  0.07s  0.03s w
emsenhub pts/0    gatekeeper.hpc.a 04:05   26.00s  0.87s  0.87s -bash
joneska  pts/3    gatekeeper.hpc.a 08:25    1:13m  0.16s  0.12s vim results_x2r
shawtaro pts/4    gatekeeper.hpc.a 08:06   58:34   0.17s  0.17s -bash
darrenc  pts/5    gatekeeper.hpc.a 07:58   51:07   0.14s  0.07s qsub -I -N pipe
```

`awk` takes a PREDICATE and a CODE BLOCK (contained within curly brackets `{}`). Without a PREDICATE, `awk` prints the whole line. I only want to see the first column, so I can tell `awk` to `print` just column `$1`:

```
[hpc:login2@~]$ w | awk 'NR>2 {print $1}' | head -5
kyclark
```

```
emsenhub
joneska
shawtaro
darrenc
```

We can see that the some users like "joneska" are logged in multiple times:

```
[hpc:login2@~]$ w | awk 'NR>2 {print $1}'
kyclark
emsenhub
joneska
shawtaro
darrenc
guven
guven
guven
joneska
dmarrone
```

Let's `uniq` that output:

```
[hpc:login2@~]$ w | awk 'NR>2 {print $1}' | uniq
kyclark
emsenhub
joneska
shawtaro
darrenc
guven
joneska
dmarrone
```

Hmm, that's not right – "joneska" is listed twice, and that is not unique. Remember that `uniq` only works *on sorted input*? So let's sort those names first:

```
[hpc:login2@~]$ w | awk 'NR>2 {print $1}' | sort | uniq
darrenc
dmarrone
emsenhub
guven
joneska
kyclark
shawtaro
```

To count how many unique users are logged in, we can use the `wc` (word count) program with the `-l` (lines) flag to count just the lines from the previous command

```
[hpc:login2@~]$ w | awk 'NR>2 {print $1}' | sort | uniq | wc -l
7
```

So what you see is that we're connecting small, well-defined programs together using pipes to connect the "standard input" (STDIN) and "standard output" (STDOUT) streams. There's a third basic file handle in Unix called"standard error" (STDERR) that we'll come across later. It's a way for programs to report problems without simply dying. You can redirect errors into a file like so:

```
$ program 2>err
$ program 1>out 2>err
```

The first example puts STDERR into a file called "err" and lets STDOUT print to the terminal. The second example captures STDOUT into a file called "out" while STDERR goes to "err."

NB: Sometimes a program will complain about things that you cannot fix, e.g., `find` may complain about file permissions that you don't care about. In those cases, you can redirect STDERR to a special filehandle called `/dev/null` where they are forgotten forever – kind of like the "memory hole" in 1984.

```
$ find / -name my-file.txt 2>/dev/null
```

Count "oo" words

On almost every Unix system, you can find `/usr/share/dict/words`. Let's use `grep` to find how many have the "oo" vowel combination. It's a long list, so I'll pipe it into "head" to see just the first five:

```
$ grep 'oo' /usr/share/dict/words | head -5
abloom
aboon
aboveproof
abrood
abrook
```

Yes, that works, so redirect those words into a file and count them. Notice the use of `!$` (bang-dollar) to reference the last argument of the previous line so that I don't have to type it again (really useful if it's a long path):

```
$ grep 'oo' /usr/share/dict/words > oo-words
$ wc -l !$
10460 oo-words
```

Let's count them directly out of `grep`:

```
$ grep 'oo' /usr/share/dict/words | wc -l
10460
```

Do any of those words additionally contain the "ow" sequence?

```
$ grep 'oo' /usr/share/dict/words | grep 'ow' | head -5
```

```
arrowroot
arrowwood
balloonflower
bloodflower
blowproof
```

How many are there?

```
$ grep 'oo' /usr/share/dict/words | grep 'ow' | wc -l
158
```

How many *do not* contain the "ow" sequence? Use **grep -v** to invert the match:

```
$ grep 'oo' /usr/share/dict/words | grep -v 'ow' | wc -l
10302
```

Do those numbers add up?

```
$ bc <<< 158+10302
10460
```

Find unclustered protein sequences

A labmate wants help finding the sequences of proteins that failed to cluster.
Here is the setup:

```
$ wget ftp://ftp.imicrobe.us/biosys-analytics/exercises/unclustered-proteins.tgz
$ tar xvf unclustered-proteins.tgz
$ cd unclustered-proteins
```

The "README" contains our instructions:

```
The file "cdhit60.3+.clstr" contains all of the GI numbers for
proteins that were clustered and put into hmm profiles.  The file
"proteins.fa" contains all proteins (the header is only the GI
number).  Extract the proteins from the "proteins.fa" file that were
not clustered.
```

If we look at the IDs in the proteins file, we'll see they are integers:

```
$ grep '>' proteins.fa | head -5
>388548806
>388548807
>388548808
>388548809
>388548810
```

Where can we find those protein IDs in the "cdhit60.3+.clstr" file?

```
$ head -5 cdhit60.3+.clstr
```

```
>Cluster_5086
0    358aa, >gi|317183610|gb|ADV... at 66.76%
1    361aa, >gi|315661179|gb|ADU... at 70.36%
2    118aa, >gi|375968555|gb|AFB... at 70.34%
3    208aa, >gi|194307477|gb|ACF... at 61.54%
```

The format of the file is similar to a FASTA file where the ">" sign at the
left-most column identifies a cluster with the following lines showing the IDs of
the sequences in the cluster. To extract just the clustered IDs, we cannot just
do `grep '>'` as we'll get both the cluster IDs and the protein IDs.

```
$ grep '>' cdhit60.3+.clstr | head -5
>Cluster_5086
0    358aa, >gi|317183610|gb|ADV... at 66.76%
1    361aa, >gi|315661179|gb|ADU... at 70.36%
2    118aa, >gi|375968555|gb|AFB... at 70.34%
3    208aa, >gi|194307477|gb|ACF... at 61.54%
```

We'll need to use a regular expression (the `-e` for "extended" on most greps, but
sometimes not required) to say that we are looking at the beginning of a line ^
for a >:

```
$ grep -e '^>' cdhit60.3+.clstr | head -5
>Cluster_5086
>Cluster_10030
>Cluster_8374
>Cluster_13356
>Cluster_7732
```

and then invert that with "-v":

```
$ grep -v '^>' cdhit60.3+.clstr | head -5
0    358aa, >gi|317183610|gb|ADV... at 66.76%
1    361aa, >gi|315661179|gb|ADU... at 70.36%
2    118aa, >gi|375968555|gb|AFB... at 70.34%
3    208aa, >gi|194307477|gb|ACF... at 61.54%
4    358aa, >gi|291292536|gb|ADD... at 68.99%
```

The integer protein IDs we want are in the third column of this output when split
on whitespace. The tool `awk` is perfect for this, and whitespace is the default
split character (as opposed to `cut` which uses tabs):

```
$ grep -ve '^>' cdhit60.3+.clstr | awk '{print $3}' | head -5
>gi|317183610|gb|ADV...
>gi|315661179|gb|ADU...
>gi|375968555|gb|AFB...
>gi|194307477|gb|ACF...
>gi|291292536|gb|ADD...
```

The protein ID is still nestled there in the second field when splitting on the

vertical bar (pipe). Again, `awk` is perfect, but we need to tell it to split on something other than the default by using the "-F" flag:

```
$ grep -ve '^>' cdhit60.3+.clstr | awk '{print $3}' | \
  awk -F'|' '{print $2}' | head -5
317183610
315661179
375968555
194307477
291292536
```

These are the protein IDs for those that were successfully clustered, so we need to capture these to a file which we can do with a redirect `>`. Since each protein might have been clustered more than once, so I should `sort | uniq` the list:

```
$ grep -ve '^>' cdhit60.3+.clstr | awk '{print $3}' | \
  awk -F"|" '{print $2}' | sort | uniq > clustered-ids.o
```

The "proteins.fa" is actually a little problematic. Some of the IDs have extra information. If you `grep '^>' proteins.fa`, you will see 220K IDs scroll by, not all of which are just integers. Let's isolate those that do not look like integers.

First we can remove the leading ">" from the FASTA header lines with this:

```
$ grep '^>' proteins.fa | sed "s/^>//"
```

If I can find a regular expression that matches what I want, then I can use `grep -v` to invert it to find the complement. `^\d+$` will do the trick. Let's break down that regex:

```
^ \d + $
1 2  3 4
```

1. start of the line
2. a digit (0-9)
3. one or more
4. end of the line

This particular regex uses extensions introduced by the Perl programming language, so we need to use the `-P` flag. Add the `-v` to invert it:

```
$ grep -e '^>' proteins.fa | sed "s/^>//" | grep -v -P '^\d+$' | head -5
26788002|emb|CAD19173.1| putative RNA helicase, partial [Agaricus bisporus virus X]
26788000|emb|CAD19172.1| putative RNA helicase, partial [Agaricus bisporus virus X]
985757046|ref|YP_009222010.1| hypothetical protein [Alternaria brassicicola fusarivirus 1]
985757045|ref|YP_009222011.1| hypothetical protein [Alternaria brassicicola fusarivirus 1]
985757044|ref|YP_009222009.1| polyprotein [Alternaria brassicicola fusarivirus 1]
```

Looking at the above output, we can see that it would be pretty easy to get rid of everything starting with the vertical bar, and `sed` is perfect for this. Note that we can tell `sed` to do more than one action by separating them with semicolons. Lastly, we need to ensure the IDs are sorted for the next step:

```
$ grep -e '^>' proteins.fa | sed "s/^>//; s/|.*//" | sort > protein-ids.o
```

To find the lines in "protein-ids.o" that are not in "clustered-ids.o", I can use
the `comm` (common) command:

```
$ comm -23 protein-ids.o clustered-ids.o > unclustered-ids.o
```

Did we get a reasonable answer?

```
$ wc -l clustered-ids.o unclustered-ids.o
  16257 clustered-ids.o
 204263 unclustered-ids.o
 220520 total
$ wc -l protein-ids.o
220520 protein-ids.o
```

Gapminder

For this exercise, look in the **biosys-analytics/data/gapminder** directory.

How many "txt" files are in the directory?

```
$ ls *.txt | wc -l
```

How many lines are in each/all of the files?

```
$ wc -l *.txt
```

You can use `cat` to spew at the entire contents of a file into your shell, but if
you'd just like to see the top of a file, you can use:

```
$ head Trinidad_and_Tobago.cc.txt
```

If you only want to see 5 lines, use `-n 5` or `-5` .

For our exercise, we'd like to combine all the files into one file we can analyze.
That's easy enough with:

```
$ cat *.cc.txt > all.txt
```

Let's use `head` to look at the top of file:

```
$ head -5 all.txt
Afghanistan    1997    22227415    Asia    41.763    635.341351
Afghanistan    2002    25268405    Asia    42.129    726.7340548
Afghanistan    2007    31889923    Asia    43.828    974.5803384
Afghanistan    1952    8425333     Asia    28.801    779.4453145
Afghanistan    1957    9240934     Asia    30.332    820.8530296
```

Hmm, there are no column headers. Let's fix that. There's one file that's pretty
different in content (it has only one line) and name ("country.cc.txt"):
```

```
$ cat country.cc.txt
country    year    pop    continent    lifeExp    gdpPercap
```

Those are the headers that you can combine to all the other files to get named columns, something very important if you want to look at the data in Excel and R/Python data frames.

```
$ rm all.txt
$ mv country.cc.txt headers
$ cat headers *.txt > all.txt
$ head -5 all.txt | column -t
country      year  pop        continent  lifeExp  gdpPercap
Afghanistan  1997  22227415   Asia       41.763   635.341351
Afghanistan  2002  25268405   Asia       42.129   726.7340548
Afghanistan  2007  31889923   Asia       43.828   974.5803384
Afghanistan  1952  8425333    Asia       28.801   779.4453145
```

Yes, that looks much better. Double-check that the number of lines in the `all.txt` match the number of lines of input:

```
$ wc -l *.cc.txt headers
$ wc -l all.txt
```

How many observations do we have for 1952? For this, we need to find all the rows where the second field is equal to "1952," and `awk` will let us do just that. Normally `awk` splits on whitespace, but we have tab-delimited so we need to use `-F"\t"`. Recipes in `awk` take the form of a CONDITIONAL and an ACTION. If the CONDITIONAL is missing, then the ACTION is applied to all lines. If the ACTION is missing, then the default is to print the entire line. Here we just provide the CONDITIONAL and then count the results:

```
$ awk -F"\t" '$2 == "1952"' all.txt | wc -l
```

How many observations for each year?

```
$ awk -F"\t" '{print $2}' all.txt | sort | uniq -c
```

How many observations are present for Africa (the fourth field is continent)?

```
$ awk -F"\t" '$4 == "Africa"' all.txt | wc -l
```

And what are the countries in Africa?

```
$ awk -F"\t" '$4 == "Africa" {print $1}' all.txt | sort | uniq
```

How many observations for for each continent?

```
$ awk -F"\t" 'NR>1 {print $4}' all.txt | sort | uniq -c
```

What was the world population in 1952? To answer this, we need to get the third column when the second column is "1952":

```
$ awk -F"\t" '$2 == "1952" {print $3}' all.txt
```

There's a problem because one of the numbers is in scientific notation:

```
$ awk -F"\t" '$2 == "1952" {print $3}' all.txt | grep [a-z]
3.72e+08
```

Let's just remove that using `grep -v` (the `-v` reverses the match), then use the `paste` command to put a "+" in between all the numbers:

```
$ awk -F"\t" '$2 == "1952" {print $3}' *.cc.txt | grep -v [a-z] | paste -sd+ -
```

and then we pipe that to the `bc` calculator:

```
$ awk -F"\t" '$2 == "1952" {print $3}' *.cc.txt | grep -v [a-z] | paste -sd+ - | bc
2034957150.999989
```

It bothers me that it's not an integer, so I'm going to use `printf` in the `awk` command to trim that:

```
$ awk -F"\t" '$2 == "1952" {printf "%d\n", $3}' *.cc.txt | grep -v [a-z] | paste -sd+ - | bc
2406957150
```

I know that's all a bit crude and absurd, but I thought you might be curious just how far you can take this.

How many observations where the life expectancy ("lifeExp," field #5) is greater than 40?

```
$ awk -F"\t" '$5 > 40' all.txt | wc -l
```

How many of those are from Africa?

```
$ awk -F"\t" '$5 > 40 && $4 == "Africa"' all.txt
```

How many countries had a life expectancy greater than 70, grouped by year?

```
$ awk -F"\t" '$5 > 70 { print $2 }' all.txt | sort | uniq -c
     5 1952
     9 1957
    16 1962
    25 1967
    30 1972
    38 1977
    44 1982
    49 1987
    54 1992
    65 1997
    75 2002
    83 2007
```

How could we add continent to this?

```
$ awk -F"\t" '$5 > 70 { print $2 ":" $4 }' all.txt | sort | uniq -c
```

As you look at the data and want to ask more complicated questions like how does gdpPercap affect lifeExp, you'll find you need more advanced tools like Python or R. Now that the data has been collated and the columns named, that will be much easier.

What if we want to add headers to each of the files?

```
$ mkdir wheaders
$ for file in *.txt; do cat headers $file > wheaders/$file; done
$ wc -l wheaders/* | head -5
      13 wheaders/Afghanistan.cc.txt
      13 wheaders/Albania.cc.txt
      13 wheaders/Algeria.cc.txt
      13 wheaders/Angola.cc.txt
      13 wheaders/Argentina.cc.txt
$ head wheaders/Vietnam.cc.txt
country     year    pop     continent    lifeExp    gdpPercap
Vietnam     1952    26246839    Asia     40.412     605.0664917
Vietnam     1957    28998543    Asia     42.887     676.2854478
Vietnam     1962    33796140    Asia     45.363     772.0491602
Vietnam     1967    39463910    Asia     47.838     637.1232887
Vietnam     1972    44655014    Asia     50.254     699.5016441
Vietnam     1977    50533506    Asia     55.764     713.5371196
Vietnam     1982    56142181    Asia     58.816     707.2357863
Vietnam     1987    62826491    Asia     62.82      820.7994449
Vietnam     1992    69940728    Asia     67.662     989.0231487
```

Exercises

Exercise: head.sh

Write a bash script called `head.sh` that mimics the `head` utility where it will print the first few lines of a file. The script should expect one required argument (the file) and a second optional argument of the number of lines, defaulting to 3. If are no arguments, it should print a "Usage" and exit *with an error code* Your program will expect to receive an argument in `$1` and maybe a second in `$2`. If the first argument is not a file, it should notify the user and exit *with an error code.* If the second argument is missing, use the value "3". Print the number of lines requested by the user by iterating over the lines in the file and exiting the loop appropriately. Do not use the actual `head` command!

```
$ ./head.sh
Usage: head.sh FILE NUM
$ ./head.sh files/issa.txt
Selected Haiku by Issa
```

```
Don't worry, spiders,
$ ./head.sh files/issa.txt 5
Selected Haiku by Issa

Don't worry, spiders,
I keep house
casually.
```

Exercise: cat_n.sh

Write a bash program called `cat_n.sh` that mimics the behavior of `cat -n`
where it will print the line number and line of an input file. If there are no
arguments, it should print a "Usage" and exit *with an error code*. Your program
will expect to receive an argument in `$1`. If the argument is not a file, it should
notify the user and exit *with an error code*. It will iterate over the lines in the
file and print the line number, a space, and the line of the file. Your output will
differ from regular `cat -n` as I won't expect you to right-align the numbers.

```
$ ./cat_n.sh
Usage: cat-n.sh FILE
$ ./cat_n.sh foo
foo is not a file
$ ./cat_n.sh files/sonnet-29.txt
1 Sonnet 29
2 William Shakespeare
3
4 When, in disgrace with fortune and men's eyes,
5 I all alone beweep my outcast state,
6 And trouble deaf heaven with my bootless cries,
7 And look upon myself and curse my fate,
8 Wishing me like to one more rich in hope,
9 Featured like him, like him with friends possessed,
10 Desiring this man's art and that man's scope,
11 With what I most enjoy contented least;
12 Yet in these thoughts myself almost despising,
13 Haply I think on thee, and then my state,
14 (Like to the lark at break of day arising
15 From sullen earth) sings hymns at heaven's gate;
16 For thy sweet love remembered such wealth brings
17 That then I scorn to change my state with kings.
```

Exercise: hello.sh

Create a bash script called `hello.sh` that accepts one or two arguments. If there are no arguments, it should print a "Usage" and exit *with an error code.* Your program will expect to receive a "greeting" in `$1` and possibly a name in `$2`; if there is no second argument, use "Human" as the default. If there are more than two arguments, print a "Usage" and exit *with an error code.* Print the greeting, a comma and space, the name, and an exclamation point.

```
$ ./hello.sh
Usage: hello.sh GREETING [NAME]
$ ./hello.sh That\'ll do pig
Usage: hello.sh GREETING [NAME]
$ ./hello.sh "That'll do" pig
That'll do, pig!
$ ./hello.sh "Top o' the morning"
Top o' the morning, Human!
$ ./hello.sh "Greetings" "Earthling"
Greetings, Earthling!
```

Exercise: gap.sh

Write a bash script called `gap.sh` that will print out the files in the `gapminder` directory. Note that to be portable for testing purposes, you will need to use a **relative** path from the directory where the script lives (hint: start with `$PWD`). If there are no arguments, print out all the *basenames* of the files in sorted order. If there is an argument, treat it like a regular expression and find files where the basename matches at the beginning of the string in a case-insensitive manner and print them in sorted order. If no files are found, print a message telling the user.

```
$ ./gap.sh | head -5
     1  Afghanistan
     2  Albania
     3  Algeria
     4  Angola
     5  Argentina
$ ./gap.sh l
     1  Lebanon
     2  Lesotho
     3  Liberia
     4  Libya
$ ./gap.sh [w-z]
     1  West_Bank_and_Gaza
     2  Yemen_Rep
     3  Zambia
```

```
     4  Zimbabwe
$ ./gap.sh x
There are no countries starting with "x"
```

Chapter 3

Minimally Competent bash Scripting

> "We build our computer (systems) the way we build our cities: over time, without a plan, on top of ruins." - Ellen Ullman

> "The most effective debugging tool is still careful thought, coupled with judiciously placed print statements." - Brian W. Kernighan, Unix for Beginners (1979)

Bash is the worst shell scripting language except for all the others. For many of the analyses you'll write, all you will need is a simple bash script, so let's figure out how to write a decent one. I'll share with you what I've found to be the minimal amount of bash I use.

Statements

All programming language have a grammar where "statements" (like "sentences") are built up from other terms. Some languages like Python and Haskell use whitespace to figure out the end of a "statement," which is usually just the right side of the window. C-like languages such as bash and Perl define the end of a statement with a colon ;. Bash is interesting because it uses both. If you hit <Enter> or type a newline in your code, Bash will execute that statement. If you want to put several commands on one line, you can separate each with a semicolon. If you want to stretch a command over more than one line, you can use a backslash \ to continue the line:

```
$ echo Hi
Hi
$ echo Hello
Hello
$ echo Hi; echo Hello
Hi
Hello
$ echo \
> Hi
Hi
```

Comments

Every language has a way to indicate text in the source code that should not be executed by the program. Many Unix/c-style languages use the **#** (hash)

sign to indicate that any text to the right should be ignored by the language, but some languages use other characters or character combinations like // in Javascript, Java, and Rust. Programmers may use comments to explain what some particularly bit of code is doing, or they may use the characters to temporarily disable some section of code. Here is an example of what you might see:

```
# cf. https://en.wikipedia.org/wiki/Factorial
sub fac(n) {
  # first check terminal condition
  if (n <= 1) {
    return 1
  }
  # no? let's recurse!
  else {
    n * fac(n - 1) # the number times one less the number
  }
}
```

It's worth investing time in an editor that can easily comment/uncomment whole sections of code. For instance, in vim, I have a function that will add or removed the appropriate comment character(s) (depending on the filetype) from the beginning of the selected section. If your editor can't do that (e.g., nano), then I suggest you find something more powerful.

Shebang

Scripting languages (sh, bash, Perl, Python, Ruby, etc.) are generally distinguished by the fact that the "program" is a regular file containing plain text that is interpreted into machine code *at the time you run it*. Other languages (c, C++, Java, Haskell, Rust) have a separate compilation step to turn their regular text source files into a binary executable. If you view a compiled file with an editor/pager, you'll see a mess that might even lock up your window. (If that happens, refer back to "Make it stop!" to kill it or just close the window and start over.)

So, basically a "script" is a plain text file that is often executable by virtue of having the executable bit(s) turned on (cf. "Permissions"). It does not have to be executable, however. It's acceptable to put some commands in a file and simply tell the appropriate program to interpret the file:

```
$ echo "echo Hello, World" > hello.sh
$ sh hello.sh
Hello, World
```

But it looks cooler to do this:

```
$ chmod +x hello.sh
$ ./hello.sh
Hello, World
```

But what's going on here?

```
$ echo 'print("Hello, World")' > hello.py
$ chmod +x hello.py
$ ./hello.py
./hello.py: line 1: syntax error near unexpected token `"Hello, World"'
./hello.py: line 1: `print("Hello, World")'
```

We put some Python code in a file and then asked our shell (which is bash) to interpret it. That didn't work. If we ask Python to run it, everything is fine:

```
$ python3 hello.py
Hello, World
```

So we just need to let the shell know that this is Python 3 code, and that is what the "shebang" (see "Pronunciations") line is for. It looks like a comment, but it's special line that the shell uses to interpret the script. I'll use an editor to add a shebang to the "hello.py" script, then I'll `cat` the file so you can see what it looks like.

```
$ cat hello.py
#!/usr/bin/env python3
print("Hello, World")
$ ./hello.py
Hello, World
```

Often the shebang line will indicate the absolute path to a program like "/bin/bash" or "/usr/local/bin/gawk," but here I used an absolute path not to Python but to the "env" program which I then passed "python3" as the argument. Why did I do that? To make this script "portable" (for certain values of "portable," cf. "It's easier to port a shell than a shell script." – Larry Wall), I prefer to use the "python3" that is found by the environment as I will usually put my preferred Python first in my `$PATH`.

Let's Make A Script!

Let's make our script say "Hello" to some people:

```
$ cat -n hello2.sh
     1  #!/usr/bin/env bash
     2
     3  NAME="Newman"
     4  echo "Hello," $NAME
     5  NAME="Jerry"
```

```
     6    echo "Hello, $NAME"
$ ./hello2.sh
Hello, Newman
Hello, Jerry
```

I've created a variable called `NAME` to hold the string "Newman" and print it. Notice there is no `$` when assigning to the variable, only when you use it. The value of `NAME` can be changed at any time. You can print it out like on line 4 as it's own argument to `echo` or inside of a string like on line 6. Notice that the version on line 4 puts a space between the arguments to `echo`.

Because all the variables from the environment (see `env`) are uppercase (e.g., `$HOME` and `$USER`), I tend to use all-caps myself, but this did lead to a problem once when I named a variable `PATH` and then overwrote the actual `PATH` and then my program stopped working entirely as it could no longer find any of the programs it needed. Just remember that everything in Unix is case-sensitive, so `$Name` is an entirely different variable from `$name`.

When assigning a variable, you can have NO SPACES around the `=` sign:

```
$ NAME1="Doge"
$ echo "Such $NAME1"
Such Doge
$ NAME2 = "Doge"
-bash: NAME2: command not found
$ echo "Such $NAME2"
Such
```

Sidebar: Catching Common Errors (set -u)

Bash is an easy language to write incorrectly. One step you can take to ensure you don't misspell variables is to add `set -u` at the top of your script. E.g., if you type `echo $HOEM` on the command line, you'll get no output or warning that you misspelled the `$HOME` variable unless you `set -u`:

```
$ echo $HOEM

$ set -u
$ echo $HOEM
-bash: HOEM: unbound variable
```

This command tells bash to complain when you use a variable that was never initialized to some value. This is like putting on your helmet. It's not a requirement (depending on which state you live in), but you absolutely should do this because there might come a day when you misspell a variable. Note that this will not save you from as error like this:

```
$ cat -n set-u-bug1.sh
```

```
     1    #!/bin/bash
     2
     3    set -u
     4
     5    if [[ $# -gt 0 ]]; then
     6      echo $THIS_IS_A_BUG; # never initialized
     7    fi
     8
     9    echo "OK";
$ ./set-u-bug1.sh
OK
$ ./set-u-bug1.sh foo
./set-u-bug1.sh: line 6: THIS_IS_A_BUG: unbound variable
```

You can see that the first execution of the script ran just fine. There is a bug on line 6, but bash didn't catch it because that line did not execute. On the second run, the error occurred, and the script blew up. (FWIW, this is a problem in Python, too.)

Here's another pernicious error:

```
$ cat -n set-u-bug2.sh
     1    #!/bin/bash
     2
     3    set -u
     4
     5    GREETING="Hi"
     6    if [[ $# -gt 0 ]]; then
     7      GRETING=$1 # misspelled
     8    fi
     9
    10    echo $GREETING
$ ./set-u-bug2.sh
Hi
$ ./set-u-bug2.sh Hello
Hi
```

We were foolishly hoping that `set -u` would prevent us from misspelling the `$GREETING`, but at line 7 we simple created a new variable called `$GRETING`. Perhaps you were hoping for more help from your language? This is why we try to limit how much bash we write.

NB: I highly recommend you use the program `shellcheck` https://www.shellcheck.net/to find errors in your bash code.

Common Patterns

This is a cut-and-paste section for you. The idea is that I will describe many common patterns that you can use directly.

Test if a variable is a file or directory

Use the `-f` or `-d` functions to test if a variable identifies a "file" or a "directory," respectively

```
if [[ -f "$ARG" ]]; then
    echo "$ARG is a file"
fi

if [[ -d "$ARG" ]]; then
    echo "$ARG is a directory"
fi
```

Use `!` to negate this:

```
if [[ ! -f "$ARG" ]]; then
    echo "$ARG is NOT a file"
fi

if [[ ! -d "$ARG" ]]; then
    echo "$ARG is a directory"
fi
```

There are many other test you can use. See `man test` for a complete list. The `-s` is handy to see if a file is empty. You can use more than one test at a time with the `&&` ("and") or `||` ("or") operator.

```
if [[ -f "$ARG" ]] && [[ -s "$ARG" ]]; then
    echo "$ARG is a file and is not empty"
fi

if [[ ! -f "$ARG" ]] || [[ ! -s "$ARG" ]]; then
    echo "$ARG is a NOT file or is empty"
fi
```

Exit your script

The `exit` function will cease all operations and immediately exit. With no argument, it will use "0" which means "zero errors"; any other value is considered a error code, so `exit 1` is commonly used indicate some unspecified error.

```
if [[ -f "$ARG" ]]; then
    wc -l "$ARG"
    exit
else
    echo "$ARG must be a file"
    exit 1
fi
```

Check the number of arguments to your program

The first argument to your script is in `$1`, the second in `$2`, and so on. The number of arguments is in `$#`, so you can check the number like this:

```
if [[ $# -eq 0 ]]; then
    echo "Usage: foo.sh ARG"
    exit 1
fi
```

The `-eq` means "equal". You can also use `-gt` or `-gte` for "greater than (or equal)" and `-lt` or `-lte` for "less than or equal".

Put the arguments into named variables

You should assign `$1` and `$2` to names that have some meaning in your program.

```
INPUT_FILE=$1
NUM_ITERATIONS=$2
```

Set default values for optional arguments

If an argument is not needed, you can assign a default value. Here we can set `NUM_ITERATIONS` to have a default value of "10":

```
INPUT_FILE=$1
NUM_ITERATIONS=${2:-10}
```

Read a file

It's common to use a `while` loop to `read` a file, line-by-line, into some `VARIABLE`. Don't use a `$` on the `while` line (assigning), do use it when you want to interpolate it:

```
while read -r LINE; do
    echo "$LINE"
done < "$FILE"
```

Use a counter variable

It's common to use the variable `i` (for "integer" maybe?) as a temporary counter, e.g., iterating over lines in a file. The syntax to increment is clunky. This will print a line number and a line of text from a file:

```
i=0
while read -r LINE; do
    i=$((i+1))
    echo $i "$LINE"
done < "$FILE"
```

Loop operations

Use `continue` to skip to the next iteration of a loop. This will print only the even lines of a file:

```
i=0
while read -r LINE; do
    i=$((i+1))
    if [[ $(expr $i % 2) -eq 0 ]]; then
        continue
    else:
        echo "$i $LINE"
    fi
done < "$FILE"
```

Use `break` to leave a loop. This will print the first 10 lines of a file:

```
i=0
while read -r LINE; do
    echo "$LINE"
    i=$((i+1))
    if [[ $i -eq 10 ]]; then
        break
    fi
done < "$FILE"
```

Capture the output of a command

Historically `bash` used backticks (the same key as the tilde on a US QWERTY keyboard) to execute a command and put the results into a variable:

```
DIR=`ls`
```

Most people now use `$()` as it stands out much better:

```
DIR=$(ls)
LINES=$(grep foo bar.txt)
```

Count the number of lines in a file

```
NUM_LINES=$(wc -l "$FILE" | awk '{print $1}')

if [[ $NUM_LINES -lt 1 ]]; then
    echo "There is noting in $FILE"
    exit 1
fi
```

Get a temporary file or directory

Sometimes you need a temporary file to store something. If the name and location of the file is unimportant, use `mktemp` to get a temporary file or `mktemp -d` to get a temporary directory.

```
TMP_FILE=$(mktemp)
cat "foo\nbar\n" > "$TMP_FILE"
```

```
TMP_DIR=$(mktemp -d)
cd "$TMP_DIR"
```

Get the last part of a file or directory name

If you have "/path/to/my/file.txt" and you just want to print "file.txt", use `basename`:

```
FILE="/path/to/my/file.txt"
basename "$FILE"
```

Or put that into a variable name to use:

```
FILE="/path/to/my/file.txt"
BASENAME=$(basename "$FILE")
echo "Basename is $BASENAME"
```

Similary `dirname` is use to get "/path/to/my" from the above:

```
FILE="/path/to/my/file.txt"
DIRNAME=$(dirname "$FILE")
echo "Dirname is $DIRNAME"
```

Print with echo and printf

The `echo` command will print messages to the screen (standard out):

```
USER="Dave"
echo "I'm sorry, $USER, I can't do that."
```

The `printf` command is useful for formatting the output. The command expects a "template" first and then all the arguments for each formatting code in the template. The percent sign `%` is used in the template to indicate the type and options, e.g., an integer right-justified and three digits wide is `%3d`. Use `man printf` to learn more. Here is an example to print the line numbers in a file more prettier:

```
i=0
while read -r LINE; do
    i=$((i+1))
    printf "%3d: %s\n" $i $LINE
done < "$FILE"
```

Capture many items into a file for looping

Bash doesn't do lists (many items in a series) very well, so I usually put lists into files; e.g., I want to find how many files are in a directory and iterate over them:

```
FILES=$(mktemp)
find "$DIR" -type f -name \*.f[aq] > "$FILES"
NUM_FILES=$(wc -l "$FILES" | awk '{print $1}')

if [[ $NUM_FILES -lt 1 ]]; then
    echo "No usable files in $DIR"
    exit 1
fi

echo "Found $NUM_FILES in $DIR"

i=0
while read -r FILENAME; do
    i=$((i+1))
    BASENAME=$(basename "$FILENAME")
    printf "%3d: %s\n" $i "$FILENAME"
done < "$FILES"
```

For Loops

Often we want to do some set of actions for all the files in a directory or all the identifiers in a file. You can use a `for` loop to iterate over the values in some command that returns a list of results:

```
$ for FILE in *.sh; do echo "FILE = $FILE"; done
FILE = args.sh
FILE = args2.sh
FILE = args3.sh
FILE = basic.sh
FILE = hello.sh
FILE = hello2.sh
FILE = hello3.sh
FILE = hello4.sh
FILE = hello5.sh
FILE = hello6.sh
FILE = named.sh
FILE = positional.sh
FILE = positional2.sh
FILE = positional3.sh
FILE = set-u-bug1.sh
FILE = set-u-bug2.sh
```

Here it is in a script:

```
$ cat -n for.sh
     1    #!/bin/bash
     2
     3    set -u
     4
     5    DIR=${1:-$PWD}
     6
     7    if [[ ! -d "$DIR" ]]; then
     8        echo "$DIR is not a directory"
     9        exit 1
    10    fi
    11
    12    i=0
    13    for FILE in $DIR/*; do
    14        let i++
    15        printf "%3d: %s\n" $i "$FILE"
    16    done
```

On line 5, I default `DIR` to the current working directory which I can find with the environmental variable `$PWD` (print working directory). I check on line 7 that the argument is actually a directory with the `-d` test (`man test`). The rest

should look familiar. Here it is in action:

```
$ ./for.sh | head
  1: /Users/kyclark/work/metagenomics-book/bash/args.sh
  2: /Users/kyclark/work/metagenomics-book/bash/args2.sh
  3: /Users/kyclark/work/metagenomics-book/bash/args3.sh
  4: /Users/kyclark/work/metagenomics-book/bash/basic.sh
  5: /Users/kyclark/work/metagenomics-book/bash/config1.sh
  6: /Users/kyclark/work/metagenomics-book/bash/config2.sh
  7: /Users/kyclark/work/metagenomics-book/bash/count-fa.sh
  8: /Users/kyclark/work/metagenomics-book/bash/for-read-file.sh
  9: /Users/kyclark/work/metagenomics-book/bash/for.sh
 10: /Users/kyclark/work/metagenomics-book/bash/functions.sh
$ ./for.sh ../problems | head
  1: ../problems/cat-n
  2: ../problems/common-words
  3: ../problems/dna
  4: ../problems/gapminder
  5: ../problems/gc
  6: ../problems/greeting
  7: ../problems/hamming
  8: ../problems/hello
  9: ../problems/proteins
 10: ../problems/tac
```

You will see many examples of using `for` to read from a file like so:

```
$ cat -n for-read-file.sh
     1    #!/usr/bin/env bash
     2
     3    FILE=${1:-'srr.txt'}
     4    for LINE in $(cat "$FILE"); do
     5        echo "LINE \"$LINE\""
     6    done
$ cat srr.txt
SRR3115965
SRR516222
SRR919365
$ ./for-read-file.sh srr.txt
LINE "SRR3115965"
LINE "SRR516222"
LINE "SRR919365"
```

But that can break badly when the file contains more than one "word" per line (as defined by the `$IFS` [input field separator]):

```
$ column -t pov-meta.tab
name             lat_lon.ll
```

```
GD.Spr.C.8m.fa      -17.92522,146.14295
GF.Spr.C.9m.fa      -16.9207,145.9965833
L.Spr.C.1000m.fa    48.6495,-126.66434
L.Spr.C.10m.fa      48.6495,-126.66434
L.Spr.C.1300m.fa    48.6495,-126.66434
L.Spr.C.500m.fa     48.6495,-126.66434
L.Spr.I.1000m.fa    48.96917,-130.67033
L.Spr.I.10m.fa      48.96917,-130.67033
L.Spr.I.2000m.fa    48.96917,-130.67033
$ ./for-read-file.sh pov-meta.tab
LINE "name"
LINE "lat_lon.ll"
LINE "GD.Spr.C.8m.fa"
LINE "-17.92522,146.14295"
LINE "GF.Spr.C.9m.fa"
LINE "-16.9207,145.9965833"
LINE "L.Spr.C.1000m.fa"
LINE "48.6495,-126.66434"
LINE "L.Spr.C.10m.fa"
LINE "48.6495,-126.66434"
LINE "L.Spr.C.1300m.fa"
LINE "48.6495,-126.66434"
LINE "L.Spr.C.500m.fa"
LINE "48.6495,-126.66434"
LINE "L.Spr.I.1000m.fa"
LINE "48.96917,-130.67033"
LINE "L.Spr.I.10m.fa"
LINE "48.96917,-130.67033"
LINE "L.Spr.I.2000m.fa"
LINE "48.96917,-130.67033"
```

While Loops

The proper way to read a file line-by-line is with `while`:

```
$ cat -n while.sh
     1    #!/usr/bin/env bash
     2
     3    FILE=${1:-'srr.txt'}
     4    while read -r LINE; do
     5        echo "LINE \"$LINE\""
     6    done < "$FILE"
$ ./while.sh srr.txt
LINE "SRR3115965"
```

```
LINE "SRR516222"
LINE "SRR919365"
$ ./while.sh meta.tab
LINE "GD.Spr.C.8m.fa     -17.92522,146.14295"
LINE "GF.Spr.C.9m.fa     -16.9207,145.9965833"
LINE "L.Spr.C.1000m.fa    48.6495,-126.66434"
LINE "L.Spr.C.10m.fa     48.6495,-126.66434"
LINE "L.Spr.C.1300m.fa    48.6495,-126.66434"
LINE "L.Spr.C.500m.fa    48.6495,-126.66434"
LINE "L.Spr.I.1000m.fa    48.96917,-130.67033"
LINE "L.Spr.I.10m.fa     48.96917,-130.67033"
LINE "L.Spr.I.2000m.fa    48.96917,-130.67033"
```

Another advantage is that `while` can break the line into fields:

```
$ cat -n while2.sh
     1    #!/usr/bin/env bash
     2
     3    FILE='meta.tab'
     4    while read -r SITE LOC; do
     5        echo "$SITE is located at \"$LOC\""
     6    done < "$FILE"
$ ./while2.sh
GD.Spr.C.8m.fa is located at "-17.92522,146.14295"
GF.Spr.C.9m.fa is located at "-16.9207,145.9965833"
L.Spr.C.1000m.fa is located at "48.6495,-126.66434"
L.Spr.C.10m.fa is located at "48.6495,-126.66434"
L.Spr.C.1300m.fa is located at "48.6495,-126.66434"
L.Spr.C.500m.fa is located at "48.6495,-126.66434"
L.Spr.I.1000m.fa is located at "48.96917,-130.67033"
L.Spr.I.10m.fa is located at "48.96917,-130.67033"
L.Spr.I.2000m.fa is located at "48.96917,-130.67033"
```

Sidebar: Saving Function Results in Files

Often I want to iterate over the results of some calculation. Here is an example
of saving the results of an operation (`find`) into a temporary file:

```
$ cat -n count-fa.sh
     1    #!/usr/bin/env bash
     2
     3    set -u
     4
     5    if [[ $# -ne 1 ]]; then
     6        printf "Usage: %s DIR\n" "$(basename "$0")"
     7        exit 1
```

```
     8    fi
     9
    10    DIR=$1
    11    TMP=$(mktemp)
    12    find "$DIR" -type f -name \*.fa > "$TMP"
    13    NUM_FILES=$(wc -l "$TMP" | awk '{print $1}')
    14
    15    if [[ $NUM_FILES -lt 1 ]]; then
    16        echo "Found no .fa files in $DIR"
    17        exit 1
    18    fi
    19
    20    NUM_SEQS=0
    21    while read -r FILE; do
    22        NUM_SEQ=$(grep -c '^>' "$FILE")
    23        NUM_SEQS=$((NUM_SEQS + NUM_SEQ))
    24        printf "%10d %s\n" "$NUM_SEQ" "$(basename "$FILE")"
    25    done < "$TMP"
    26
    27    rm "$TMP"
    28
    29    echo "Done, found $NUM_SEQS sequences in $NUM_FILES files."
$ ./count-fa.sh ../problems
        23 anthrax.fa
         9 burk.fa
Done, found 32 sequences in 2 files.
```

Line 11 uses the `mktemp` function to give us the name of a temporary file, then I `find` all the files ending in ".fa" or ".fasta" and put that into the temporary file. I could them to make sure I found something. Then I read from the tempfile and use the `FILE` name to count the number of times I see a greater-than sign at the beginning of a line.

Getting Data Into Your Program: Arguments

We would like to get the NAME from the user rather than having it hardcoded in the script. I'll show you three ways our script can take in data from outside:

1. Command-line arguments, both positional (i.e., the first one, the second one, etc.) or named (e.g., `-n NAME`)
2. The environment
3. Reading a configuration file

First we'll cover the command-line arguments which are available through a few variables:

- `$#`: The number (think "#" == number) of arguments
- `$@`: All the arguments in a single string
- `$0`: The name of the script
- `$1, $2`: The first argument, the second argument, etc.

A la:

```
$ cat -n args.sh
     1    #!/usr/bin/env bash
     2
     3    echo "Num of args     : \"$#\""
     4    echo "String of args : \"$@\""
     5    echo "Name of program: \"$0\""
     6    echo "First arg       : \"$1\""
     7    echo "Second arg      : \"$2\""
$ ./args.sh
Num of args     : "0"
String of args : ""
Name of program: "./args.sh"
First arg       : ""
Second arg      : ""
$ ./args.sh foo
Num of args     : "1"
String of args : "foo"
Name of program: "./args.sh"
First arg       : "foo"
Second arg      : ""
$ ./args.sh foo bar
Num of args     : "2"
String of args : "foo bar"
Name of program: "./args.sh"
First arg       : "foo"
Second arg      : "bar"
```

If you would like to iterate over all the arguments, you can use `$@` like so:

```
$ cat -n args2.sh
     1    #!/usr/bin/env bash
     2
     3    if [[ $# -lt 1 ]]; then
     4        echo "There are no arguments"
     5    else
     6        i=0
     7        for ARG in "$@"; do
     8            let i++
     9            echo "$i: $ARG"
    10        done
    11    fi
```

```
$ ./args2.sh
There are no arguments
$ ./args2.sh foo
1: foo
$ ./args2.sh foo bar "baz quux"
1: foo
2: bar
3: baz quux
```

Here I'm throwing in a conditional at line 3 to check if the script has any arguments. If the number of arguments (`$#`) is less than (`-lt`) 1, then let the user know there is nothing to show; otherwise (`else`) do the next block of code. The `for` loop on line 7 works by splitting the argument string (`$@`) on spaces just like the command line does. Both `for` and `while` loops require the `do/done` pair to delineate the block of code (some languages use `{}`, Haskell and Python use only indentation). Along those lines, line 11 is the close of the `if` – "if" spell backwards; the close of a `case` statement in bash is `esac`.

The other bit of magic I threw in was a counter variable (which I always use lowercase `i` ["integer"], `j` if I needed an inner-counter and so on) which is initialized to "0" on line 6. I increment it, I could have written `$i=$(($i + 1))`, but it's easier to use the `let i++` shorthand. Lastly, notice that "baz quux" seen as a single argument because it was placed in quotes; otherwise arguments are separated by spaces.

Sidebar: Make It Pretty (or else)

Note that indentation doesn't matter as the program below works, but, honestly, which one is easier for you to read?

```
$ cat -n args3.sh
     1	#!/usr/bin/env bash
     2
     3	if [[ $# -lt 1 ]]; then
     4	echo "There are no arguments"
     5	else
     6	i=0
     7	for ARG in "$@"; do
     8	let i++
     9	echo "$i: $ARG"
    10	done
    11	fi
$ ./args3.sh foo bar
1: foo
2: bar
```

48

Our First Argument

AT LAST, let's return to our "hello" script!

```
$ cat -n hello3.sh
     1    #!/usr/bin/env bash
     2
     3    echo "Hello, $1!"
$ ./hello3.sh Captain
Hello, Captain!
```

This should make perfect sense now. We are simply saying "hello" to the first argument, but what happens if we provide no arguments?

```
$ ./hello3.sh
Hello, !
```

Checking the Number of Arguments

Well, that looks bad. We should check that the script has the proper number of arguments which is 1:

```
$ cat -n hello4.sh
     1    #!/usr/bin/env bash
     2
     3    if [[ $# -ne 1 ]]; then
     4        printf "Usage: %s NAME\n" "$(basename "$0")"
     5        exit 1
     6    fi
     7
     8    echo "Hello, $1!"
$ ./hello4.sh
Usage: hello4.sh NAME
$ ./hello4.sh Captain
Hello, Captain!
$ ./hello4.sh Captain Picard
Usage: hello4.sh NAME
```

Line 3 checks if the number of arguments is not equal (`-ne`) to 1 and prints a help message to indicate proper "usage." Importantly, it also will `exit` the program with a value which is not zero to indicate that there was an error. (NB: An exit value of "0" indicates 0 errors.) Line 4 uses `printf` rather than `echo` so I can do some fancy substitution so that the results of calling the `basename` function on the `$0` (name of the program) is inserted at the location of the `%s` (a string value, cf. man pages for "printf" and "basename").

Here is an alternate way to write this script:

```
$ cat -n hello5.sh
     1    #!/usr/bin/env bash
     2
     3    if [[ $# -eq 1 ]]; then
     4        NAME=$1
     5        echo "Hello, $NAME!"
     6    else
     7        printf "Usage: %s NAME\n" "$(basename "$0")"
     8        exit 1
     9    fi
```

Here I check on line 3 if there is just one argument, and the `else` is devoted
to handling the error; however, I prefer to check for all possible errors at the
beginning and `exit` the program quickly. This also has the effect of keeping my
code as far left on the page as possible.

Sidebar: Saving Function Results

In the previous script, you may have noticed `$(basename "$0")`. I was passing
the script name (`$0`) to the function `basename` and then passing that to the
`printf` function. To call a function in bash and save the results into a variable
or use the results as an argument, we can use either backticks (") (under the ~
on a US keyboard) or `$()`. I find backticks to be too similar to single quotes, so
I prefer the latter. To demonstrate:

```
$ ls | head
args.sh*
args2.sh*
args3.sh*
basic.sh*
hello.sh*
hello2.sh*
hello3.sh*
hello4.sh*
hello5.sh*
hello6.sh*
$ FILES=`ls | head`
$ echo $FILES
args.sh args2.sh args3.sh basic.sh hello.sh hello2.sh hello3.sh hello4.sh hello5.sh hello6.s
```

Here is a script that shows:

1. Calling `basename` and having the result print out (line 5)
2. Using `$()` to capture the results of `basename` into a variable (line 8)
3. Using `$()` to call `basename` as the second argument to `echo`
4. Showing that `$()` can be interpolated **inside a string**
5. Using `$()` to call `basename` as an argument to `printf`

50

```
$ cat -n functions.sh
     1    #!/usr/bin/env bash
     2
     3    # call function
     4    echo -n "1: BASENAME: "
     5    basename "$0"
     6
     7    # put function results into variable
     8    BASENAME=$(basename "$0")
     9    echo "2: BASENAME: $BASENAME"
    10
    11    # use results of function as argument to another function
    12    echo "3: BASENAME:" "$(basename "$0")"
    13    echo "4: BASENAME: $(basename "$0")"
    14    printf "5: BASENAME: %s\n" "$(basename "$0")"
$ ./functions.sh
1: BASENAME: functions.sh
2: BASENAME: functions.sh
3: BASENAME: functions.sh
4: BASENAME: functions.sh
5: BASENAME: functions.sh
```

Providing Default Argument Values

Here is how you can provide a default value for an argument with `:-`:

```
$ cat -n hello6.sh
     1    #!/usr/bin/env bash
     2
     3    echo "Hello, ${1:-Stranger}!"
$ ./hello6.sh
Hello, Stranger!
$ ./hello6.sh Govnuh
Hello, Govnuh!
```

Arguments From The Environment

You can also use look in the environment for argument values. For instance, we could accept the NAME as either the first argument to the script ($1) or the $USER from the environment:

```
$ cat -n hello7.sh
     1    #!/usr/bin/env bash
     2
     3    NAME=${1:-$USER}
```

51

```
    4    [[ -z "$NAME" ]] && NAME='Stranger'
    5    echo "Hello, $NAME
$ ./hello7.sh
Hello, kyclark
$ ./hello7.sh Barbara
Hello, Barbara
```

What's interesting is that you can temporarily over-ride an environmental variable like so:

```
$ USER=Bart ./hello7.sh
Hello, Bart
$ ./hello7.sh
Hello, kyclark
```

Exporting Values to the Environment

Notice that I can set `USER` for the first run to "Bart," but the value returns to "kyclark" on the next run. I can permanently set a value in the environment by using the `export` command. Here is a version of the script that looks for an environmental variable called `WHOM` (please do override your `$USER` name in the environment as things will break):

```
$ cat -n hello8.sh
    1    #!/usr/bin/env bash
    2
    3    echo "Hello, ${WHOM:-Marie}"
$ ./hello8.sh
Hello, Marie
```

As before I can set it temporarily:

```
$ WHOM=Doris ./hello8.sh
Hello, Doris
$ ./hello8.sh
Hello, Marie
```

Now I will `export` `WHOM` so that it persists:

```
$ WHOM=Doris
$ export WHOM
$ ./hello8.sh
Hello, Doris
$ ./hello8.sh
Hello, Doris
```

To remove `WHOM` from the environment, use `unset`:

```
$ unset WHOM
```

```
$ ./hello8.sh
Hello, Marie
```

Some programs rely heavily on environmental variables (e.g., Centrifuge, TACC LAUNCHER) for arguments. Here is a short script to illustrate how you would use such a program:

```
$ cat -n hello9.sh
     1    #!/usr/bin/env bash
     2
     3    WHOM="Who's on first" ./hello8.sh
     4    WHOM="What's on second"
     5    export WHOM
     6    ./hello8.sh
     7    WHOM="I don't know's on third" ./hello8.sh
$ ./hello9.sh
Hello, Who's on first
Hello, What's on second
Hello, I don't know's on third
```

Required and Optional Arguments

Now we're going to accept two arguments, "GREETING" and "NAME" while providing defaults for both:

```
$ cat -n positional.sh
     1    #!/usr/bin/env bash
     2
     3    set -u
     4
     5    GREETING=${1:-Hello}
     6    NAME=${2:-Stranger}
     7
     8    echo "$GREETING, $NAME"
$ ./positional.sh
Hello, Stranger
$ ./positional.sh Howdy
Howdy, Stranger
$ ./positional.sh Howdy Padnuh
Howdy, Padnuh
$ ./positional.sh "" Pahnuh
Hello, Pahnuh
```

You notice that if I want to use the default argument for the greeting, I have to pass an empty string "".

What if I want to require at least one argument?

```
$ cat -n positional2.sh
     1    #!/usr/bin/env bash
     2
     3    set -u
     4
     5    if [[ $# -lt 1 ]]; then
     6        printf "Usage: %s GREETING [NAME]\n" "$(basename "$0")"
     7        exit 1
     8    fi
     9
    10    GREETING=$1
    11    NAME=${2:-Stranger}
    12
    13    echo "$GREETING, $NAME"
$ ./positional2.sh "Good Day"
Good Day, Stranger
$ ./positional2.sh "Good Day" "Kind Sir"
Good Day, Kind Sir
```

It's also important to note the subtle hints given to the user in the "Usage"
statement. [NAME] has square brackets to indicate that it is an option, but
GREETING does not to say it is required. As noted before I wanted to use the
GREETING "Good Day," so I had to put it in quotes so that the shell would
not interpret them as two arguments. Same with the NAME "Kind Sir."

```
$ ./positional2.sh Good Day Kind Sir
Good, Day
```

Not Too Few, Not Too Many (Goldilocks)

Hmm, maybe we should detect that the script had too many arguments?

```
$ cat -n positional3.sh
     1    #!/usr/bin/env bash
     2
     3    set -u
     4
     5    if [[ $# -lt 1 ]] || [[ $# -gt 2 ]]; then
     6        printf "Usage: %s GREETING [NAME]\n" "$(basename "$0")"
     7        exit 1
     8    fi
     9
    10    GREETING=$1
    11    NAME=${2:-Stranger}
    12
    13    printf "%s, %s\n" "$GREETING" "$NAME"
```

```
$ ./positional3.sh Good Day Kind Sir
Usage: positional3.sh GREETING [NAME]
$ ./positional3.sh "Good Day" "Kind Sir"
Good Day, Kind Sir
```

To check for too many arguments, I added an "OR" (the double pipes `||`) and another conditional ("AND" is `&&`). I also changed line 13 to use a `printf` command to highlight the importance of quoting the arguments *inside the script* so that bash won't get confused. Try it without those quotes and try to figure out why it's doing what it's doing. I highly recommend using the program "shellcheck" (https://github.com/koalaman/shellcheck) to find mistakes like this. Also, consider using more powerful/helpful/sane languages – but that's for another discussion.

Named Arguments To The Rescue

I hope maybe by this point you're thinking that the script is getting awfully complicated just to allow for a combination of required an optional arguments all given in a particular order. You can manage with 1-3 positional arguments, but, after that, we really need to have named arguments and/or flags to indicate how we want to run the program. A named argument might be `-f mouse.fa` to indicate the value for the `-f` ("file," probably) argument is "mouse.fa," whereas a flag like `-v` might be a yes/no ("Boolean," if you like) indicator that we do or do not want "verbose" mode. You've encountered these with programs like `ls -l` to indicate you want the "long" directory listing or `ps -u $USER` to indicate the value for `-u` is the `$USER`.

The best thing about named arguments is that they can be provided in any order:

```
$ ./named.sh -n Patch -g "Good Boy"
Good Boy, Patch!
```

Some may have values, some may be flags, and you can easily provide good defaults to make it easy for the user to provide the bare minimum information to run your program. Here is a version that has named arguments:

```
$ cat -n named.sh
     1	#!/usr/bin/env ash
     2	
     3	set -u
     4	
     5	GREETING=""
     6	NAME="Stranger"
     7	EXCITED=0
     8	
     9	function USAGE() {
```

55

```
10          printf "Usage:\n  %s -g GREETING [-e] [-n NAME]\n\n" $(basename $0)
11          echo "Required arguments:"
12          echo " -g GREETING"
13          echo
14          echo "Options:"
15          echo " -n NAME ($NAME)"
16          echo " -e Print exclamation mark (default yes)"
17          echo
18          exit ${1:-0}
19      }
20
21      [[ $# -eq 0 ]] && USAGE 1
22
23      while getopts :g:n:eh OPT; do
24        case $OPT in
25          h)
26            USAGE
27            ;;
28          e)
29            EXCITED=1
30            ;;
31          g)
32            GREETING="$OPTARG"
33            ;;
34          n)
35            NAME="$OPTARG"
36            ;;
37          :)
38            echo "Error: Option -$OPTARG requires an argument."
39            exit 1
40            ;;
41          \?)
42            echo "Error: Invalid option: -${OPTARG:-""}"
43            exit 1
44        esac
45      done
46
47      [[ -z "$GREETING" ]] && USAGE 1
48      PUNCTUATION="."
49      [[ $EXCITED -ne 0 ]] && PUNCTUATION="!"
50
51      echo "$GREETING, $NAME$PUNCTUATION"
```

When run without arguments or with the **-h** flag, it produces a help message.

```
$ ./named.sh
Usage:
```

```
   named.sh -g GREETING [-e] [-n NAME]

Required arguments:
 -g GREETING

Options:
 -n NAME (Stranger)
 -e Print exclamation mark (default yes)
```

Our script just got much longer but also more flexible. I've written a hundred shell scripts with just this as the template, so you can, too. Go search for how `getopt` works and copy-paste this for your bash scripts, but the important thing to understand about `getopt` is that flags that take arguments have a `:` after them (`g:` == "-g something") and ones that do not, well, do not (`h` == "-h" == "please show me the help page). Both the"h" and "e" arguments are flags:

```
$ ./named.sh -n Patch -g "Good Boy"
Good Boy, Patch.
$ ./named.sh -n Patch -g "Good Boy" -e
Good Boy, Patch!
```

I've introduced a new function called `USAGE` that prints out the "Usage" statement so that it can be called when:

- the script is run with no arguments (line 21)
- the script is run with the "-h" flag (lines 25-26)
- the script is run with bad input (line 47)

I initialized the NAME to "Stranger" (line 6) and then let the user know in the "Usage" what the default value will be. When checking the GREETING in line 44, I'm actually checking that the length of the value is greater than zero because it's possible to run the script like this:

```
$ ./named01.sh -g ""
```

Which would technically pass muster but does not actually meet our requirements.

Reading a Configuration File

The last way I'll show you to get data into your program is to read a configuration file. This builds on the earlier example of using `export` to put values into the environment:

```
$ cat -n config1.sh
     1	    export NAME="Merry Boy"
     2	    export GREETING="Good morning"
$ cat -n read-config.sh
     1	    #!/usr/bin/env bash
     2
```

```
     3     source config1.sh
     4     echo "$GREETING, $NAME!"
$ ./read-config.sh
Good morning, Merry Boy!
```

To make this more flexible, let's pass the config file as an argument:

```
$ cat -n read-config2.sh
     1     #!/usr/bin/env bash
     2
     3     CONFIG=${1:-config1.sh}
     4     if [[ ! -f "$CONFIG" ]]; then
     5         echo "Bad config \"$CONFIG\""
     6         exit 1
     7     fi
     8
     9     source $CONFIG
    10     echo "$GREETING, $NAME!"
$ ./read-config2.sh
Good morning, Merry Boy!
$ cat -n config2.sh
     1     export NAME="François"
     2     export GREETING="Salut"
$ ./read-config2.sh config2.sh
Salut, François!
$ ./read-config2.sh foo
Bad config "foo"
```

A Full Bag of Tricks

Lastly I'm going to show you how to create some sane defaults, make missing
directories, find user input, transform that input, and report back to the user.
Here's a script that takes an IN_DIR, counts the lines of all the files therein,
and reports said line counts into an optional OUT_DIR.

```
     1     #!/usr/bin/env bash
     2
     3     set -u
     4
     5     IN_DIR=""
     6     OUT_DIR="$PWD/$(basename "$0" '.sh')-out"
     7
     8     function lc() {
     9         wc -l "$1" | awk '{print $1}'
    10     }
```

```
11
12    function USAGE() {
13        printf "Usage:\n  %s -i IN_DIR -o OUT_DIR\n\n" "$(basename "$0")"
14
15        echo "Required arguments:"
16        echo " -i IN_DIR"
17        echo "Options:"
18        echo " -o OUT_DIR"
19        echo
20        exit "${1:-0}"
21    }
22
23    [[ $# -eq 0 ]] && USAGE 1
24
25    while getopts :i:o:h OPT; do
26        case $OPT in
27            h)
28                USAGE
29                ;;
30            i)
31                IN_DIR="$OPTARG"
32                ;;
33            o)
34                OUT_DIR="$OPTARG"
35                ;;
36            :)
37                echo "Error: Option -$OPTARG requires an argument."
38                exit 1
39                ;;
40            \?)
41                echo "Error: Invalid option: -${OPTARG:-""}"
42                exit 1
43        esac
44    done
45
46    if [[ -z "$IN_DIR" ]]; then
47        echo "IN_DIR is required"
48        exit 1
49    fi
50
51    if [[ ! -d "$IN_DIR" ]]; then
52        echo "IN_DIR \"$IN_DIR\" is not a directory."
53        exit 1
54    fi
55
56    echo "Started $(date)"
```

```
57
58    FILES_LIST=$(mktemp)
59    find "$IN_DIR" -type f -name \*.sh > "$FILES_LIST"
60    NUM_FILES=$(lc "$FILES_LIST")
61
62    if [[ $NUM_FILES -gt 0 ]]; then
63        echo "Will process NUM_FILES \"$NUM_FILES\""
64
65        [[ ! -d $OUT_DIR ]] && mkdir -p "$OUT_DIR"
66
67        i=0
68        while read -r FILE; do
69            BASENAME=$(basename "$FILE")
70            let i++
71            printf "%3d: %s\n" $i "$BASENAME"
72            wc -l "$FILE" > "$OUT_DIR/$BASENAME"
73        done < "$FILES_LIST"
74
75        rm "$FILES_LIST"
76        echo "See results in OUT_DIR \"$OUT_DIR\""
77    else
78        echo "No files found in \"$IN_DIR\""
79    fi
80
81    echo "Finished $(date)"
```

The IN_DIR argument is required (lines 46-49), and it must be a directory (lines 51-54). If the user does not supply an OUT_DIR, I will create a reasonable default using the current working directory and the name of the script plus "-out" (line 6). One thing I love about bash is that I can call functions inside of strings, so OUT_DIR is a string (it's in double quotes) of the variable $PWD, the character "/", and the result of the function call to `basename` where I'm giving the optional second argument ".sh" that I want removed from the first argument, and then the string "-out".

At line 58, I create a temporary file to hold the names of the files I need to process. A line 59, I look for the files in IN_DIR that need to be processed. You can read the manpage for `find` and think about what your script might need to find (".fa" files greater than 0 bytes in size last modified since some date, etc.). At line 60, I call my `lc` (line count) function to see how many files I found. If I found more than 0 files (line 62), then I move ahead with processing. I check to see if the OUT_DIR needs to be created (line 65), and then create a counter variable ("i") that I'll use to number the files as I process them. At line 68, I start a `while` loop to iterate over the input from redirecting *in* from the temporary file holding the file names (line 73, `< "$FILES_LIST"`). Then a `printf` to let the user know which file we're processing, then a simple command (`wc`) but where you might choose to BLAST the sequence file to a database of

pathogens to determine how deadly the sample is. When I'm done, I clean up the temp file (line 75).

The alternate path when I find no input files (line 77-79) is to report that fact. Bracketing the main processing logic are "Started/Finished" statements so I can see how long my script took. When you start your coding career, you will usually sit and watch your code run before you, but eventually you'll submit the your jobs to an HPC queue where they will be run for you on a separate machine when the resources become available.

The above is, I would say, a minimally competent bash script. If you can understand everything in there, then you know enough to be dangerous and should move on to learning more powerful languages – like Python!

Chapter 4

Introduction to Python

> "Any fool can write code that a computer can understand. Good programmers write code that humans can understand." - Martin Fowler

Hello

Let's use our familiar "Hello, World!" to get started:

```
$ cat -n hello.py
     1  #!/usr/bin/env python3
     2
     3  print('Hello, World!')
```

The first thing to notice is the "shebang" on line 1. I'm going to use `env` to find the first `python3` binary in the user's `$PATH`. In `bash`, we could use either `echo` or `printf` to print to the terminal (or a file). In Python, we have `print()` noting that we must use parentheses now to invoke functions.

Variables

It's not so interesting to just say "Hello, World!" all the time. Let's make a program that will say "hello" to some value that we pass in. This value can change each time we run the program, so it's common to call this a "variable."

Let's use the REPL (Read-Evaluate-Print-Loop, pronounced "reh-pul") to play with variables. Type `python` (or `python3` or `ipython`) to get into a REPL:

```
>>> name = 'Duderino'
>>> print('Hello,', name)
Hello, Duderino
```

Here I'm showing that we can create variable called `name` by assigning it some value like "Duderino." Unlike `bash`, we don't have to worry about spaces around the =. You can put any number of spaces around the equal sign, but it's most common (and readable) to put just one on each side. Notice that `print` will accept more than one argument and will put spaces between the arguments. You can tell it to use some other "separator" by indicating the `sep` keyword argument. Notice the Pythonic style is that there are *no* spaces around the = for keywords:

```
>>> print('Hello', name, sep=', ')
Hello, Duderino
```

It's not easy to tell, but `print` is also putting a newline on the end. We can change that with the `end` keyword argument:

```
>>> print('Hello', name, sep=', ', end='!')
Hello, Duderino!>>>
```

Unlike in `bash`, we cannot use a variable directly in a `print` statement or we get the equivalent of George Burns telling Gracie "Say 'Good night,' Gracie" and she says "Good night, Gracie!":

```
>>> print('Hello, name')
Hello, name
```

We could to use the `+` operator to concatenate it to the literal string "Hello,":

```
>>> print('Hello, ' + name)
Hello, Duderino
```

Types: Strings and Numbers

As you might expect, the "plus" operator `+` is also used to perform numeric addition:

```
>>> n = 10
>>> n + 1
11
```

The `name` variable above is of the type `str` (string) because we put the value in quotes (single or double, it doesn't matter).

```
>>> name = 'The Dude'
>>> type(name)
<class 'str'>
```

Numbers don't have quotes. Number can be integers (`int`) or floating-point numbers (`float`) if they have a decimal somewhere or you write them in scientific notation:

```
>>> type(10)
<class 'int'>
>>> type(3.14)
<class 'float'>
>>> type(1.)
<class 'float'>
>>> type(2.864e-10)
<class 'float'>
```

The "plus" operator behaves completely differently with different *types* of arguments as long as the arguments are both strings or both numbers. Things go wobbly when you mix them:

```
>>> 'Hello, ' + 'Mr. Lebowski'
'Hello, Mr. Lebowski'
>>> 1 + 2
3
>>> 1 + 'Mr. Lebowski'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: unsupported operand type(s) for +: 'int' and 'str'
```

Lists

Before we go further, I will introduce a different variable type called a "list" as we are going to need that immediately. You create a list by putting values in `[]` (square brackets or just "brackets"):

```
>>> vals = ['foo', 'bar', 'baz']
>>> vals
['foo', 'bar', 'baz']
```

You can get the length of a list using the `len` function:

```
>>> len(vals)
3
```

But note that, like so many other languages, Python starts counting at `0`, so the first element is in the "zeroth" position. The value at position `1` is actually the *second* value.

```
>>> vals[0]
'foo'
>>> vals[1]
'bar'
```

We'll talk much more about lists in the next chapter. I needed to tell you that so I could tell you this next bit.

Command-line Arguments: sys.argv is a list

Now let's get our "hello" program to greet an argument passed from the command line. We discussed in the `bash` section that programs can take *positional* arguments, e.g. `ls` can accept the name of the directory you wish to list or `wc` can take the name of a file to count. *Positional* arguments mean the first

argument, the second argument, and so on. In the command `ls ~`, the `~` (tilde which means `$HOME` in `bash`) is the one and only positional argument. In the command `ls /bin /usr/bin/`, there are two positional arguments, `/bin` and `/usr/bin/`.

Named options have some sort of prefix, e.g., `find` can take a `-maxdepth` argument to indicate how many levels deep to search. Lastly, commands may also take flags like the `-l` flags to `ls` that indicates you wish to see the "long" listing.

To get access to the positional arguments to our program, we need to `import sys` which is a package of code that will interact with the system. Those arguments will be a *list*:

```
$ cat -n hello_arg.py
     1    #!/usr/bin/env python3
     2
     3    import sys
     4
     5    args = sys.argv
     6    print('Hello, ' + args[1] + '!')
```

From the `sys` module, we call the `argv` function to get the "argument vector." This is a list, and, like `bash`, the name of the script is the first argument (in the zeroth position) – `args[0]`. That means the first *actual* "argument" to the script is in `args[1]`.

```
$ ./hello_arg.py Geddy
Hello, Geddy!
```

But there is a problem if we fail to pass any arguments:

```
$ ./hello_arg.py
Traceback (most recent call last):
  File "./hello_arg.py", line 6, in <module>
    print('Hello, ' + args[1] + '!')
IndexError: list index out of range
```

We tried to access something in `args` that doesn't exist, and so the entire program came to a halt ("crashed"). As in `bash`, we need to check how many arguments we have. Before I show you how to do that, let me explain something about *slicing* lists. Inside the `[]`, you can indicate a start and stop positon like so:

```
>>> vals
['foo', 'bar', 'baz']
>>> vals[1:]
['bar', 'baz']
>>> vals[1:2]
['bar']
```

In the last example, you see that the stop position is not inclusive. Even though you've seen that Python gets very upset by asking for a position in a list that does not exist, it has no problem giving you nothing when you ask for a *slice* that doesn't exist:

```
>>> vals[1000:]
[]
```

So we can use that to ask for `sys.argv[1:]` to get all the *actual* arguments to our program, skipping over the name of the program itself:

```
$ cat -n hello_arg2.py
     1  #!/usr/bin/env python3
     2
     3  import sys
     4
     5  args = sys.argv[1:]
     6
     7  if len(args) < 1:
     8      print('Usage:', sys.argv[0], 'NAME')
     9      sys.exit(1)
    10
    11  name = args[0]
    12  print('Hello, ' + name + '!')
```

If there are fewer than 1 argument, then we print a usage statement and use `sys.exit` to send the operating system a non-zero exit status, just like in `bash`. It works much better now:

```
$ ./hello_arg2.py
Usage: ./hello_arg2.py NAME
$ ./hello_arg2.py Alex
Hello, Alex!
```

Here is the same functionality but using some new functions, `str.format` so we can introduce a different way to join strings, and `os.path.basename` so we can get the name of the program without any leading path information like `./`:

```
$ cat -n hello_arg3.py
     1  #!/usr/bin/env python3
     2
     3  import sys
     4  import os
     5
     6  args = sys.argv[1:]
     7
     8  if len(args) != 1:
     9      script = os.path.basename(sys.argv[0])
    10      print('Usage: {} NAME'.format(script))
```

```
    11        sys.exit(1)
    12
    13  name = args[0]
    14  print('Hello, {}!'.format(name))
$ ./hello_arg3.py
Usage: hello_arg3.py NAME
$ ./hello_arg3.py Neil
Hello, Neil!
```

The main() thing

Many languages (e.g., Perl, Rust, Haskell) have the idea of a `main` module/function where all the processing starts. If you define a `main` function using `def main`, most people reading your code would understand that the program *ought* to begin there. (I say "ought" because Python won't actually make that happen. You still have to *call* the `main` function to make your program run!) I usually put my `main` first and then call it at the end of the script with this `__name__ == '__main__'` business. This looks a bit of a hack, but it is fairly Pythonic.

```
$ cat -n hello_arg4.py
     1  #!/usr/bin/env python3
     2
     3  import sys
     4  import os
     5
     6  def main():
     7      args = sys.argv[1:]
     8
     9      if len(args) != 1:
    10          script = os.path.basename(sys.argv[0])
    11          print('Usage: {} NAME'.format(script))
    12          sys.exit(1)
    13
    14      name = args[0]
    15      print('Hello, {}!'.format(name))
    16
    17
    18  if __name__ == '__main__':
    19      main()
$ ./hello_arg4.py
Usage: hello_arg4.py NAME
$ ./hello_arg4.py '2013 Rock and Roll Hall of Fame Inductees'
Hello, 2013 Rock and Roll Hall of Fame Inductees!
```

Function Order

Note that you cannot put call to `main()` before `def main` because you cannot call a function that hasn't been defined (lexically) in the program yet. To add insult to injury, this is a **run-time error** – meaning the mistake isn't caught by the compiler when the program is parsed into byte-code; instead the program just crashes.

```
$ cat -n func-def-order.py
     1    #!/usr/bin/env python3
     2
     3    print('Starting the program')
     4    foo()
     5    print('Ending the program')
     6
     7    def foo():
     8        print('This is foo')
$ ./func-def-order.py
Starting the program
Traceback (most recent call last):
  File "./func-def-order.py", line 4, in <module>
    foo()
NameError: name 'foo' is not defined
```

To contrast:

```
$ cat -n func-def-order2.py
     1    #!/usr/bin/env python3
     2
     3    def foo():
     4        print('This is foo')
     5
     6    print('Starting the program')
     7    foo()
     8    print('Ending the program')
$ ./func-def-order2.py
Starting the program
This is foo
Ending the program
```

Handle All The Args!

If we like, we can greet to any number of arguments:

```
$ cat -n hello_arg5.py
     1  #!/usr/bin/env python3
```

```
    2
    3  import sys
    4  import os
    5
    6  def main():
    7      names = sys.argv[1:]
    8
    9      if len(names) < 1:
   10          script = os.path.basename(sys.argv[0])
   11          print('Usage: {} NAME [NAME2 ...]'.format(script))
   12          sys.exit(1)
   13
   14      print('Hello, {}!'.format(', '.join(names)))
   15
   16  if __name__ == '__main__':
   17      main()
$ ./hello_arg5.py
Usage: hello_arg5.py NAME [NAME2 ...]
$ ./hello_arg5.py Geddy Alex Neil
Hello, Geddy, Alex, Neil!
```

Notice on line 14 to see how we can `join` all the arguments on a comma + space.

Conditionals

So far we've been using an `if` condition to see if we have enough arguments. If you want to test for more than one condition, you can use `elif` (else if) and `else` ("otherwise" or the "default" branch if all others fail). Here we'll use the `input` function to present the user with a prompt and get their input:

```
$ cat -n if-else.py
    1  #!/usr/bin/env python3
    2
    3  name = input('What is your name? ')
    4  age = int(input('Hi, ' + name + '. What is your age? '))
    5
    6  if age < 0:
    7      print("That isn't possible.")
    8  elif age < 18:
    9      print('You are a minor.')
   10  else:
   11      print('You are an adult.')
$ ./if-else.py
What is your name? Ken
Hi, Ken. What is your age? -4
```

```
That isn't possible.
$ ./if-else.py
What is your name? Lincoln
Hi, Lincoln. What is your age? 29
You are an adult.
```

On line 3, we can put the first answer directly into the **name** variable; however, on line 4, I need to convert the answer to an integer with **int** because I will need to compare it numerically, cf:

```
>>> 4 < 5
True
>>> '4' < 5
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: unorderable types: str() < int()
>>> int('4') < 5
True
```

Things go very badly if we blindly try to coerce a string into an **int**:

```
$ ./if-else.py
What is your name? Doreen
Hi, Doreen. What is your age? Ageless
Traceback (most recent call last):
  File "./if-else.py", line 4, in <module>
    age = int(input('Hi, ' + name + '. What is your age? '))
ValueError: invalid literal for int() with base 10: 'Ageless'
```

Later we'll talk about how to avoid problems like this.

Loops

As in bash, we can use **for** loops in Python. Here's another way to greet all the people:

```
$ cat -n hello_arg6.py
     1  #!/usr/bin/env python3
     2
     3  import sys
     4  import os
     5
     6  def main():
     7      names = sys.argv[1:]
     8
     9      if len(names) < 1:
    10          prg = os.path.basename(sys.argv[0])
```

```
11            print('Usage: {} NAME [NAME2 ...]'.format(prg))
12            sys.exit(1)
13
14       for name in names:
15            print('Hello, ' + name + '!')
16
17
18  if __name__ == '__main__':
19      main()
$ ./hello_arg6.py
Usage: hello_arg6.py NAME [NAME2 ...]
$ ./hello_arg6.py Salt Peppa
Hello, Salt!
Hello, Peppa!
```

You can use a `for` loop on anything that is like a list! A string is a list of characters:

```
>>> for letter in "abc":
...     print(letter)
...
a
b
c
```

The `range` function returns something that can be "iterated" like a list:

```
>>> for number in range(0, 5):
...     print(number)
...
0
1
2
3
4
```

Lists, of course:

```
>>> for word in ['foo', 'bar']:
...     print(word)
...
foo
bar
```

You can use the `str.split` function to split a string (the default is to split on spaces):

```
>>> for word in 'We hold these truths'.split():
...     print(word)
...
```

```
We
hold
these
truths
```

And we can use the `open` function to open a file and read each line using a `for` loop:

```
>>> for line in open('input1.txt'):
...     print(line, end='')
...
this is
some text
from a file.
```

The last example either needs to suppress the newline from `print` or do `rstrip()` on the line to remove it as the text coming from the file has a newline.

Stubbing New Programs with new_py.py

Every program we've seen so far has had the same basic structure:

- Shebang
- import modules
- define main()
- call main()

Additionally we keep having to write the same few lines of code to get the arguments from `sys.argv[1:]` and then test that we have the right number and then print a usage and `sys.exit(1)`. Rather than type all that boilerplate or copy-paste from other programs, let's use a program to help us create new programs.

Included in the `bin` directory of the GitHub repo, there is a program called `new_py.py` that will stub out all this code for you. Make sure you either add that directory to you `$PATH` or copy that program into your existing `$PATH`, e.g., I like to have `$HOME/.local/bin` for programs like this:

```
$ which new_py.py
/Users/kyclark/.local/bin/new_py.py
```

Now run it with no arguments. As you might expect, it gives you a usage statement:

```
$ new_py.py
usage: new_py.py [-h] [-a] [-f] program
new_py.py: error: the following arguments are required: program
```

Give it the name of a new program (either with or without `.py`):

```
$ new_py.py foo
Done, see new script "foo.py."
$ cat -n foo.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date   : 2019-05-14
     5  Purpose: Rock the Casbah
     6  """
     7
     8  import os
     9  import sys
    10
    11
    12  # --------------------------------------------------
    13  def main():
    14      args = sys.argv[1:]
    15
    16      if len(args) != 1:
    17          print('Usage: {} ARG'.format(os.path.basename(sys.argv[0])))
    18          sys.exit(1)
    19
    20      arg = args[0]
    21
    22      print('Arg is "{}"'.format(arg))
    23
    24
    25  # --------------------------------------------------
    26  if __name__ == '__main__':
    27      main()
```

What happens if you try to initialize a script when one already exists with that name?

```
$ new_py.py foo
"foo.py" exists.  Overwrite? [yN] n
Will not overwrite. Bye!
```

Unless you answer "y", the script will not be overwritten. You could also use the `-f|--force` flag to force the overwritting of an existing file.

Introducing argparse

Run `new_py.py` with `-h` or `--help` to see all the options:

```
$ new_py.py -h
```

```
usage: new_py.py [-h] [-a] [-f] program

Create Python script

positional arguments:
  program         Program name

optional arguments:
  -h, --help      show this help message and exit
  -a, --argparse  Use argparse (default: False)
  -f, --force     Overwrite existing (default: False)
```

Hey, what is `--argparse` about? Let's try it! I will combine the two short flag `-a` and `-f` into `-fa` to "force" a new script that uses the `argparse` module to give us positional arguments, named arguments, and flags.

```
$ new_py.py -fa foo
Done, see new script "foo.py."
$ cat -n foo.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date   : 2019-05-14
     5  Purpose: Rock the Casbah
     6  """
     7
     8  import argparse
     9  import sys
    10
    11
    12  # --------------------------------------------------
    13  def get_args():
    14      """get command-line arguments"""
    15      parser = argparse.ArgumentParser(
    16          description='Argparse Python script',
    17          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    18
    19      parser.add_argument(
    20          'positional', metavar='str', help='A positional argument')
    21
    22      parser.add_argument(
    23          '-a',
    24          '--arg',
    25          help='A named string argument',
    26          metavar='str',
    27          type=str,
    28          default='')
```

```
29
30      parser.add_argument(
31          '-i',
32          '--int',
33          help='A named integer argument',
34          metavar='int',
35          type=int,
36          default=0)
37
38      parser.add_argument(
39          '-f', '--flag', help='A boolean flag', action='store_true')
40
41      return parser.parse_args()
42
43
44  # -------------------------------------------------
45  def warn(msg):
46      """Print a message to STDERR"""
47      print(msg, file=sys.stderr)
48
49
50  # -------------------------------------------------
51  def die(msg='Something bad happened'):
52      """warn() and exit with error"""
53      warn(msg)
54      sys.exit(1)
55
56
57  # -------------------------------------------------
58  def main():
59      """Make a jazz noise here"""
60      args = get_args()
61      str_arg = args.arg
62      int_arg = args.int
63      flag_arg = args.flag
64      pos_arg = args.positional
65
66      print('str_arg = "{}"'.format(str_arg))
67      print('int_arg = "{}"'.format(int_arg))
68      print('flag_arg = "{}"'.format(flag_arg))
69      print('positional = "{}"'.format(pos_arg))
70
71
72  # -------------------------------------------------
73  if __name__ == '__main__':
74      main()
```

That's a good bit more code, but the advantage here is that we can now get quite detailed help documentation and very specific behavior from our arguments, e.g., one argument needs to be a string while another needs to be a number while another is a true/false flag. You can immediately run the program that was just created and see the usage:

```
$ ./foo.py -h
usage: foo.py [-h] [-a str] [-i int] [-f] str

Argparse Python script

positional arguments:
  str                 A positional argument

optional arguments:
  -h, --help          show this help message and exit
  -a str, --arg str   A named string argument (default: )
  -i int, --int int   A named integer argument (default: 0)
  -f, --flag          A boolean flag (default: False)
```

All this without writing a line of Python! Quite useful.

In my experience, perhaps 20-50% of the effort to solve most of the exercises can be handled by using `argparse` well. You can specify an exact number of positional arguments, you can specify named arguments that must be constrained to a list of choices, you can force one argument to be an `int` and another to be a `float`, you can get Boolean flags or ensure that arguments are existing files that can be opened and read. I urge you to read the documentation for `argparse` thoroughly. I often find the REPL is quite useful for this:

```
>>> import argparse
>>> help(argparse)
```

Exercises

Now is the time for you to write your own programs. Go into the `exercises` directory and solve `hello_a`. The way you are intended to work is to read the README, then use `new_py.py` to stub out the named program, then use `make test` to pass the tests.

Exercise: hello

If you `ls` the directory, you should see something like this:

```
$ ls
Makefile      README.md      README.pdf     solution.py* test.py*
```

The README says:

Write a Python program named `hello.py` that warmly greets the names you provide. When there are two names, join them with "and." When there are three or more, join them on commas (INCLUDING THE OXFORD WE ARE NOT SAVAGES) and "and." If no names are supplied, print a usage.

```
$ ./hello.py
Usage: hello.py NAME [NAME...]
$ ./hello.py Alice
Hello to the 1 of you: Alice!
$ ./hello.py Mike Carol
Hello to the 2 of you: Mike and Carol!
$ ./hello.py Greg Peter Bobby Marcia Jane Cindy
Hello to the 6 of you: Greg, Peter, Bobby, Marcia, Jane, and Cindy!
```

The name of the program should be `hello.py`, so do this:

```
$ new_py.py hello
Done, see new script "hello.py."
$ cat -n hello.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date   : 2019-05-14
     5  Purpose: Rock the Casbah
     6  """
     7
     8  import os
     9  import sys
    10
    11
    12  # --------------------------------------------------
    13  def main():
    14      args = sys.argv[1:]
    15
    16      if len(args) != 1:
    17          print('Usage: {} ARG'.format(os.path.basename(sys.argv[0])))
    18          sys.exit(1)
    19
    20      arg = args[0]
    21
    22      print('Arg is "{}"'.format(arg))
    23
    24
    25  # --------------------------------------------------
    26  if __name__ == '__main__':
    27      main()
```

Next you will run `make test`, but let's take a moment to understand what this command is doing. The `make` program will look for a file called `Makefile` (or `makefile`) and will then look for a "target" called `test`. Look at the `Makefile`:

```
$ cat -n Makefile
     1  .PHONY: doc test
     2
     3  doc:
     4      pandoc README.md -o README.pdf
     5
     6  test:
     7      pytest -v test.py
```

You see there is a line with `test:` that starts on the far left side of the file and there is some text indented *with a tab* underneath it. The indented bit will be run when we say `make test`, so you could just run `pytest -v test.py` directly if you wish or if your system doesn't have `make` installed. If you don't have `pytest` installed, then probably this will do it:

```
$ python3 -m pip install pytest
```

Or if you have the Anaconda distribution (highly recommended), then you could do this:

```
$ conda install pytest
```

Now, run `make test` and you should see that you are already passing one test – the usage! That is because `new_py.py` automatically adds that code for you. Of course, it's not accurate code:

```
$ ./hello.py
Usage: hello.py ARG
```

Our program is supposed to take multiple arguments, and we've been given the expected usage in the README, so first correct that.

Now look at the first failure:

```
=================================== FAILURES ===================================
_____ test_01 _____

    def test_01():
        """runs hello"""

        rv, out = getstatusoutput('{} {}'.format(prg, 'Alice'))
        assert rv == 0
>       assert out.rstrip() == 'Hello to the 1 of you: Alice!'
E       assert 'Arg is "Alice"' == 'Hello to the 1 of you: Alice!'
E         - Arg is "Alice"
E         + Hello to the 1 of you: Alice!
```

The lines starting with E are the errors we need to fix. It's telling us that the output from the program out should equal 'Hello to the 1 of you: Alice!' but instead it says 'Arg is "Alice"'. Confirm this is the case:

```
$ ./hello.py Alice
Arg is "Alice"
```

OK, go fix that.

```
$ ./hello.py Alice
Hello to the 1 of you: Alice!
```

And now run make test. You should find you now pass two tests! Here's the next failure:

```
================================== FAILURES ==================================
_____ test_02 _____

    def test_02():
        """runs hello"""

        rv, out = getstatusoutput('{} {}'.format(prg, ' Mike Carol'))
        assert rv == 0
>       assert out.rstrip() == 'Hello to the 2 of you: Mike and Carol!'
E       AssertionError: assert '' == 'Hello to the 2 of you: Mike and Carol!'
E           + Hello to the 2 of you: Mike and Carol!
```

Now when given the arguments "Mike" and "Carol", we need to fix the greeting. When you are passing that test, then move on to the next failure. In this way, you will eventually have a program that meets the specifications. At any point you can inspect the solution.py that I have provided if you get stuck or want to compare your version with what I wrote.

For what it's worth, this is the idea of "Test-Driven Development" which you can read about on the Internet.

Exercise: article.py

Make a Python program called article.py that prepends before the given "word" (it may not really be a word) the article "a" if the word begins with a consonant or "an" if it begins with a vowel. If given no arguments, the program should provide a usage statement.

```
$ ./article.py
Usage: article.py WORD
$ ./article.py foo
a foo
$ ./article.py oof
an oof
```

Chapter 5

Strings, Lists, and Tuples

"Good programming is good writing." - John Shore

There's some overlap among Python's strings, lists, and tuples. In a way, you could think of strings as lists of characters. Many list operations work exactly the same over strings like subscripting to get a particular item. We can ask for the first (or "zeroth") element from a string:

```
>>> name = 'Curly'
>>> name[0]
'C'
```

Or from a list:

```
>>> names = ['Larry', 'Moe', 'Curly', 'Shemp']
>>> names[0]
'Larry'
```

"Slice" operations let you take a range of items. Notice that we can operate on a string literal (in quotes):

```
>>> names[2:4]
['Curly', 'Shemp']
>>> 'Curly'[2:4]
'rl'
```

Functions like `join` that take lists can also work on strings:

```
>>> ', '.join(names)
'Larry, Moe, Curly, Shemp'
>>> ', '.join(names[0])
'L, a, r, r, y'
```

You can ask if a list contains a certain member, and you can also ask if a string contains a certain character or substring:

```
>>> 'Moe' in names
True
>>> 'r' in 'Larry'
True
>>> 'url' in 'Curly'
True
>>> 'x' in 'Larry'
False
>>> 'Joe' in names
False
```

You can iterate with a `for` loop over both the items in a list:

```
>>> names = ['Larry', 'Moe', 'Curly', 'Shemp']
>>> for name in names:
...     print(name)
...
Larry
Moe
Curly
Shemp
```

Or the characters in a word:

```
>>> for letter in 'Curly':
...     print(letter)
...
C
u
r
l
y
```

Just as in bash, we can create a counter, increment it inside our loop, and print the element number before the element:

```
>>> names = ['Larry', 'Moe', 'Curly', 'Shemp']
>>> i = 0
>>> for name in names:
...     i += 1
...     print(i, name)
...
1 Larry
2 Moe
3 Curly
4 Shemp
```

Because we so often want this behavior, there is a function called **enumerate** that takes a list/string and returns the index/position along with the item/character. Since it's so annoying to deal with zero-offset counting, we can tell **enumerate** to **start** at 1:

```
>>> names = ['Larry', 'Moe', 'Curly', 'Shemp']
>>> for i, name in enumerate(names, start=1):
...     print('{:3} {}'.format(i, name))
...
  1 Larry
  2 Moe
  3 Curly
  4 Shemp
```

```
>>> for i, letter in enumerate('Curly', start=1):
...     print('{:3} {}'.format(i, letter))
...
  1 C
  2 u
  3 r
  4 l
  5 y
```

You can turn a list around with the `reversed` function:

```
>>> reversed(names)
<list_reverseiterator object at 0x109e490f0>
```

What we have here is a failure to communicate. You expected to see the list of names in the reverse order, but what you got was a promise from Python to give you that list when you actually need it. In the REPL, we can use the `list` function:

```
>>> list(reversed(names))
['Shemp', 'Curly', 'Moe', 'Larry']
```

So you can also use that to reverse a word:

```
>>> list(reversed('cat'))
['t', 'a', 'c']
```

OK, well, I wanted the word "tac" and not the list of letters in "tac"! We can put them back into a word by calling the `join` *function* of the *string element* that we want to put between the letter (which is an empty string). Notice that I don't have to use `list` because the `join` function will iterate on the `reversed` result:

```
>>> ''.join(reversed('cat'))
'tac'
```

Identifying Palindromes

Hmm, quite a bit of work to turn a word around. Still, could be useful, for example in finding CRISPR (clustered regularly interspaced short palindromic repeats) sequences? Here is a simple program to determine if a given string is a palindrome which is a string that is the same forwards and backwards.

```
$ cat -n word_is_palindrome.py
     1  #!/usr/bin/env python3
     2  """Report if the given word is a palindrome"""
     3
     4  import sys
```

```
 5  import os
 6
 7
 8  def main():
 9      """main"""
10      args = sys.argv[1:]
11
12      if len(args) != 1:
13          print('Usage: {} STR'.format(os.path.basename(sys.argv[0])))
14          sys.exit(1)
15
16      word = args[0]
17      rev = ''.join(reversed(word))
18
19      print('"{}" is{} a palindrome.'.format(
20          word, '' if word.lower() == rev.lower() else ' NOT'))
21
22
23  if __name__ == '__main__':
24      main()
$ ./word_is_palindrome.py
Usage: word_is_palindrome.py STR
$ ./word_is_palindrome.py foobar
"foobar" is NOT a palindrome.
$ ./word_is_palindrome.py foof
"foof" is a palindrome.
```

Finding Palindromes in a File

We can expand our palindrome program to one that searches in a file:

```
$ cat -n find_palindromes.py
 1  #!/usr/bin/env python3
 2  """Report if the given word is a palindrome"""
 3
 4  import sys
 5  import os
 6
 7
 8  def main():
 9      """main"""
10      args = sys.argv[1:]
11      min_length = 2
12
```

```
13        if len(args) != 1:
14            print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
15            sys.exit(1)
16
17        file = args[0]
18
19        if not os.path.isfile(file):
20            print('"{}" is not a file'.format(file))
21            sys.exit(1)
22
23        for line in open(file):
24            for word in line.lower().split():
25                # skip short words
26                if len(word) > min_length:
27                    rev = ''.join(reversed(word))
28                    if rev == word:
29                        print(word)
30
31
32  if __name__ == '__main__':
33      main()
```

```
$ cat input.txt
My sister Anna was looking for a kayak around noon when her mom
suggested she might rather enjoy a new racecar instead.
$ ./find_palindromes.py
Usage: find_palindromes.py FILE
$ ./find_palindromes.py foo
"foo" is not a file
$ ./find_palindromes.py input.txt
anna
kayak
noon
mom
racecar
```

Here we've added the ability to take as an argument the name of a file. By include `import os`, we can ask the OS (operating system) if the argument we were given is the name of an existing file. If this returns `False`, then we print an error message and exit with a non-zero exit code. Then we can use a `for` loop to iterate over each line in the `open` file handle, then another `for` loop to iterate over each `word` in the result of lowercasing and splitting the line of text. We skip short words and then print the current `word` if it is a palindrome.

Tetranucleotide Composition

A common operation in bioinformatics is to determine sequence composition
and is the first problem (id = DNA) to solve at the rosalind.info website. Here is
a program to find the frequencies of the DNA bases (A, C, T, G):

```
$ cat -n dna1.py
     1  #!/usr/bin/env python3
     2  """Tetra-nucleotide counter"""
     3
     4  import sys
     5  import os
     6
     7
     8  def main():
     9      """main"""
    10      args = sys.argv[1:]
    11
    12      if len(args) != 1:
    13          print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    14          sys.exit(1)
    15
    16      dna = args[0]
    17      num_a, num_c, num_g, num_t = 0, 0, 0, 0
    18
    19      for base in dna:
    20          if base == 'a' or base == 'A':
    21              num_a += 1
    22          elif base == 'c' or base == 'C':
    23              num_c += 1
    24          elif base == 'g' or base == 'G':
    25              num_g += 1
    26          elif base == 't' or base == 'T':
    27              num_t += 1
    28
    29      print(' '.join([str(num_a), str(num_c), str(num_g), str(num_t)]))
    30
    31
    32  if __name__ == '__main__':
    33      main()
$ cat dna.txt
AGCTTTTCATTCTGACTGCAACGGGCAATATGTCTCTGTGTGGATTAAAAAAAAGAGTGTCTGATAGCAGC
$ ./dna1.py `cat dna.txt`
3 2 1 1
```

We initialize four variables to count each DNA base. Just as we can use a `for`

loop to iterate through a list, we can iterate through each base in the input DNA string, checking both the upper- and lowercase version of each base to determine which variable to increment.

To save typing, let's force the input sequence to lowercase:

```
$ cat -n dna2.py
     1  #!/usr/bin/env python3
     2  """Tetra-nucleotide counter"""
     3
     4  import sys
     5  import os
     6
     7  def main():
     8      """main"""
     9      args = sys.argv[1:]
    10
    11      if len(args) != 1:
    12          print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    13          sys.exit(1)
    14
    15      dna = args[0]
    16
    17      num_a, num_c, num_g, num_t = 0, 0, 0, 0
    18
    19      for base in dna.lower():
    20          if base == 'a':
    21              num_a += 1
    22          elif base == 'c':
    23              num_c += 1
    24          elif base == 'g':
    25              num_g += 1
    26          elif base == 't':
    27              num_t += 1
    28
    29      print(' '.join([str(num_a), str(num_c), str(num_g), str(num_t)]))
    30
    31  if __name__ == '__main__':
    32      main()
```

There are better ways than this to count the characters, but we'll save this until we talk about dictionaries.

Lastly, let's use the `format` method to get rid of those pesky `str` calls:

```
$ cat -n dna3.py
     1  #!/usr/bin/env python3
     2  """Tetra-nucleotide counter"""
```

```
 3
 4  import sys
 5  import os
 6
 7  def main():
 8      """main"""
 9      args = sys.argv[1:]
10
11      if len(args) != 1:
12          print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
13          sys.exit(1)
14
15      dna = args[0]
16
17      num_a, num_c, num_g, num_t = 0, 0, 0, 0
18
19      for base in dna.lower():
20          if base == 'a':
21              num_a += 1
22          elif base == 'c':
23              num_c += 1
24          elif base == 'g':
25              num_g += 1
26          elif base == 't':
27              num_t += 1
28
29      print('{} {} {} {}'.format(num_a, num_c, num_g, num_t))
30
31  if __name__ == '__main__':
32      main()
```

If you're having trouble seeing the differences from `dna2.py` to `dna3.py`, try using `diff`:

```
$ diff dna2.py dna3.py
29c29
<     print(' '.join([str(num_a), str(num_c), str(num_g), str(num_t)]))
---
>     print('{} {} {} {}'.format(num_a, num_c, num_g, num_t))
```

Run-length Encoding

Along the lines of counting characters in a string, we can write a very simple string compression program that encodes repetitions of characters. Note that this program shows how the argument can be a string or the name of a file.

```
$ cat -n run_length.py
     1  #!/usr/bin/env python3
     2  """Compress text/DNA by marking repeated letters"""
     3
     4  import os
     5  import sys
     6
     7
     8  def main():
     9      """main"""
    10      args = sys.argv[1:]
    11
    12      if len(args) != 1:
    13          print('Usage: {} ARG'.format(os.path.basename(sys.argv[0])))
    14          sys.exit(1)
    15
    16      # If the argument is a file, the text should be the file contents
    17      arg = args[0]
    18      text = ''
    19      if os.path.isfile(arg):
    20          text = ''.join(open(arg).read().split())
    21      else:
    22          text = arg.strip()
    23
    24      # Make sure we have something
    25      if len(text) == 0:
    26          print('No usable text')
    27          sys.exit(1)
    28
    29      counts = []
    30      count = 0
    31      prev = None
    32      for letter in text:
    33          # We are at the start
    34          if prev is None:
    35              prev = letter
    36              count = 1
    37          # This letter is the same as before
    38          elif letter == prev:
    39              count += 1
    40          # This is a new letter, so record the count
    41          # of the previous letter and reset the counter
    42          else:
    43              counts.append((prev, count))
    44              count = 1
    45              prev = letter
```

88

```
46
47          # get the last letter after we fell out of the loop
48          counts.append((prev, count))
49
50          for letter, num in counts:
51              print('{}{}'.format(letter, '' if num == 1 else num), end='')
52
53          print()
54
55
56   if __name__ == '__main__':
57          main()
$ ./run_length.py AAAAACCCGATT
A5C3GAT2
$ ./run_length.py dna.txt
AGCT4CAT2CTGACTGCA2CG3CA2TATGTCTCTGTGTG2AT2A7GAGTGTCTGATAGCAGC
```

This program makes use of a `counts` list to keep track of each letter we saw. We
add a "tuple" to the list:

```
>>> counts = []
>>> counts.append(('A', 3))
>>> counts
[('A', 3)]
>>> counts.append(('C', 1))
>>> counts
[('A', 3), ('C', 1)]
```

Tuples are similar to lists, but they are immutable:

```
>>> tup = ('white', 'dog')
>>> tup[1]
'dog'
>>> tup[1] = 'cat'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: 'tuple' object does not support item assignment
```

You see they are subscripted like strings and lists, but you cannot change a value
inside a tuple. Tuples are not limited to pairs:

```
>>> tup = ('white', 'dog', 'bird')
>>> tup[-1]
'bird'
```

tac (reverse a file)

We all know and love the venerable `cat` program, but do you know about `tac`?
It prints a file in reverse. We can use lists in Python to read a file into list and
`reverse` it:

```
$ cat -n tac1.py
     1  #!/usr/bin/env python3
     2  """print a file in reverse"""
     3
     4  import sys
     5  import os
     6
     7
     8  def main():
     9      """main"""
    10      args = sys.argv[1:]
    11      if len(args) != 1:
    12          print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
    13          sys.exit(1)
    14
    15      file = args[0]
    16      if not os.path.isfile(file):
    17          print('"{}" is not a file'.format(file))
    18          sys.exit(1)
    19
    20      lines = []
    21      for line in open(file):
    22          lines.append(line)
    23
    24      lines.reverse()
    25
    26      for line in lines:
    27          print(line, end='')
    28
    29
    30  if __name__ == '__main__':
    31      main()
$ cat quick.txt
The quick brown
fox jumps
over the
lazy dog.
$ ./tac1.py quick.txt
lazy dog.
over the
```

```
fox jumps
The quick brown
```

We initialize a new list on line 17, then read through the file line-by-line and call the `append` method to add the line to the end of our list. Then we call `reverse` *function* on the list to mutate the list **IN PLACE**:

```
>>> names
['Larry', 'Moe', 'Curly', 'Shemp']
>>> names.reverse()
>>> names
['Shemp', 'Curly', 'Moe', 'Larry']
```

After `reverse` we see that the `names` are permanently changed. We can put them back with another call:

```
>>> names.reverse()
>>> names
['Larry', 'Moe', 'Curly', 'Shemp']
```

If we had simply wanted to use them in a reversed order **WITHOUT ALTERING THE ACTUAL LIST**, we could call the `reversed` function:

```
>>> list(reversed(names))
['Shemp', 'Curly', 'Moe', 'Larry']
>>> names
['Larry', 'Moe', 'Curly', 'Shemp']
```

It's easy to read an entire file directly into a list with `readlines` (this preserves newlines), but you should be sure that you have at least as much memory on your machine as the file is big. Compare these various ways to read an entire file. `read` will give you the contents as one string, and newlines will be present to denote the end of each line:

```
>>> open('input.txt').read()
'first line\nsecond line\nthird line\nfourth line\n'
```

Whereas `readlines` will return a list of strings broken on the newlines (but not removing them):

```
>>> open('input.txt').readlines()
['first line\n', 'second line\n', 'third line\n', 'fourth line\n']
```

Calling `read().splitlines()` will suck in the whole file, then break on the newlines, removing them in the process:

```
>>> open('input.txt').read().splitlines()
['first line', 'second line', 'third line', 'fourth line']
```

Similarly, you can `read().split()` to break all the input on spaces to get the words:

```
>>> open('input.txt').read().split()
['first', 'line', 'second', 'line', 'third', 'line', 'fourth', 'line']
```

Here is a version that uses `readlines()`:

```
$ cat -n tac2.py
     1  #!/usr/bin/env python3
     2  """print a file in reverse"""
     3
     4  import sys
     5  import os
     6
     7
     8  def main():
     9      """main"""
    10      args = sys.argv[1:]
    11      if len(args) != 1:
    12          print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
    13          sys.exit(1)
    14
    15      file = args[0]
    16      if not os.path.isfile(file):
    17          print('"{}" is not a file'.format(file))
    18          sys.exit(1)
    19
    20      lines = open(file).readlines()
    21      lines.reverse()
    22
    23      for line in lines:
    24          print(line, end='')
    25
    26
    27  if __name__ == '__main__':
    28      main()
```

This version uses the `reversed` function:

```
$ cat -n tac3.py
     1  #!/usr/bin/env python3
     2  """print a file in reverse"""
     3
     4  import sys
     5  import os
     6
     7
     8  def main():
     9      """main"""
    10      args = sys.argv[1:]
```

```
11          if len(args) != 1:
12              print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
13              sys.exit(1)
14
15          file = args[0]
16          if not os.path.isfile(file):
17              print('"{}" is not a file'.format(file))
18              sys.exit(1)
19
20          lines = open(file).readlines()
21
22          for line in reversed(lines):
23              print(line, end='')
24
25
26  if __name__ == '__main__':
27      main()
```

And finally I will introduce the **with/open** convention that you will see in Python.
Notice that I will call the variable **fh** to stand for "file handle" which is not the
same as the file name (which is a string):c

```
$ cat -n tac4.py
     1  #!/usr/bin/env python3
     2  """print a file in reverse"""
     3
     4  import sys
     5  import os
     6
     7
     8  def main():
     9      """main"""
    10      args = sys.argv[1:]
    11      if len(args) != 1:
    12          print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
    13          sys.exit(1)
    14
    15      file = args[0]
    16      if not os.path.isfile(file):
    17          print('"{}" is not a file'.format(file))
    18          sys.exit(1)
    19
    20      with open(file) as fh:
    21          lines = fh.readlines()
    22          for line in reversed(lines):
    23              print(line, end='')
    24
```

```
    25
    26  if __name__ == '__main__':
    27      main()
```

Picnic

Here is a little memory game you might have played with your bored siblings on family car trips:

```
$ ./picnic.py
What are you bringing? [! to quit] chips
We'll have chips.
What else are you bringing? [! to quit] ham sammich
We'll have chips and ham sammich.
What else are you bringing? [! to quit] co-cola
We'll have chips, ham sammich, and co-cola.
What else are you bringing? [! to quit] cupcakes
We'll have chips, ham sammich, co-cola, and cupcakes.
What else are you bringing? [! to quit] apples
We'll have chips, ham sammich, co-cola, cupcakes, and apples.
What else are you bringing? [! to quit] !
Bye.
```

Each person introduces a new item, and the other person has to remember all the previous items and add a new one. This is a classic "stack" that can be implemented with lists and an infinite loop:

```
$ cat -n picnic.py
     1  #!/usr/bin/env python3
     2  """What are you bringing to the picnic?"""
     3
     4
     5  # --------------------------------------------------
     6  def joiner(items):
     7      """properly conjuct items"""
     8      num_items = len(items)
     9      if num_items == 0:
    10          return ''
    11      elif num_items == 1:
    12          return items[0]
    13      elif num_items == 2:
    14          return ' and '.join(items)
    15      else:
    16          items[-1] = 'and ' + items[-1]
    17          return ', '.join(items)
```

```
18
19
20  # ----------------------------------------------------
21  def main():
22      """main"""
23      items = []
24
25      while True:
26          item = input('What {}are you bringing? [! to quit] '.format(
27              'else ' if items else ''))
28          if item == '!':
29              break
30          elif len(item.strip()) > 0:
31              if item in items:
32                  print('You said "{}" already.'.format(item))
33              else:
34                  items.append(item)
35                  # pass a copy because it gets mutated!
36                  print("We'll have {}.".format(joiner(items.copy())))
37
38      print('Bye.')
39
40
41  # ----------------------------------------------------
42  if __name__ == '__main__':
43      main()
```

One bug that got me in writing this program was in calling my `joiner` with
the list of `items`. I was surprised to learn that in mutating the list (by adding
"and" before the last item when there are more than 2 items), I was mutating
the *original list*! I learned to pass `items.copy()` so as to work on a copy of the
data and not the actual list.

Insult Generator

Sometimes (especially when writing games) you may want a random selection
from a list of items. Here is an insult generator that draws from the fabulous
vocabulary of Shakespeare:

```
$ cat -n insult.py
     1  #!/usr/bin/env python3
     2  """Shakespearean insult generator"""
     3
     4  import sys
     5  import random
```

```
 6
 7  adjectives = """
 8  scurvy old filthy scurilous lascivious foolish rascaly gross rotten
 9  corrupt foul loathsome irksome heedless unmannered whoreson cullionly
10  false filthsome toad-spotted caterwauling wall-eyed insatiate vile
11  peevish infected sodden-witted lecherous ruinous indistinguishable
12  dishonest thin-faced slanderous bankrupt base detestable rotten
13  dishonest lubbery
14  """.strip().split()
15
16  nouns = """
17  knave coward liar swine villain beggar slave scold jolthead whore
18  barbermonger fishmonger carbuncle fiend traitor block ape braggart
19  jack milksop boy harpy recreant degenerate Judas butt cur Satan ass
20  coxcomb dandy gull minion ratcatcher maw fool rogue lunatic varlet
21  worm
22  """.strip().split()
23
24
25  def main():
26      """main"""
27      args = sys.argv[1:]
28      num = int(args[0]) if len(args) > 0 and args[0].isdigit() else 5
29
30      for _ in range(num):
31          adjs = [random.choice(adjectives) for _ in range(3)]
32          noun = random.choice(nouns)
33          print('You {} {}!'.format(', '.join(adjs), noun))
34
35
36  if __name__ == '__main__':
37      main()
$ ./insult.py
You bankrupt, cullionly, detestable milksop!
You foul, indistinguishable, false Satan!
You lascivious, scurilous, bankrupt villain!
You lascivious, lecherous, rotten jack!
You toad-spotted, base, foolish Satan!
$ ./insult.py 3
You detestable, cullionly, wall-eyed scold!
You peevish, caterwauling, caterwauling traitor!
You thin-faced, foul, dishonest Judas!
```

Notice how the program takes an optional argument that I expect to be an integer. On line 24, I test both that there is an argument present and that it isdigit() before attempting to use it as a number. The real work is done by

the `random.choice` function to grab my adjectives and noun. The `"""` operator lets us write strings with newlines, then we `split` the long string into words. This is a common idiom in Python.

Synthetic Biology

Lists could represent biological entities such as promotor, coding, and terminator regions. Let's say we wanted to design synthetic microbes where we tested all possible permutations of these regions with each other to see if we were able to increase production of a desired enzyme. Since the operation is N^3, I will only show the output for 2 genes:

```
$ ./recomb.py
Usage: recomb.py NUM_GENES
$ ./recomb.py foo
"foo" is not an integer
$ ./recomb.py 2
N = "2"
    1: P1 - C1 - T1
    2: P1 - C1 - T2
    3: P1 - C2 - T1
    4: P1 - C2 - T2
    5: P2 - C1 - T1
    6: P2 - C1 - T2
    7: P2 - C2 - T1
    8: P2 - C2 - T2
```

Here is the Python code. Note that I replaced repeated calls to `print/sys.exit(1)` with a new function I call `die`. I use the `product` function from `itertools` to get all the combinations of the promoter, coding, and terminator regions.

```
$ cat -n recomb.py
     1  #!/usr/bin/env python3
     2  """Show recominations"""
     3
     4  import os
     5  import sys
     6  from itertools import product
     7
     8  def die(msg):
     9      """print and exit with an error"""
    10      print(msg)
    11      sys.exit(1)
    12
```

```
13  def main():
14      """main"""
15      args = sys.argv[1:]
16
17      if len(args) != 1:
18          die('Usage: {} NUM_GENES'.format(os.path.basename(sys.argv[0])))
19
20      if not args[0].isdigit():
21          die('"{}" is not an integer'.format(args[0]))
22
23      num_genes = int(args[0])
24      min_num = 2
25      max_num = 10
26      if not min_num <= num_genes <= max_num:
27          die('NUM_GENES must be between {} and {}'.format(min_num, max_num))
28
29      promotors, coding, term = ([], [], [])
30      for n in [str(i) for i in range(1, num_genes + 1)]:
31          promotors.append('P' + n)
32          coding.append('C' + n)
33          term.append('T' + n)
34
35      print('N = "{}"'.format(num_genes))
36      for i, combo in enumerate(product(promotors, coding, term), start=1):
37          print('{:4}: {}'.format(i, ' - '.join(combo)))
38
39
40  if __name__ == '__main__':
41      main()
```

Line 30 is an example of a list comprehension. They read awkwardly because the
`for` loop that introduces the variable happens after you start seeing the variable:

```
>>> [n for n in range(1, 4)]
[1, 2, 3]
>>> [n ** n for n in range(1, 4)]
[1, 4, 27]
```

You can add a predicate for item selection to the end:

```
>>> [n ** 2 for n in range(1, 5) if n % 2 == 0]
[4, 16]
```

That `%` operator is the remainder after division, so if the remainder after dividing
a number by 2 is 0, then the number is even.

Here is the comprehensions in the program (lines 23-25):

```
$ cat -n recomb2.py
```

```
 1  #!/usr/bin/env python3
 2  """Show recominations"""
 3
 4  import os
 5  import sys
 6  from itertools import product
 7
 8
 9  def die(msg):
10      """print and exit with an error"""
11      print(msg)
12      sys.exit(1)
13
14
15  def main():
16      """main"""
17      args = sys.argv[1:]
18
19      if len(args) != 1:
20          die('Usage: {} NUM_GENES'.format(os.path.basename(sys.argv[0])))
21
22      if not args[0].isdigit():
23          die('"{}" does not look like an integer'.format(args[0]))
24
25      num_genes = int(args[0])
26      min_num = 2
27      max_num = 10
28      if not min_num <= num_genes <= max_num:
29          die('NUM_GENES must be between {} and {}'.format(min_num, max_num))
30
31      ranger = range(1, num_genes + 1)
32      promotors = ['P' + str(n + 1) for n in ranger]
33      coding = ['C' + str(n + 1) for n in ranger]
34      term = ['T' + str(n + 1) for n in ranger]
35
36      print('N = "{}"'.format(num_genes))
37      for i, combo in enumerate(product(promotors, coding, term), start=1):
38          print('{:4}: {}'.format(i + 1, ' - '.join(combo)))
39
40
41  if __name__ == '__main__':
42      main()
```

But lines 32-34 are identical with the exception of the character I'm using, so I can put that code into a little function:

```
$ cat -n recomb3.py
```

```
 1  #!/usr/bin/env python3
 2  """Show recominations"""
 3
 4  import os
 5  import sys
 6  from itertools import product
 7
 8
 9  def die(msg):
10      """print and exit with an error"""
11      print(msg)
12      sys.exit(1)
13
14
15  def main():
16      """main"""
17      args = sys.argv[1:]
18
19      if len(args) != 1:
20          die('Usage: {} NUM_GENES'.format(os.path.basename(sys.argv[0])))
21
22      if not args[0].isdigit():
23          die('"{}" does not look like an integer'.format(args[0]))
24
25      num_genes = int(args[0])
26      if not 2 <= num_genes <= 10:
27          die('NUM_GENES must be greater than 1, less than 10')
28
29      def gen(prefix):
30          return [prefix + str(n) for n in range(1, num_genes + 1)]
31
32      print('N = "{}"'.format(num_genes))
33      combos = product(gen('P'), gen('C'), gen('T'))
34      for i, combo in enumerate(combos, start=1):
35          print('{:4}: {}'.format(i, ' - '.join(combo)))
36
37
38  if __name__ == '__main__':
39      main()
```

Now all the repeated code is in the gen function, and I simply call that for each
character I want.

Exercise: head.py

Create a Python program called `head.py` that expects one or two arguments. If there are no arguments, print a "Usage" statement. The first argument is required and much be a regular file; if it is not, print " is not a file" and exit *with an error code.* The second argument is optional. If given, it must be a positive number (non-zero); if it is not, then print "lines () must be a positive number". If no argument is provided, use a default value of 3. You can expect that the test will only give you a value that can be safely converted to a number using the `int` function. If given good input, it should act like the normal `head` utility and print the expected number of lines from the given file.

```
$ ./head.py
Usage: head.py FILE [NUM_LINES]
$ ./head.py foo
foo is not a file
$ ./head.py files/issa.txt
Selected Haiku by Issa

Don't worry, spiders,
$ ./head.py files/issa.txt 5
Selected Haiku by Issa

Don't worry, spiders,
I keep house
casually.
```

Exercises

Exercise: cat_n.py

Create a Python program called `cat_n.py` that expects exactly one argument which is a regular file and prints usage statement if either condition fails. It should print each line of the file argument preceeded by the line number which is right-justified in spaces and a colon. You may the format '{:5}: {}' to make it look exactly like the output below, but the test is just checking for a leading space, some number(s), a colon, and the line of text.

```
$ ./cat_n.py
Usage: cat_n.py FILE
$ ./cat_n.py foo
foo is not a file
$ ./cat_n.py files/sonnet-29.txt
    1: Sonnet 29
    2: William Shakespeare
```

```
 3:
 4: When, in disgrace with fortune and men's eyes,
 5: I all alone beweep my outcast state,
 6: And trouble deaf heaven with my bootless cries,
 7: And look upon myself and curse my fate,
 8: Wishing me like to one more rich in hope,
 9: Featured like him, like him with friends possessed,
10: Desiring this man's art and that man's scope,
11: With what I most enjoy contented least;
12: Yet in these thoughts myself almost despising,
13: Haply I think on thee, and then my state,
14: (Like to the lark at break of day arising
15: From sullen earth) sings hymns at heaven's gate;
16: For thy sweet love remembered such wealth brings
17: That then I scorn to change my state with kings.
```

Exercise: gc.py

Write a Python program called `gc.py` that takes a single positional argument which should be a file. Die with a warning if the argument is not a file. For each line in the file, print the line number and the percentage of the characters on that line that are a "G" or "C" (case-insensitive).

```
$ ./gc.py
usage: gc.py [-h] FILE
gc.py: error: the following arguments are required: FILE
$ ./gc.py foo
"foo" is not a file
$ ./gc.py samples/sample1.txt
  1:   9%
  2:  19%
  3:  19%
  4:  22%
  5:  32%
  6:  21%
```

Exercise: vowel_counter.py

Write a Python program called `vowel_counter.py` that counts the number of vowels in a single string. Be sure your subject and verb agree in number, and use proper plurals.

```
$ ./vowel_counter.py
Usage: vowel_counter.py STRING
$ ./vowel_counter.py for
```

```
There is 1 vowel in "for."
$ ./vowel_counter.py elliptical
There are 4 vowels in "elliptical."
$ ./vowel_counter.py YYZ
There are 0 vowels in "YYZ."
```

Exercise: grid.py

Write a Python program called `grid.py` that will create a square grid from the
argument given on the command line. The program will expect one positional
argument; if the number of arguments is not exactly one, print a "usage" state-
ment and exit *with an error code*. The test suite will only provide integer values,
so you can assume it is safe to use `int` to convert the input from a string to an
integer. The number provided must be in the range of 2 to 9 (inclusive). If it is
not, print "NUM () must be between 1 and 9" and exit *with an error code*. You
will square the given number to create a grid (so think of the number as how
many rows and columns).

```
$ ./grid.py
Usage: grid.py NUM
$ ./grid.py 1
NUM (1) must be between 1 and 9
$ ./grid.py -1
NUM (-1) must be between 1 and 9
$ ./grid.py 10
NUM (10) must be between 1 and 9
$ ./grid.py 2
  1  2
  3  4
$ ./grid.py 5
  1  2  3  4  5
  6  7  8  9 10
 11 12 13 14 15
 16 17 18 19 20
 21 22 23 24 25
$ ./grid.py 8
  1  2  3  4  5  6  7  8
  9 10 11 12 13 14 15 16
 17 18 19 20 21 22 23 24
 25 26 27 28 29 30 31 32
 33 34 35 36 37 38 39 40
 41 42 43 44 45 46 47 48
 49 50 51 52 53 54 55 56
 57 58 59 60 61 62 63 64
```

Exercise: histy.py

Write a Python program called `histy.py` that takes one or more integer values
as positional argumements and creates a histogram of the values in their sorted
order using the `-c|--character` (default "|") repeated by the number. Also
accept a `-m|--minimum` option that is the minimum value to print a number as
well as a `-s|--scale` option that will be used to scale the number of characters
printed so that large numbers won't line wrap in your terminal.

```
$ ./histy.py
usage: histy.py [-h] [-c str] [-m int] [-s int] int [int ...]
histy.py: error: the following arguments are required: int
$ ./histy.py -h
usage: histy.py [-h] [-c str] [-m int] [-s int] int [int ...]

Histogrammer

positional arguments:
  int                   Inputs

optional arguments:
  -h, --help            show this help message and exit
  -c str, --character str
                        Character to represent (default: |)
  -m int, --minimum int
                        Minimum value to print (default: 1)
  -s int, --scale int   Scale inputs (default: 1)
$ ./histy.py 3 1 2
  1 |
  2 ||
  3 |||
$ ./histy.py 300 100 200 -s 100
100 |
200 ||
300 |||
$ ./histy.py 300 100 200 -s 100 -c '#'
100 #
200 ##
300 ###
$ ./histy.py 300 100 200 -s 100 -c '#' -m 150
200 ##
300 ##
```

Exercise: column.py

Create a Python program called `column.py` that takes a list of words and creates a columnar output of each word and their length. If given no words as positional, command-line arguments, print a usage statement. For the output, first print a header of "word" and "len", then lines which are the width of the longest word and the longest numbers with a minimum for each of the column headers themselves. The words should be left-justified in the first column and the numbers should be right-justified in the second column.

```
$ ./column.py
Usage: column.py WORD [WORD...]
$ ./column.py a an the
word  len
----  ---
a       1
an      2
the     3
$ ./column.py `cat out/1.in`
word              len
----------------  ---
Iphis               5
cyclone             7
dare                4
umbraculiferous    15
indescribableness  17
prattling           9
pediculine         10
pondwort            8
lava                4
adipoma             7
```

Exercise: rot13.py

Write a Python program called `rot13.py` that will encrypt/decrypt input text by shifting the text by a given `-s|--shift` argument or will move each character halfway through the alphabet, e.g., "a" becomes "n," "b" becomes "o," etc. The text to rotate should be provided as a single positional argument to your program and can either be a text file, text on the command line, or `-` to indicate STDIN so that you can round-trip data through your program to ensure you are encrypting and decrypting properly.

```
$ ./rot13.py
usage: rot13.py [-h] [-s int] str
rot13.py: error: the following arguments are required: str
$ ./rot13.py -h
```

```
usage: rot13.py [-h] [-s int] str

Argparse Python script

positional arguments:
  str                   Input text, file, or "-" for STDIN

optional arguments:
  -h, --help            show this help message and exit
  -s int, --shift int   Shift arg (default: 0)
$ ./rot13.py AbCd
NoPq
$ ./rot13.py AbCd -s 2
CdEf
$ ./rot13.py fox.txt
Gur dhvpx oebja sbk whzcf bire gur ynml qbt.
$ ./rot13.py fox.txt | ./rot13.py -
The quick brown fox jumps over the lazy dog.
$ ./rot13.py -s 3 fox.txt | ./rot13.py -s -3 -
The quick brown fox jumps over the lazy dog.
```

Exercise: transpose.py

Write a Python program called `transpose.py` that will read a file in ABC nota-
tion (https://en.wikipedia.org/wiki/ABC_notation) and transpose the melody
line up or down by a given `-s|--shift` argument. Like the `rot13` exercise,
it might be helpful to think of the space of notes (`ABCDEFG`) as a list which
you can roll through. For instance, if you have the note `c` and want to trans-
pose up a (minor) third (`-s 3`), you would make the new note `e`; similarly if
you have the note `F` and you go up a (major) third, you get `A`. You will not
need to worry about the actual number of semitones that you are being asked
to shift, as the previous example showed that we might be shifting by a ma-
jor/minor/augmented/diminished/pure interval. The purpose of the exercise is
simply to practice with lists.

```
$ ./transpose.py
usage: transpose.py [-h] [-s int] FILE
transpose.py: error: the following arguments are required: FILE
$ ./transpose.py -h
usage: transpose.py [-h] [-s int] FILE

Tranpose ABC notation

positional arguments:
  FILE                  Input file
```

```
optional arguments:
  -h, --help           show this help message and exit
  -s int, --shift int  Interval to shift (default: 2)
$ ./transpose.py foo
"foo" is not a file
$ ./transpose.py songs/legacy.abc -s 1
--shift "1" must be between 2 and 8
$ ./transpose.py songs/legacy.abc
<score lang="ABC">
X:1
T:The Legacy Jig
M:6/8
L:1/8
R:jig
K:A
AGA CBC | aga abc | AGA CBC | e2B BGE |
AGA CBC | aga abc | baf feC |1 eCB BGE :|2 eCB BCe |:
fgf feC | eCB BCe | fgf feC | aeC BCe |
fgf feC | e2e efg | agf feC |1 eCB BCe :|2 eCB BGE |]
</score>
```

Exercise: tictactoe.py

Write a Python program named `tictactoe.py` that accepts four named arguments:

- `-s|--state`: The state of the board (type `str`, default "........." [9 dots])
- `-p|--player`: The player to modify the state (type `str`, valid "X" or "O", no default)
- `-c|--cell`: The cell to alter (type `int`, valid 1-9, default None)
- `-h|--help`: Indication to print "usage" and exit (no error)

Your program must do the following:

- Print a usage with the `-h` or `--help` argument and exit
- Ensure the `--state` of the board is exactly 9 characters (for the 9 cells) and comprised only of ".", "X", and "O"
- Print a board with the given state where each character of the `--state` is the contents of each cell; when the state for a cell is ".", print the number of the cell
- If provided only `--player` or `--cell`, print the message "Must provide both –player and –cell" and exit *with an error code*
- If provided `--player`, ensure it is only an "X" or "O"; if not, print the expected message and exit *with an error code*

- If provided `--cell`, ensure it is in 1-9; if not, print the expected message and exit *with an error code*
- If provided both `--player` and `--cell`, modify the `--state` to change the given "cell" to the value of "player"; e.g., if `--state` is ".........." `--cell` is 1, and `--player` is "X", the state should change to "X......." and then print the board with the new state

```
$ ./tictactoe.py -h
usage: tictactoe.py [-h] [-s str] [-p str] [-c str]

Tic-Tac-Toe board

optional arguments:
  -h, --help           show this help message and exit
  -s str, --state str   Board state (default: ........)
  -p str, --player str  Player (default: )
  -c str, --cell str    Cell to apply -p (default: )
```

Print an "empty" (no cells selected) grid with no arguments:

```
$ ./tictactoe.py
-------------
| 1 | 2 | 3 |
-------------
| 4 | 5 | 6 |
-------------
| 7 | 8 | 9 |
-------------
```

Die if given a bad `--state` (too short, not made entirely of "-XO"). Note the feedback where the bad input from the user is repeated in the error message:

```
$ ./tictactoe.py -s abcdefghi
State "abcdefghi" must be 9 characters of only ., X, O
$ ./tictactoe.py --state XXO
State "XXO" must be 9 characters of only ., X, O
```

Die on invalid player:

```
$ ./tictactoe.py -p A
Invalid player "A", must be X or O
```

Die on invalid cell:

```
$ ./tictactoe.py -c 10
Invalid cell "10", must be 1-9
$ ./tictactoe.py -c 0
Invalid cell "0", must be 1-9
```

Die on a cell argument that is not an `int` (note that `argparse` will create this if you indicate the "type" as `int` for the argument):

```
$ ./tictactoe.py -c foo
usage: tictactoe.py [-h] [-s str] [-p str] [-c int]
tictactoe.py: error: argument -c/--cell: invalid int value: 'foo'
```

Print a valid state:

```
$ ./tictactoe.py -s XXO..X.OO
-------------
| X | X | O |
-------------
| 4 | 5 | X |
-------------
| 7 | O | O |
-------------
```

Mutate an initial state as described by --cell and --player. For instance, starting from the default state, change cell 2 to "X":

```
$ ./tictactoe.py -p X -c 2
-------------
| 1 | X | 3 |
-------------
| 4 | 5 | 6 |
-------------
| 7 | 8 | 9 |
-------------
```

From a state of "XXO..X.OO", change cell 7 to an "O":

```
$ ./tictactoe.py -s XXO..X.OO -p O -c 7
-------------
| X | X | O |
-------------
| 4 | 5 | X |
-------------
| O | O | O |
-------------
```

Die if the --cell is already taken:

```
$ ./tictactoe.py -s XXO..X.OO -p O -c 1
Cell 1 already taken
```

Chapter 6

Python Dictionaries

> Sometimes I feel like my job is deeply meaningful and then I remember
> that at the end of the day most of what I do is asking students to
> read error messages from compilers. – Kristopher Micinski

Python has a data type called a "dictionary" that allows you to associate some
"key" (often a string but it could be a number or even a tuple) to some "value"
(which can be anything such as a string, number, tuple, list, set, or another
dictionary). The same data structure in other languages is also called a map,
hash, and associative array.

You can define the define a dictionary with all the key/value pairs using the `{}`
braces:

```
>>> patch = {'species': 'dog', 'age': 4}
>>> patch
{'species': 'dog', 'age': 4}
```

Or you can use the `dict` function and "keyword" arguments (which, in Pythonic
style, do not use spaces around the `=` but the whitespace is not actually signifi-
cant!):

```
>>> patch = dict(species='dog', age=4)
>>> patch
{'species': 'dog', 'age': 4}
```

You might be tempted to use the `{}` curly brackets to access the keys (e.g., if you
were coming from Perl or you thought the language might be somehow internally
consistent), but Python uses the `[]` square brackets to access dictionary fields
just like arrays:

```
>>> patch['species']
'dog'
```

Since a dictionary key may be an integer, it can lead to dictionaries looking like
arrays:

```
>>> patch[0] = 'food'
>>> patch[0]
'food'
```

Note that the data types of keys of the dictionary, like lists, may be heterogenous:

```
>>> patch
{'species': 'dog', 'age': 4, 0: 'food'}
>>> list(map(type, patch.keys()))
[<class 'str'>, <class 'str'>, <class 'int'>]
```

As may be the values:

```
>>> type(patch['species'])
<class 'str'>
>>> patch['age']
4
>>> type(patch['age'])
<class 'int'>
>>> patch['likes'] = ['walking', 'running', 'car trips', 'treats', 'pets']
>>> patch
{'species': 'dog', 'age': 4, 0: 'food', 'likes': ['walking', 'running', 'car trips', 'treats
<class 'list'>
>>> list(map(type, patch.values()))
[<class 'str'>, <class 'int'>, <class 'str'>, <class 'list'>]
```

You can directly use the dictionary values like the data types they are. Here we
join the list that is in the likes slot:

```
>>> 'Patch is {} and likes {}'.format(patch['age'], ', '.join(patch['likes']))
'Patch is 4 and likes walking, running, car trips, treats, pets'
```

If you want to know if a key exists, in just as we did for list membership:

```
>>> 'likes' in patch
True
>>> 'dislikes' in patch
False
```

Just as you should not request a list position that does not exist in the list, you
should not ask for a key that does not exist in a dictionary or you program will
asplode at runtime:

```
>>> patch['dislikes']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'dislikes'
```

Better to check first:

```
>>> if 'dislikes' in patch:
...     print(patch['dislikes'])
... else:
...     print('Patch likes everything!')
...
Patch likes everything!
```

Or use the get method of the dictionary:

```
>>> patch.get('dislikes')
```

Wait, what did we get?

```
>>> type(patch.get('dislikes'))
<class 'NoneType'>
```

To find all the methods you can call on a dictionary, in the REPL type:

```
>>> help(dict)
```

Type q to "quit" the help. Use / to initiate a search, e.g., "/pop" to see how you can pop similar to the method in the list class.

Bridge of Death

Let's write a script to play with a dictionary:

```
$ cat -n bridge_of_death.py
     1  #!/usr/bin/env python3
     2
     3  person = {}
     4  print(person)
     5
     6  print('\n'.join([
     7      'Stop!', 'Who would cross the Bridge of Death',
     8      'Must answer me these questions three,',
     9      '\'ere the other side he see.'
    10  ]))
    11
    12  for field in ['name', 'quest', 'favorite color']:
    13      person[field] = input('What is your {}? '.format(field))
    14      print(person)
    15
    16  if person['favorite color'].lower() == 'blue':
    17      print('Right, off you go.')
    18  else:
    19      print('You have been eaten by a grue.')
```

And here it is in action:

```
$ ./bridge_of_death.py
{}
Stop!
Who would cross the Bridge of Death
must answer me these questions three,
ere the other side he see.
What is your name? Sir Lancelot of Camelot
{'name': 'Sir Lancelot of Camelot'}
What is your quest? To seek the Holy Grail
{'name': 'Sir Lancelot of Camelot', 'quest': 'To seek the Holy Grail'}
```

```
What is your favorite color? Blue
{'name': 'Sir Lancelot of Camelot', 'quest': 'To seek the Holy Grail', 'favorite color': 'Bl
Right, off you go.
$ ./bridge_of_death.py
{}
Stop!
Who would cross the Bridge of Death
must answer me these questions three,
ere the other side he see.
What is your name? Sir Galahad of Camelot
{'name': 'Sir Galahad of Camelot'}
What is your quest? I seek the Holy Grail
{'name': 'Sir Galahad of Camelot', 'quest': 'I seek the Holy Grail'}
What is your favorite color? Blue. No yello--
{'name': 'Sir Galahad of Camelot', 'quest': 'I seek the Holy Grail', 'favorite color': 'Blue
You have been eaten by a grue.
```

Gashlycrumb

Dictionaries are perfect for looking up some bit of information by some value:

```
$ ./gashlycrumb.py c
C is for Clara who wasted away.
$ ./gashlycrumb.py t
T is for Titus who flew into bits.
$ cat -n gashlycrumb.py
     1    #!/usr/bin/env python3
     2    """dictionary lookup"""
     3
     4    import os
     5    import sys
     6
     7    args = sys.argv[1:]
     8
     9    if len(args) != 1:
    10        print('Usage: {} LETTER'.format(os.path.basename(sys.argv[0])))
    11        sys.exit(1)
    12
    13    letter = args[0].upper()
    14
    15    text = """
    16    A is for Amy who fell down the stairs.
    17    B is for Basil assaulted by bears.
    18    C is for Clara who wasted away.
```

```
    19     D is for Desmond thrown out of a sleigh.
    20     E is for Ernest who choked on a peach.
    21     F is for Fanny sucked dry by a leech.
    22     G is for George smothered under a rug.
    23     H is for Hector done in by a thug.
    24     I is for Ida who drowned in a lake.
    25     J is for James who took lye by mistake.
    26     K is for Kate who was struck with an axe.
    27     L is for Leo who choked on some tacks.
    28     M is for Maud who was swept out to sea.
    29     N is for Neville who died of ennui.
    30     O is for Olive run through with an awl.
    31     P is for Prue trampled flat in a brawl.
    32     Q is for Quentin who sank on a mire.
    33     R is for Rhoda consumed by a fire.
    34     S is for Susan who perished of fits.
    35     T is for Titus who flew into bits.
    36     U is for Una who slipped down a drain.
    37     V is for Victor squashed under a train.
    38     W is for Winnie embedded in ice.
    39     X is for Xerxes devoured by mice.
    40     Y is for Yorick whose head was bashed in.
    41     Z is for Zillah who drank too much gin.
    42     """
    43
    44     lookup = {}
    45     for line in text.splitlines():
    46         if line:
    47             lookup[line[0]] = line
    48
    49     if letter in lookup:
    50         print(lookup[letter])
    51     else:
    52         print('I do not know "{}"'.format(letter))
$ ./gashlycrumb.py
Usage: gashlycrumb.py LETTER
$ ./gashlycrumb.py a
A is for Amy who fell down the stairs.
$ ./gashlycrumb.py b
B is for Basil assaulted by bears.
$ ./gashlycrumb.py 8
I do not know "8"
```

On line 47, we create the `lookup` using the first character of the line (`line[0]`). On line 49, we look to see if we have that letter in the `lookup`, printing the line of text if we do or complaining if we don't.

If we return to our previous chapter's DNA base counter, we can use dictionaries for this:

```
$ cat -n dna3.py
     1  #!/usr/bin/env python3
     2  """Tetra-nucleotide counter"""
     3
     4  import sys
     5  import os
     6
     7  args = sys.argv[1:]
     8
     9  if len(args) != 1:
    10      print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  dna = args[0]
    14
    15  count = {}
    16
    17  for base in dna.lower():
    18      if not base in count:
    19          count[base] = 0
    20
    21      count[base] += 1
    22
    23  counts = []
    24  for base in "acgt":
    25      num = count[base] if base in count else 0
    26      counts.append(str(num))
    27
    28  print(' '.join(counts))
$ cat dna.txt
AGCTTTTCATTCTGACTGCAACGGGCAATATGTCTCTGTGTGGATTAAAAAAAGAGTGTCTGATAGCAGC
$ ./dna3.py `cat dna.txt`
20 12 17 21
```

But why? Well, this has the great advantage of not having to declare four variables to count the four bases. True, we're only checking (in line 24) for those four, but we can now count all the letters in any string.

Notice that we create a new dict on line 15 with empty curlies {}. In line 18, we have to check if the base exists in the dict; if it doesn't, we initialize it to 0, and then we increment it by one. In line 25, we have to be careful when asking for a key that doesn't exist:

```
>>> patch['dislikes']
Traceback (most recent call last):
```

```
  File "<stdin>", line 1, in <module>
KeyError: 'dislikes'
```

If we were counting a string of DNA like "AAAAAA," then there would be no C, G or T to report, so we have to use an **if/then** expression:

```
>>> seq = 'AAAAAA'
>>> counts = {}
>>> for base in seq:
...    if not base in counts:
...       counts[base] = 0
...    counts[base] += 1
...
>>> counts
{'A': 6}
>>> counts['G']
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 'G'
>>> g = counts['G'] if 'G' in counts else 0
```

Or we can use the **get** method of a dictionary to safely get a value by a key even if the key doesn't exist:

```
>>> counts.get('G')
>>> type(counts.get('G'))
<class 'NoneType'>
```

If you look at "dna4.py," you'll see it's exactly the same as "dna3.py" with this exception:

```
    23  counts = []
    24  for base in "acgt":
    25      num = count.get(base, 0)
    26      counts.append(str(num))
```

The **get** method will not blow up your program, and it accepts an optional second argument for the default value when nothing is present:

```
>>> cat.get('likes')
>>> type(cat.get('likes'))
<class 'NoneType'>
>>> cat.get('likes', 'Cats like nothing')
'Cats like nothing'
```

Sidebar: Truthiness

Note that you might be tempted to write:

```
>>> cat.get('likes') or 'Cats like nothing'
'Cats like nothing'
```

Which appears to do the same thing, but compare with this:

```
>>> d = {'x': 0, 'y': '', 'z': None}
>>> for k in sorted(d.keys()):
...    print('{} = "{}"'.format(k, d.get(k) or 'NA'))
...
x = "NA"
y = "NA"
z = "NA"
>>> for k in sorted(d.keys()):
...    print('{} = "{}"'.format(k, d.get(k, 'NA')))
...
x = "0"
y = ""
z = "None"
```

This is a minor but potentially pernicious error due to Python's idea of Truthiness (tm):

```
>>> 1 == True
True
>>> 0 == False
True
```

The integer 1 is not actually the same thing as the boolean value True, but Python will treat it as such. Vice verse for 0 and False. The only true way to get around this is to explicitly check for None:

```
>>> for k in sorted(d.keys()):
...    val = d.get(k)
...    print('{} = "{}"'.format(k, 'NA' if val is None else val))
...
x = "0"
y = ""
z = "NA"
```

To get around the check, we could initialize the dict:

```
$ cat -n dna5.py
     1     #!/usr/bin/env python3
     2     """Tetra-nucleotide counter"""
     3
     4     import sys
     5     import os
     6
     7     args = sys.argv[1:]
     8
```

117

```
     9    if len(args) != 1:
    10        print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    11        sys.exit(1)
    12
    13    dna = args[0]
    14
    15    count = {'a': 0, 'c': 0, 'g': 0, 't': 0}
    16
    17    for base in dna.lower():
    18        if base in count:
    19            count[base] += 1
    20
    21    counts = []
    22    for base in "acgt":
    23        counts.append(str(count[base]))
    24
    25    print(' '.join(counts))
```

Back To Our Program

Now when we check on line 18, we're only going to count bases that we initialized;
further, we can then just use the keys method to get the bases:

```
$ cat -n dna5.py
     1    #!/usr/bin/env python3
     2    """Tetra-nucleotide counter"""
     3
     4    import sys
     5    import os
     6
     7    args = sys.argv[1:]
     8
     9    if len(args) != 1:
    10        print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    11        sys.exit(1)
    12
    13    dna = args[0]
    14
    15    count = {'a': 0, 'c': 0, 'g': 0, 't': 0}
    16
    17    for base in dna.lower():
    18        if base in count:
    19            count[base] += 1
    20
    21    counts = []
```

```
    22      for base in sorted(count.keys()):
    23          counts.append(str(count[base]))
    24
    25      print(' '.join(counts))
```

This kind of checking and initializing is so common that there is a standard module to define a dictionary with a default value. Unsurprisingly, it is called "defaultdict":

```
$ cat -n dna6.py
     1      #!/usr/bin/env python3
     2      """Tetra-nucleotide counter"""
     3
     4      import sys
     5      import os
     6      from collections import defaultdict
     7
     8      args = sys.argv[1:]
     9
    10      if len(args) != 1:
    11          print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    12          sys.exit(1)
    13
    14      dna = args[0]
    15
    16      count = defaultdict(int)
    17
    18      for base in dna.lower():
    19          count[base] += 1
    20
    21      counts = []
    22      for base in "acgt":
    23          counts.append(str(count[base]))
    24
    25      print(' '.join(counts))
```

On line 16, we create a `defaultdict` with the `int` type (not in quotes) for which the default value will be zero:

```
>>> from collections import defaultdict
>>> counts = defaultdict(int)
>>> counts['a']
0
```

Finally, I will show you the `Counter` that will do all the base-counting for you, returning a `defaultdict`:

```
>>> from collections import Counter
>>> c = Counter('AACTAC')
```

```
>>> c['A']
3
>>> c['G']
0
```

And here is it in the script:

```
 $ cat -n dna7.py
      1    #!/usr/bin/env python3
      2    """Tetra-nucleotide counter"""
      3
      4    import sys
      5    import os
      6    from collections import Counter
      7
      8    args = sys.argv[1:]
      9
     10    if len(args) != 1:
     11        print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
     12        sys.exit(1)
     13
     14    dna = args[0]
     15
     16    count = Counter(dna.lower())
     17
     18    counts = []
     19    for base in "acgt":
     20        counts.append(str(count[base]))
     21
     22    print(' '.join(counts))
```

So we can take that and create a program that counts all characters either from the command line or a file:

```
$ cat -n char_count1.py
      1    #!/usr/bin/env python3
      2    """Character counter"""
      3
      4    import sys
      5    import os
      6    from collections import Counter
      7
      8    args = sys.argv
      9
     10    if len(args) != 2:
     11        print('Usage: {} INPUT'.format(os.path.basename(args[0])))
     12        sys.exit(1)
     13
```

```
    14    arg = args[1]
    15    text = ''
    16    if os.path.isfile(arg):
    17        text = ''.join(open(arg).read().splitlines())
    18    else:
    19        text = arg
    20
    21    count = Counter(text.lower())
    22
    23    for letter, num in count.items():
    24        print('{} {:5}'.format(letter, num))
$ ./char_count1.py input.txt
a    20
g    17
c    12
t    21
```

Methods

The `keys` from a dict are in no particular order:

```
>>> c = Counter('AAACTAGGGACTGA')
>>> c
Counter({'A': 6, 'G': 4, 'C': 2, 'T': 2})
>>> c.keys()
dict_keys(['A', 'C', 'T', 'G'])
```

If you want them sorted, you must be explicit:

```
>>> sorted(c.keys())
['A', 'C', 'G', 'T']
```

Note that, unlike a list, you cannot call `sort` which makes sense as that will try to sort a list in-place:

```
>>> c.keys().sort()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'dict_keys' object has no attribute 'sort'
```

You can also just call `values` to get those:

```
>>> c.values()
dict_values([6, 2, 2, 4])
```

Often you will want to go through the `items` in a dict and do something with the key and value:

```
>>> for base, count in c.items():
...    print('{} = {}'.format(base, count))
...
A = 6
C = 2
T = 2
G = 4
```

But if you want to have the `keys` in a particular order, you can do this:

```
>>> for base in sorted(c.keys()):
...    print('{} = {}'.format(base, c[base]))
...
A = 6
C = 2
G = 4
T = 2
```

Or you can notice that `items` returns a list of tuples:

```
>>> c.items()
dict_items([('A', 6), ('C', 2), ('T', 2), ('G', 4)])
```

And you can call `sorted` on that:

```
>>> sorted(c.items())
[('A', 6), ('C', 2), ('G', 4), ('T', 2)]
```

Which means this will work:

```
>>> for base, count in sorted(c.items()):
...    print('{} = {}'.format(base, count))
...
A = 6
C = 2
G = 4
T = 2
```

Note that `sorted` will sort by the first elements of all the tuples, then by the second, and so forth:

```
>>> genes = [('Indy', 4), ('Boss', 2), ('Lush', 10), ('Boss', 4), ('Lush', 1)]
>>> sorted(genes)
[('Boss', 2), ('Boss', 4), ('Indy', 4), ('Lush', 1), ('Lush', 10)]
```

If we want to sort the bases instead by their frequency, we have to use some trickery like a list comprehension to first reverse the tuples:

```
>>> [(x[1], x[0]) for x in c.items()]
[(6, 'A'), (2, 'C'), (2, 'T'), (4, 'G')]
>>> sorted([(x[1], x[0]) for x in c.items()])
[(2, 'C'), (2, 'T'), (4, 'G'), (6, 'A')]
```

But what is particularly nifty about Counters is that they have built-in methods
to help you with such actions:

```
>>> c.most_common(2)
[('A', 6), ('G', 4)]
>>> c.most_common()
[('A', 6), ('G', 4), ('C', 2), ('T', 2)]
```

You should read the documentation to learn more ([https://docs.python.org/3/library/collections.html](https://do

Character Counter with the works

Finally, I'll show you a version of the character counter that takes some other
arguments to control how to show the results:

```
$ cat -n char_count2.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : Ken Youens-Clark <kyclark@email.arizona.edu>
     4  Date   : 2019-02-06
     5  Purpose: Character Counter
     6  """
     7
     8  import argparse
     9  import os
    10  import sys
    11  from collections import Counter
    12
    13
    14  # --------------------------------------------------
    15  def get_args():
    16      """get command-line arguments"""
    17      parser = argparse.ArgumentParser(
    18          description='Character counter',
    19          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    20
    21      parser.add_argument('input', help='Filename or string to count', type=str)
    22
    23      parser.add_argument(
    24          '-c',
    25          '--charsort',
    26          help='Sort by character',
    27          dest='charsort',
    28          action='store_true')
    29
```

```python
30      parser.add_argument(
31          '-n',
32          '--numsort',
33          help='Sort by number',
34          dest='numsort',
35          action='store_true')
36
37      parser.add_argument(
38          '-r',
39          '--reverse',
40          help='Sort in reverse order',
41          dest='reverse',
42          action='store_true')
43
44      return parser.parse_args()
45
46
47  # --------------------------------------------------
48  def warn(msg):
49      """Print a message to STDERR"""
50      print(msg, file=sys.stderr)
51
52
53  # --------------------------------------------------
54  def die(msg='Something bad happened'):
55      """warn() and exit with error"""
56      warn(msg)
57      sys.exit(1)
58
59
60  # --------------------------------------------------
61  def main():
62      """Make a jazz noise here"""
63      args = get_args()
64      input_arg = args.input
65      charsort = args.charsort
66      numsort = args.numsort
67      revsort = args.reverse
68
69      if charsort and numsort:
70          die('Please choose one of --charsort or --numsort')
71
72      if not charsort and not numsort:
73          charsort = True
74
75      text = ''
```

```
 76        if os.path.isfile(input_arg):
 77            text = ''.join(open(input_arg).read().splitlines())
 78        else:
 79            text = input_arg
 80
 81        count = Counter(text.lower())
 82
 83        if charsort:
 84            letters = sorted(count.keys())
 85            if revsort:
 86                letters.reverse()
 87
 88            for letter in letters:
 89                print('{} {:5}'.format(letter, count[letter]))
 90        else:
 91            pairs = sorted([(x[1], x[0]) for x in count.items()])
 92            if revsort:
 93                pairs.reverse()
 94
 95            for n, char in pairs:
 96                print('{} {:5}'.format(char, n))
 97
 98
 99  # --------------------------------------------------
100  if __name__ == '__main__':
101      main()
```

Acronym Finder

Similar to the `gashlycrumb.py` program that looked up a line of text for a given letter, we could randomly create meanings for a given acronym:

```
$ ./bacryonym.py NSF
NSF =
 - Nonrepresentationalism Staunchness Forever
 - Naturing Significantly Fontal
 - Nonclinical Solecistical Folkmoter
 - Nonhumanist Scaledrake Fellani
 - Naumk Sulpha Fause
$ ./bacryonym.py FBI
FBI =
 - Folksiness Boxmaker Interviewer
 - Flavorless Bumbler Incorruption
 - Flusterate Bakuninism Isopilocarpine
```

```
 - Freshen Bondsman Indigene
 - Fluotantalate Bornyl Interligamentous
```

That is just using the standard dictionary to look up words, so we could make it more interesting by using the works of Shakespeare:

```
$ ./bacryonym.py -w shakespeare.txt FBI
FBI =
 - Furthermore Burnet Instigation
 - Favor Bursting Insisting
 - Flower Beart Immanity
 - Fearfully Borne Itmy
 - Fooleries Blunts Intoxicates
```

Here is the Python for that:

```
$ cat -n bacronym.py
     1  #!/usr/bin/env python3
     2  """Make guesses about acronyms"""
     3
     4  import argparse
     5  import sys
     6  import os
     7  import random
     8  import re
     9  from collections import defaultdict
    10
    11
    12  # --------------------------------------------------
    13  def get_args():
    14      """get arguments"""
    15      parser = argparse.ArgumentParser(
    16          description='Explain acronyms',
    17          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    18
    19      parser.add_argument('acronym', help='Acronym', type=str, metavar='STR')
    20
    21      parser.add_argument(
    22          '-n',
    23          '--num',
    24          help='Maximum number of definitions',
    25          type=int,
    26          metavar='NUM',
    27          default=5)
    28      parser.add_argument(
    29          '-w',
    30          '--wordlist',
    31          help='Dictionary/word file',
```

```python
32            type=str,
33            metavar='STR',
34            default='/usr/share/dict/words')
35      parser.add_argument(
36            '-x',
37            '--exclude',
38            help='List of words to exclude',
39            type=str,
40            metavar='STR',
41            default='a,an,the')
42      return parser.parse_args()
43
44
45  # --------------------------------------------------
46  def main():
47      """main"""
48      args = get_args()
49      acronym = args.acronym
50      wordlist = args.wordlist
51      limit = args.num
52      goodword = r'^[a-z]{2,}$'
53      badwords = set(re.split(r'\s*,\s*', args.exclude.lower()))
54
55      if not re.match(goodword, acronym.lower()):
56          print('"{}" must be >1 in length, only use letters'.format(acronym))
57          sys.exit(1)
58
59      if not os.path.isfile(wordlist):
60          print('"{}" is not a file.'.format(wordlist))
61          sys.exit(1)
62
63      seen = set()
64      words_by_letter = defaultdict(list)
65      for word in open(wordlist).read().lower().split():
66          clean = re.sub('[^a-z]', '', word)
67          if not clean:  # nothing left?
68              continue
69
70          if re.match(goodword,
71                      clean) and clean not in seen and clean not in badwords:
72              seen.add(clean)
73              words_by_letter[clean[0]].append(clean)
74
75      len_acronym = len(acronym)
76      definitions = []
77      for i in range(0, limit):
```

127

```
78              definition = []
79              for letter in acronym.lower():
80                  possible = words_by_letter.get(letter, [])
81                  if len(possible) > 0:
82                      definition.append(
83                          random.choice(possible).title() if possible else '?')
84
85              if len(definition) == len_acronym:
86                  definitions.append(' '.join(definition))
87
88          if len(definitions) > 0:
89              print(acronym.upper() + ' =')
90              for definition in definitions:
91                  print(' - ' + definition)
92          else:
93              print('Sorry I could not find any good definitions')
94
95
96      # ------------------------------------------------
97      if __name__ == '__main__':
98          main()
```

Sequence Similarity

We can use dictionaries to count how many words are in common between any
two texts. Since I'm only trying to see if a word is present, I can use a `set` which
is like a `dict` where the values are just "1." Here is the code:

```
$ cat -n common_words.py
     1      #!/usr/bin/env python3
     2      """Count words in common between two files"""
     3
     4      import os
     5      import re
     6      import sys
     7      import string
     8
     9      # ------------------------------------------------
    10      def main():
    11          files = sys.argv[1:]
    12
    13          if len(files) != 2:
    14              msg = 'Usage: {} FILE1 FILE2'
    15              print(msg.format(os.path.basename(sys.argv[0])))
```

```
16              sys.exit(1)
17
18          for file in files:
19              if not os.path.isfile(file):
20                  print('"{}" is not a file'.format(file))
21                  sys.exit(1)
22
23          file1, file2 = files[0], files[1]
24          words1 = uniq_words(file1)
25          words2 = uniq_words(file2)
26          common = words1.intersection(words2)
27          num_common = len(common)
28          msg = 'There {} {} word{} in common between "{}" and "{}."'
29          print(msg.format('is' if num_common == 1 else 'are',
30                           num_common,
31                           '' if num_common == 1 else 's',
32                           os.path.basename(file1),
33                           os.path.basename(file2)))
34
35          for i, word in enumerate(sorted(common)):
36              print('{:3}: {}'.format(i + 1, word))
37
38      # ---------------------------------------------------
39      def uniq_words(file):
40          regex = re.compile('[' + string.punctuation + ']')
41          words = set()
42          for line in open(file):
43              for word in [regex.sub('', w) for w in line.lower().split()]:
44                  words.add(word)
45
46          return words
47
48      # ---------------------------------------------------
49      if __name__ == '__main__':
50          main()
```

Let's see it in action using a common nursery rhyme and a poem by William
Blake (1757-1827):

```
$ cat mary-had-a-little-lamb.txt
Mary had a little lamb,
It's fleece was white as snow,
And everywhere that Mary went,
The lamb was sure to go.
$ cat little-lamb.txt
Little Lamb, who made thee?
Dost thou know who made thee?
```

```
Gave thee life, & bid thee feed
By the stream & o'er the mead;
Gave thee clothing of delight,
Softest clothing, wooly, bright;
Gave thee such a tender voice,
Making all the vales rejoice?
Little Lamb, who made thee?
Dost thou know who made thee?
Little Lamb, I'll tell thee,
Little Lamb, I'll tell thee,
He is called by thy name,
For he calls himself a Lamb.
He is meek, & he is mild;
He became a little child.
I a child, & thou a lamb,
We are called by his name.
Little Lamb, God bless thee!
Little Lamb, God bless thee!
$ ./common_words.py mary-had-a-little-lamb.txt little-lamb.txt
There are 4 words in common between "mary-had-a-little-lamb.txt" and "little-lamb.txt."
  1: a
  2: lamb
  3: little
  4: the
```

Well, that's pretty uninformative. Sure "a" and "the" are shared, but we don't much care about those. And while "little" and "lamb" are present, it hardly tells us about how prevalent they are. In the nursery rhyme, they occur a total of 3 times, but they make up a significant portion of the Blake poem. Let's try to work in word frequency:

```
$ cat -n common_words2.py
     1  #!/usr/bin/env python3
     2  """Count words/frequencies in two files"""
     3
     4  import os
     5  import re
     6  import sys
     7  import string
     8  from collections import defaultdict
     9
    10  # -------------------------------------------------
    11  def word_counts(file):
    12      """Return a dictionary of words/counts"""
    13      words = defaultdict(int)
    14      regex = re.compile('[' + string.punctuation + ']')
    15      for line in open(file):
```

```
16              for word in [regex.sub('', w) for w in line.lower().split()]:
17                  words[word] += 1
18
19          return words
20
21      # -------------------------------------------------
22      def main():
23          """Start here"""
24          args = sys.argv[1:]
25
26          if len(args) != 2:
27              msg = 'Usage: {} FILE1 FILE2'
28              print(msg.format(os.path.basename(sys.argv[0])))
29              sys.exit(1)
30
31          for file in args[0:2]:
32              if not os.path.isfile(file):
33                  print('"{}" is not a file'.format(file))
34                  sys.exit(1)
35
36          file1 = args[0]
37          file2 = args[1]
38          words1 = word_counts(file1)
39          words2 = word_counts(file2)
40          common = set(words1.keys()).intersection(set(words2.keys()))
41          num_common = len(common)
42          verb = 'is' if num_common == 1 else 'are'
43          plural = '' if num_common == 1 else 's'
44          msg = 'There {} {} word{} in common between "{}" ({}) and "{}" ({}).'
45          tot1 = sum(words1.values())
46          tot2 = sum(words2.values())
47          print(msg.format(verb, num_common, plural, file1, tot1, file2, tot2))
48
49          if num_common > 0:
50              fmt = '{:>3} {:20} {:>5} {:>5}'
51              print(fmt.format('#', 'word', '1', '2'))
52              print('-' * 36)
53              shared1, shared2 = 0, 0
54              for i, word in enumerate(sorted(common)):
55                  c1 = words1[word]
56                  c2 = words2[word]
57                  shared1 += c1
58                  shared2 += c2
59                  print(fmt.format(i + 1, word, c1, c2))
60
61              print(fmt.format('', '-----', '--', '--'))
```

```
62               print(fmt.format('', 'total', shared1, shared2))
63               print(fmt.format('', 'pct',
64                            int(shared1/tot1 * 100), int(shared2/tot2 * 100)))
65
66       # -------------------------------------------------
67       if __name__ == '__main__':
68           main()
```

And here it is in action:

```
$ ./common_words2.py mary-had-a-little-lamb.txt little-lamb.txt
There are 4 words in common between "mary-had-a-little-lamb.txt" (22) and "little-lamb.txt"
  # word                    1     2
----------------------------------
  1 a                       1     5
  2 lamb                    2     8
  3 little                  1     7
  4 the                     1     3
    -----                  --    --
    total                   5    23
    pct                    22    20
```

It is interesting (to me, at least) that the shared content actually works out to about the same proportion no matter the direction. Imagine comparing a large genome to a smaller one – what is a significant portion of shared sequence space from the smaller genome might be only a small fraction of the larger one. Here we see that just those few words make up an equivalent proportion of both texts because of how repeated the words are in the Blake poem.

This is all pretty good as long as the words are spelled the same, but take the two texts here that show variations between British and American English:

```
$ cat british.txt
I went to the theatre last night with my neighbour and had a litre of
beer, the colour and flavour of which put us into such a good humour
that we forgot our labours.  We set about to analyse our behaviour,
organise our thoughts, recognise our faults, catalogue our merits, and
generally have a dialogue without pretence as a licence to improve
ourselves.
$ cat american.txt
I went to the theater last night with my neighbor and had a liter of
beer, the color and flavor of which put us into such a good humor that
we forgot our labors.  We set about to analyze our behavior, organize
our thoughts, recognize our faults, catalog our merits, and generally
have a dialog without pretense as a license to improve ourselves.
$ ./common_words2.py british.txt american.txt
There are 34 words in common between "british.txt" (63) and "american.txt" (63).
  # word                    1     2
```

```
----------------------------------
 1 a                         4     4
 2 about                     1     1
 3 and                       3     3
 4 as                        1     1
 5 beer                      1     1
 6 faults                    1     1
 7 forgot                    1     1
 8 generally                 1     1
 9 good                      1     1
10 had                       1     1
11 have                      1     1
12 i                         1     1
13 improve                   1     1
14 into                      1     1
15 last                      1     1
16 merits                    1     1
17 my                        1     1
18 night                     1     1
19 of                        2     2
20 our                       5     5
21 ourselves                 1     1
22 put                       1     1
23 set                       1     1
24 such                      1     1
25 that                      1     1
26 the                       2     2
27 thoughts                  1     1
28 to                        3     3
29 us                        1     1
30 we                        2     2
31 went                      1     1
32 which                     1     1
33 with                      1     1
34 without                   1     1
   -----                    --    --
   total                    48    48
   pct                      76    76
```

Obviously we will miss all those words because the are not spelled exactly the
same. Neither are genomes. So we need a way to decide if two words or sequences
are similar enough. One way is through sequence alignment:

```
l a b o u r      c a t a l o g u e      p r e t e n c e      l i t r e
| | | |   |      | | | | | | |          | | | | | |   |      | | |
l a b o   r      c a t a l o g          p r e t e n s e      l i t e r
```

Try writing a sequence alignment program (no, really!), and you'll find it's really quite difficult. Decades of research have gone into Smith-Waterman and BLAST and BLAT and LAST and more. Alignment works very well, but it's computationally expensive. We need a faster approximation of similarity. Enter k-mers!

A k-mer is a `k` length of "mers" or contiguous sequence (think "polymers"). Here are the 3/4-mers in my last name:

```
$ ./kmer_tiler.py youens
There are 4 3-mers in "youens."
youens
you
 oue
  uen
   ens
$ ./kmer_tiler.py youens 4
There are 3 4-mers in "youens."
youens
youe
 ouen
  uens
```

If instead looking for shared "words" we search for k-mers, we will find very different results, and the length of the k-mer matters. For instance, the first 3-mer in my name, "you" can be found 81 times in my local dictionary, but the 4-mer "youe" not at all. The longer the k-mer, the greater the specificity. Let's try our English variations with a k-mer counter:

```
$ ./common_kmers.py british.txt american.txt
There are 112 kmers in common between "british.txt" (127) and "american.txt" (127).
  # kmer                 1    2
-----------------------------------
  1 abo                  2    2
  2 all                  1    1
...
111 whi                  1    1
112 wit                  2    2
    -----               --   --
    total               142  133
    pct                 86   86
```

Our word counting program thought these two texts only 76% similar, but our kmer counter thinks they are 86% similar.

Chapter 7

Parsing with Python

> Programming is the art of doing one thing at a time. – Michael Feathers

We'll use the term "parsing" to mean deriving meaning from structured text. In this chapter, we'll look at parsing command-line arguments and common file file formats in bioinformatics like CSV, FASTA/Q, and GFF.

Command-line Arguments

If you've been using `new_py.py -a` to create new programs, you've already been using a parser – one that uses the `argparse` module to derive meaning from command-line arguments that may or may not have flags or be defined by positions. Let's create a new program and see how it works:

```
$ new_py.py -a test
Done, see new script "test.py."
```

If you check out the new script, it has a `get_args` function that will show you how to create named arguments for strings, integers, booleans, and positional arguments:

```
 1  #!/usr/bin/env python3
 2  """
 3  Author : kyclark
 4  Date   : 2019-02-19
 5  Purpose: Rock the Casbah
 6  """
 7
 8  import argparse
 9  import sys
10
11
12  # --------------------------------------------------
13  def get_args():
14      """get command-line arguments"""
15      parser = argparse.ArgumentParser(
16          description='Argparse Python script',
17          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
18
19      parser.add_argument(
20          'positional', metavar='str', help='A positional argument')
21
```

```
22      parser.add_argument(
23          '-a',
24          '--arg',
25          help='A named string argument',
26          metavar='str',
27          type=str,
28          default='')
29
30      parser.add_argument(
31          '-i',
32          '--int',
33          help='A named integer argument',
34          metavar='int',
35          type=int,
36          default=0)
37
38      parser.add_argument(
39          '-f', '--flag', help='A boolean flag', action='store_true')
40
41      return parser.parse_args()
42
43
44  # --------------------------------------------------
45  def warn(msg):
46      """Print a message to STDERR"""
47      print(msg, file=sys.stderr)
48
49
50  # --------------------------------------------------
51  def die(msg='Something bad happened'):
52      """warn() and exit with error"""
53      warn(msg)
54      sys.exit(1)
55
56
57  # --------------------------------------------------
58  def main():
59      """Make a jazz noise here"""
60      args = get_args()
61      str_arg = args.arg
62      int_arg = args.int
63      flag_arg = args.flag
64      pos_arg = args.positional
65
66      print('str_arg = "{}"'.format(str_arg))
67      print('int_arg = "{}"'.format(int_arg))
```

```
68        print('flag_arg = "{}"'.format(flag_arg))
69        print('positional = "{}"'.format(pos_arg))
70
71
72  # --------------------------------------------------
73  if __name__ == '__main__':
74        main()
```

If you run without any arguments or with `-h|--help`, you get a usage statement:

```
$ ./test.py
usage: test.py [-h] [-a str] [-i int] [-f] str
test.py: error: the following arguments are required: str
[cholla@~/work/biosys-analytics/lectures/09-python-parsing]$ ./test.py -h
usage: test.py [-h] [-a str] [-i int] [-f] str

Argparse Python script

positional arguments:
  str                 A positional argument

optional arguments:
  -h, --help          show this help message and exit
  -a str, --arg str   A named string argument (default: )
  -i int, --int int   A named integer argument (default: 0)
  -f, --flag          A boolean flag (default: False)
```

And the `argparse` module is able to turn the command line arguments into useful information:

```
$ ./test.py -a foo -i 42 -f ABCDE
str_arg = "foo"
int_arg = "42"
flag_arg = "True"
positional = "ABCDE"
```

If you try to write the code to parse `-a foo -i 42 -f ABCDE`, you will quickly appreciate how much effort using this module will save you!

CSV Files

"CSV" stands for "comma-separated values" and describes structured text that looks like:

```
foo,bar,baz
flip,burp,quux
```

More generally, these are values that are separated by some marker. Commas are typical but can cause problems when a comma can be a legitimate value, e.g., in addresses or formatted numbers, so tabs are often used as delimiters. Tab-delimited files may have the extension ".tsv," ".dat," ".tab", or ".txt." Usually CSV files have ".csv" and are especially common in the R/Pandas world.

Delimited text files are a standard way to distribute non/semi-hierarchical data – e.g., records that can be represented each on one line. (When you get into data that have relationships, e.g., parents/children, then structures like XML and JSON are more appropriate, which is not to say that people haven't sorely abused this venerable format, e.g., GFF3.) Let's first take a look at the `csv` module in Python to parse the output from Centrifuge (http://www.ccb.jhu.edu/software/centrifuge/). Despite the name, this module parses any line-oriented, delimited text, not just CSV files.

For this, we'll use some data from a study from Yellowstone National Park (https://www.imicrobe.us/#/samples/1378). For each input file, Centrifuge creates two tab-delimited output files:

1. a file ("YELLOWSTONE_SMPL_20723.sum") showing the taxonomy ID for each read it was able to classify and
2. a file ("YELLOWSTONE_SMPL_20723.tsv") of the complete taxonomy information for each taxonomy ID.

One record from the first looks like this:

```
readID        : Yellowstone_READ_00007510
seqID         : cid|321327
taxID         : 321327
score         : 640000
2ndBestScore  : 0
hitLength     : 815
queryLength   : 839
numMatches    : 1
```

One from the second looks like this:

```
name          : synthetic construct
taxID         : 32630
taxRank       : species
genomeSize    : 26537524
numReads      : 19
numUniqueReads : 19
abundance     : 0.0
```

Let's write a program that shows a table of the number of records for each "taxID":

```
$ cat -n read_count_by_taxid.py
     1    #!/usr/bin/env python3
```

```
 2    """Counts by taxID"""
 3
 4    import csv
 5    import os
 6    import sys
 7    from collections import defaultdict
 8
 9    args = sys.argv[1:]
10
11    if len(args) != 1:
12        print('Usage: {} SAMPLE.SUM'.format(os.path.basename(sys.argv[0])))
13        sys.exit(1)
14
15    sum_file = args[0]
16
17    _, ext = os.path.splitext(sum_file)
18    if not ext == '.sum':
19        print('File extention "{}" is not ".sum"'.format(ext))
20        sys.exit(1)
21
22    counts = defaultdict(int)
23    with open(sum_file) as csvfile:
24        reader = csv.DictReader(csvfile, delimiter='\t')
25        for row in reader:
26            taxID = row['taxID']
27            counts[taxID] += 1
28
29    print('\t'.join(['count', 'taxID']))
30    for taxID, count in counts.items():
31        print('\t'.join([str(count), taxID]))
```

As always, it prints a "usage" statement when run with no arguments. It also uses the `os.path.splitext` function to get the file extension and make sure that it is ".sum." Finally, if the input looks OK, then it uses the `csv.DictReader` module to parse each record of the file into a dictionary:

```
$ ./read_count_by_taxid.py
Usage: read_count_by_taxid.py SAMPLE.SUM
$ ./read_count_by_taxid.py YELLOWSTONE_SMPL_20723.tsv
File extention ".tsv" is not ".sum"
$ ./read_count_by_taxid.py YELLOWSTONE_SMPL_20723.centrifuge.sum
count   taxID
6432    321327
80   321332
19   32630
```

That's a start, but most people would rather see the a species name rather than

the NCBI taxonomy ID, so we'll need to go look up the taxIDs in the ".tsv" file:

```
$ cat -n read_count_by_tax_name.py
     1    #!/usr/bin/env python3
     2    """Counts by tax name"""
     3
     4    import csv
     5    import os
     6    import sys
     7    from collections import defaultdict
     8
     9    args = sys.argv[1:]
    10
    11    if len(args) != 1:
    12        print('Usage: {} SAMPLE.SUM'.format(os.path.basename(sys.argv[0])))
    13        sys.exit(1)
    14
    15    sum_file = args[0]
    16
    17    basename, ext = os.path.splitext(sum_file)
    18    if not ext == '.sum':
    19        print('File extention "{}" is not ".sum"'.format(ext))
    20        sys.exit(1)
    21
    22    tsv_file = basename + '.tsv'
    23    if not os.path.isfile(tsv_file):
    24        print('Cannot find expected TSV "{}"'.format(tsv_file))
    25        sys.exit(1)
    26
    27    tax_name = {}
    28    with open(tsv_file) as csvfile:
    29        reader = csv.DictReader(csvfile, delimiter='\t')
    30        for row in reader:
    31            tax_name[row['taxID']] = row['name']
    32
    33    counts = defaultdict(int)
    34    with open(sum_file) as csvfile:
    35        reader = csv.DictReader(csvfile, delimiter='\t')
    36        for row in reader:
    37            taxID = row['taxID']
    38            counts[taxID] += 1
    39
    40    print('\t'.join(['count', 'taxID']))
    41    for taxID, count in counts.items():
    42        name = tax_name.get(taxID) or 'NA'
    43        print('\t'.join([str(count), name]))
```

```
$ ./read_count_by_tax_name.py YELLOWSTONE_SMPL_20723.sum
count    taxID
6432     Synechococcus sp. JA-3-3Ab
80       Synechococcus sp. JA-2-3B'a(2-13)
19       synthetic construct
```

tabchk.py

A huge chunk of my time is spent doing ETL operations – extract, transform,
load – meaning someone sends me data (Excel or delimited-text, JSON/XML),
and I put it into some sort of database. I usually want to inspect the data to
see what it looks like, and it's hard to see the data when it's in columnar format
like this:

```
$ head oceanic_mesopelagic_zone_biome.csv
Analysis,Pipeline version,Sample,MGnify ID,Experiment type,Assembly,ENA run,ENA WGS sequence
MGYA00005220,2.0,ERS490373,MGYS00000410,metagenomic,,ERR599044,
MGYA00005081,2.0,ERS490507,MGYS00000410,metagenomic,,ERR599005,
MGYA00005208,2.0,ERS492680,MGYS00000410,metagenomic,,ERR598999,
MGYA00005133,2.0,ERS490633,MGYS00000410,metagenomic,,ERR599154,
MGYA00005272,2.0,ERS488769,MGYS00000410,metagenomic,,ERR599062,
MGYA00005209,2.0,ERS490714,MGYS00000410,metagenomic,,ERR599124,
MGYA00005243,2.0,ERS493822,MGYS00000410,metagenomic,,ERR599051,
MGYA00005117,2.0,ERS491980,MGYS00000410,metagenomic,,ERR599132,
MGYA00005135,2.0,ERS493705,MGYS00000410,metagenomic,,ERR599152,
```

I'd rather see it formatted vertically:

```
$ tabchk.py oceanic_mesopelagic_zone_biome.csv
// ****** Record 1 ****** //
Analysis             : MGYA00005220
Pipeline version     : 2.0
Sample               : ERS490373
MGnify ID            : MGYS00000410
Experiment type      : metagenomic
Assembly             :
ENA run              : ERR599044
ENA WGS sequence set :
```

Sometimes I have many more fields and lots of missing values, so I can use the
-d flag to the program indicates to show a "dense" matrix, i.e., leave out the
empty fields:

```
$ tabchk.py -d oceanic_mesopelagic_zone_biome.csv
// ****** Record 1 ****** //
Analysis         : MGYA00005220
Pipeline version : 2.0
```

```
Sample           : ERS490373
MGnify ID        : MGYS00000410
Experiment type  : metagenomic
ENA run          : ERR599044
```

Here is the `tabchk.py` program I wrote to do that. The program is generally useful, so I added it to the main `bin` directory of the repo so that you can use that if you have already added it to your `$PATH`.

```python
 1  #!/usr/bin/env python3
 2  """
 3  Author:  Ken Youens-Clark <kyclark@email.arizona.edu>
 4  Purpose: Check the first/few records of a delimited text file
 5  """
 6
 7  import argparse
 8  import csv
 9  import os
10  import re
11  import sys
12
13
14  # --------------------------------------------------
15  def get_args():
16      """Get command-line arguments"""
17      parser = argparse.ArgumentParser(
18          description='Check a delimited text file',
19          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
20
21      parser.add_argument('file', metavar='FILE', help='Input file')
22
23      parser.add_argument(
24          '-s',
25          '--sep',
26          help='Field separator',
27          metavar='str',
28          type=str,
29          default='')
30
31      parser.add_argument(
32          '-f',
33          '--field_names',
34          help='Field names (no header)',
35          metavar='str',
36          type=str,
37          default='')
38
```

```
39      parser.add_argument(
40          '-l',
41          '--limit',
42          help='How many records to show',
43          metavar='int',
44          type=int,
45          default=1)
46
47      parser.add_argument(
48          '-d',
49          '--dense',
50          help='Not sparse (skip empty fields)',
51          action='store_true')
52
53      parser.add_argument(
54          '-n',
55          '--number',
56          help='Show field number (e.g., for awk)',
57          action='store_true')
58
59      parser.add_argument(
60          '-N',
61          '--no_headers',
62          help='No headers in first row',
63          action='store_true')
64
65      return parser.parse_args()
66
67
68  # -------------------------------------------------
69  def main():
70      """main"""
71      args = get_args()
72      file = args.file
73      limit = args.limit
74      sep = args.sep
75      dense = args.dense
76      show_numbers = args.number
77      no_headers = args.no_headers
78
79      if not os.path.isfile(file):
80          print('"{}" is not a file'.format(file))
81          sys.exit(1)
82
83      if not sep:
84          _, ext = os.path.splitext(file)
```

```
 85             if ext == '.csv':
 86                 sep = ','
 87             else:
 88                 sep = '\t'
 89
 90         with open(file) as csvfile:
 91             dict_args = {'delimiter': sep}
 92
 93             if args.field_names:
 94                 regex = re.compile(r'\s*,\s*')
 95                 names = regex.split(args.field_names)
 96                 if names:
 97                     dict_args['fieldnames'] = names
 98
 99             if args.no_headers:
100                 num_flds = len(csvfile.readline().split(sep))
101                 dict_args['fieldnames'] = list(
102                     map(lambda i: 'Field' + str(i), range(1, num_flds + 1)))
103                 csvfile.seek(0)
104
105             reader = csv.DictReader(csvfile, **dict_args)
106
107             for i, row in enumerate(reader, start=1):
108                 vals = dict(
109                     [x for x in row.items() if x[1] != '']) if dense else row
110                 flds = vals.keys()
111                 longest = max(map(len, flds))
112                 fmt = '{:' + str(longest + 1) + '}: {}'
113                 print('// ****** Record {} ****** //'.format(i))
114                 n = 0
115                 for key, val in vals.items():
116                     n += 1
117                     show = fmt.format(key, val)
118                     if show_numbers:
119                         print('{:3} {}'.format(n, show))
120                     else:
121                         print(show)
122
123                 if i == limit:
124                     break
125
126
127 # ------------------------------------------------
128 if __name__ == '__main__':
129     main()
```

144

BLAST's tab-delimited output (`-outfmt 6`) does not include headers, so I have this alias:

```
alias blast6chk='tabchk.py -f "qseqid,sseqid,pident,length,mismatch,gapopen,qstart,qend,ssta
```

tabget.py

Here's a program that extracts columns from a delimited text file using the column names instead of the number (yes, I know we could just use `awk`):

```
$ cat -n tabget.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : Ken Youens-Clark <kyclark@gmail.com>
     4  Date   : 2018-11-16
     5  Purpose: Get fields from a tab/csv file
     6  """
     7
     8  import argparse
     9  import csv
    10  import os
    11  import re
    12  import sys
    13
    14
    15  # --------------------------------------------------
    16  def get_args():
    17      """get command-line arguments"""
    18      parser = argparse.ArgumentParser(
    19          description='Argparse Python script',
    20          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    21
    22      parser.add_argument(
    23          'file', nargs='+', metavar='FILE', help='Input file(s)')
    24
    25      parser.add_argument(
    26          '-d',
    27          '--delimiter',
    28          help='Field delimiter',
    29          metavar='str',
    30          type=str,
    31          default='')
    32
    33      parser.add_argument(
    34          '-f',
    35          '--field',
```

```
36              help='Name of field(s)',
37              metavar='str',
38              type=str,
39              default='')
40
41      return parser.parse_args()
42
43
44  # --------------------------------------------------
45  def warn(msg):
46      """Print a message to STDERR"""
47      print(msg, file=sys.stderr)
48
49
50  # --------------------------------------------------
51  def die(msg='Something bad happened'):
52      """warn() and exit with error"""
53      warn(msg)
54      sys.exit(1)
55
56
57  # --------------------------------------------------
58  def main():
59      """Make a jazz noise here"""
60      args = get_args()
61      files = args.file
62      default_delim = args.delimiter
63      field_names = re.split('\s*,\s*', args.field)
64
65      for file in files:
66          with open(file, 'rt') as fh:
67              delim = default_delim
68              if not delim:
69                  _, ext = os.path.splitext(file)
70                  if ext == '.csv':
71                      delim = ','
72                  else:
73                      delim = '\t'
74
75              reader = csv.DictReader(fh, delimiter=delim)
76
77              print(delim.join(field_names))
78
79              for row in reader:
80                  flds = list(map(lambda f: row[f], field_names))
81                  print(delim.join(flds))
```

```
82
83
84  # --------------------------------------------------
85  if __name__ == '__main__':
86      main()
```

tab2json.py

At some point I must have needed to turn a flat, delimited text file into a
hierarchical, JSON structured, but I cannot at this moment remember why.
Anyway, here's a program that will do that.

```
$ cat -n tab2json.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:  Ken Youens-Clark <kyclark@gmail.com>
     4  Purpose: Convert a delimited text file to JSON
     5  """
     6
     7  import argparse
     8  import csv
     9  import json
    10  import os
    11  import re
    12  import sys
    13
    14
    15  # --------------------------------------------------
    16  def get_args():
    17      """get args"""
    18      parser = argparse.ArgumentParser(
    19          description='Argparse Python script',
    20          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    21
    22      parser.add_argument(
    23          'tabfile', metavar='str', nargs='+', help='A positional argument')
    24
    25      parser.add_argument(
    26          '-s',
    27          '--sep',
    28          help='Field separator',
    29          metavar='str',
    30          type=str,
    31          default='\t')
    32
```

147

```
33      parser.add_argument(
34          '-o',
35          '--outdir',
36          help='Output dir',
37          metavar='str',
38          type=str,
39          default='')
40
41      parser.add_argument(
42          '-i',
43          '--indent',
44          help='Indent level',
45          metavar='int',
46          type=int,
47          default=2)
48
49      parser.add_argument(
50          '-n',
51          '--normalize_headers',
52          help='Normalize headers',
53          action='store_true')
54
55      return parser.parse_args()
56
57
58  # --------------------------------------------------
59  def main():
60      """main"""
61      args = get_args()
62      indent_level = args.indent
63      out_dir = args.outdir
64      fs = args.sep
65      norm_hdr = args.normalize_headers
66      tabfiles = args.tabfile
67
68      if len(tabfiles) < 1:
69          print('No input files')
70          sys.exit(1)
71
72      if indent_level < 0:
73          indent_level = 0
74
75      if out_dir and not os.path.isdir(out_dir):
76          os.makedirs(out_dir)
77
78      for i, tabfile in enumerate(tabfiles, start=1):
```

```
 79              basename = os.path.basename(tabfile)
 80              filename, _ = os.path.splitext(basename)
 81              dirname = os.path.dirname(os.path.abspath(tabfile))
 82              print('{:3}: {}'.format(i, basename))
 83              write_dir = out_dir if out_dir else dirname
 84              out_path = os.path.join(write_dir, filename + '.json')
 85              out_fh = open(out_path, 'wt')
 86
 87              with open(tabfile) as fh:
 88                  reader = csv.DictReader(fh, delimiter=fs)
 89                  if norm_hdr:
 90                      reader.fieldnames = list(map(normalize, reader.fieldnames))
 91                  out_fh.write(json.dumps(list(reader), indent=indent_level))
 92
 93
 94  # --------------------------------------------------
 95  def normalize(hdr):
 96      return re.sub(r'[^A-Za-z0-9_]', '', hdr.lower().replace(' ', '_'))
 97
 98
 99  # --------------------------------------------------
100  if __name__ == '__main__':
101      main()
```

FASTA

Now let's finally get into parsing good, old FASTA files. We're going to need to install the BioPython (http://biopython.org/) module to get a FASTA parser. This should work for you:

```
$ python3 -m pip install biopython
```

For this exercise, I'll use a few reads from the Global Ocean Sampling Expedition (https://imicrobe.us/#/samples/578). You can download the full file with this command:

```
$ iget /iplant/home/shared/imicrobe/projects/26/samples/578/CAM_SMPL_GS108.fa
```

Since that file is 725M, I've added a sample to the repo in the examples directory.

```
$ head -5 CAM_SMPL_GS108.fa
>CAM_READ_0231669761 /library_id="CAM_LIB_GOS108XLRVAL-4F-1-400" /sample_id="CAM_SMPL_GS108'
ATTTACAATAATTTAATAAAATTAACTAGAAATAAAATATTGTATGAAAATATGTTAAAT
AATGAAAGTTTTTCAGATCGTTTAATAATATTTTTCTTCCATTTTGCTTTTTTCTAAAAT
TGTTCAAAAACAAACTTCAAAGGAAAATCTTCAAAATTTACATGATTTTATATTTAAACA
AATAGAGTTAAGTATAAGAGAAATTGGATATGGTGATGCTTCAATAAATAAAAAAATGAA
```

The format of a FASTA file is:

- A record starts with a header row which has > as the first character on a line
- The string following the > up until the first whitespace is the record ID
- Anything following the ID up to the newline can be the "description," but here we see this space has been set up as key/value pairs of metadata
- Any line after a header that does not start with > is the sequence. The sequence may be one long line or many shorter lines.

We **could** write our own FASTA parser, and we would definitely learn much along the way, but let's not and instead use the BioPython `SeqIO` (sequence input-output) module to read and write all the different formats. FASTA is one of the most common, but other formats may include FASTQ (FASTA but with "Quality" scores for the base calls), GenBank, EMBL, and more. See https://biopython.org/wiki/SeqIO for an exhaustive list.

There is a useful program called `seqmagick` that will give you information like the following:

```
$ seqmagick info *.fa
name               alignment   min_len   max_len   avg_len   num_seqs
CAM_SMPL_GS108.fa FALSE             47       594    369.65        499
CAM_SMPL_GS112.fa FALSE             50       624    383.50        500
```

You can install it like so:

```
$ python -m pip install seqmagick
```

Let's write a toy program to mimic part of the output. We'll skip the "alignment" and just do min/max/avg lengths, and the number of sequences. You can pretty much copy and paste the example code from http://biopython.org/wiki/SeqIO. Here is the output from our script, `seqmagique.py`:

```
$ ./seqmagique.py *.fa
name               min_len   max_len   avg_len    num_seqs
CAM_SMPL_GS108.fa       47       594 369.45            500
CAM_SMPL_GS112.fa       50       624 383.50            500
```

The code to produce this builds on our earlier skills of lists and dictionaries as we will parse each file and save a dictionary of stats into a list, then we will iterate over that list at the end to show the output.

```
$ cat -n seqmagique.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:  Ken Youens-Clark <kyclark@email.arizona.edu>
     4  Purpose: Mimic seqmagick, print stats on FASTA sequences
     5  """
     6
     7  import os
     8  import sys
```

```
 9  import numpy as np
10  from Bio import SeqIO
11
12  files = sys.argv[1:]
13
14  if not files:
15      print('Usage: {} F1.fa [F2.fa...]'.format(os.path.basename(sys.argv[0])))
16      sys.exit(1)
17
18  info = []
19  for file in files:
20      lengths = []
21      for record in SeqIO.parse(file, 'fasta'):
22          lengths.append(len(record.seq))
23
24      info.append({
25          'name': os.path.basename(file),
26          'min_len': min(lengths),
27          'max_len': max(lengths),
28          'avg_len': '{:.2f}'.format(np.mean(lengths)),
29          'num_seqs': len(lengths)
30      })
31
32  if info:
33      longest_file_name = max([len(f['name']) for f in info])
34      fmt = '{:' + str(longest_file_name) + '} {:10} {:10} {:10} {:10}'
35      flds = ['name', 'min_len', 'max_len', 'avg_len', 'num_seqs']
36      print(fmt.format(*flds))
37      for rec in info:
38          print(fmt.format(*[rec[fld] for fld in flds]))
39  else:
40      print('I had trouble parsing your data')
```

FASTA subset

Sometimes you may only want to use part of a FASTA file, e.g., you want the
first 1000 sequences to test some code, or you have samples that vary wildly
in size and you want to sub-sample them down to an equal number of reads.
Here is a Python program that will write the first N samples to a given output
directory:

```
$ cat -n subset_fastx.py
    1  #!/usr/bin/env python3
    2  """
    3  Author:  Ken Youens-Clark <kyclark@email.arizona.edu>
```

```
 4   Purpose: Subset FASTA/Q files
 5   """
 6
 7   import argparse
 8   import os
 9   import sys
10   from Bio import SeqIO
11
12
13   # --------------------------------------------------
14   def get_args():
15       """get args"""
16       parser = argparse.ArgumentParser(
17           description='Split FASTA files',
18           formatter_class=argparse.ArgumentDefaultsHelpFormatter)
19
20       parser.add_argument('file', help='Input file', metavar='FILE')
21
22       parser.add_argument(
23           '-f',
24           '--infmt',
25           help='Input file format',
26           type=str,
27           metavar='FMT',
28           choices=['fasta', 'fastq'],
29           default='fasta')
30
31       parser.add_argument(
32           '-F',
33           '--outfmt',
34           help='Output file format',
35           type=str,
36           metavar='FMT',
37           default=None)
38
39       parser.add_argument(
40           '-n',
41           '--num',
42           help='Number of records per file',
43           type=int,
44           metavar='NUM',
45           default=500000)
46
47       parser.add_argument(
48           '-o',
49           '--outdir',
```

```
50              help='Output directory',
51              type=str,
52              metavar='DIR',
53              default='subset')
54
55      return parser.parse_args()
56
57  # --------------------------------------------------
58  def warn(msg):
59      """Print a message to STDERR"""
60      print(msg, file=sys.stderr)
61
62
63  # --------------------------------------------------
64  def die(msg='Something bad happened'):
65      """warn() and exit with error"""
66      warn(msg)
67      sys.exit(1)
68
69
70  # --------------------------------------------------
71  def main():
72      """main"""
73      args = get_args()
74      in_file = args.file
75      in_fmt = args.infmt
76      out_fmt = args.outfmt if args.outfmt else args.infmt
77      out_dir = args.outdir
78      num_seqs = args.num
79
80      if not os.path.isfile(in_file):
81          die('--file "{}" is not a file'.format(in_file))
82
83      if os.path.dirname(os.path.abspath(in_file)) == os.path.abspath(out_dir):
84          die('--outdir "{}" cannot be the same as input files'.format(out_dir))
85
86      if num_seqs < 1:
87          die("--num cannot be less than one")
88
89      if not os.path.isdir(out_dir):
90          os.mkdir(out_dir)
91
92      basename = os.path.basename(in_file)
93      out_file = os.path.join(out_dir, basename)
94      out_fh = open(out_file, 'wt')
95      num_written = 0
```

```
 96
 97        for record in SeqIO.parse(in_file, in_fmt):
 98            SeqIO.write(record, out_fh, out_fmt)
 99            num_written += 1
100
101            if num_written == num_seqs:
102                break
103
104        print('Done, wrote {} sequence{} to "{}"'.format(
105            num_written, '' if num_written == 1 else 's', out_file))
106
107
108    # --------------------------------------------------
109    if __name__ == '__main__':
110        main()
```

Here is a version that will randomly select some percentage of the reads from the input file. I had to write this version because we had created an artificial metagenome from a set of known organisms, and I was testing a program with input of various numbers of reads. I did not realize at first that, in creating the artificial set, reads from each organism had been added in blocks. Since I was taking all my reads from the top of the file down, I was mostly getting just the first few species. Randomly selecting reads when there are potentially millions of records is a bit tricky, so I decided to use a non-deterministic approach where I just roll the dice and see if the number I get on each read is less than the percentage of reads I want to take. This program will also stop at a given number of reads so you could use it to randomly subset an unevenly sized number of samples down to the same number of reads per sample.

```
$ cat -n random_subset.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:  Ken Youens-Clark <kyclark@email.arizona.edu>
     4  Purpose: Probabalistically subset FASTQ/A
     5  """
     6
     7  import argparse
     8  import os
     9  import re
    10  import sys
    11  from random import randint
    12  from Bio import SeqIO
    13
    14
    15  # --------------------------------------------------
    16  def get_args():
    17      """get args"""
```

```
18      parser = argparse.ArgumentParser(
19          description='Randomly subset FASTQ',
20          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
21
22      parser.add_argument('file', metavar='FILE', help='FASTQ/A file')
23
24      parser.add_argument(
25          '-p',
26          '--pct',
27          help='Percent of reads',
28          metavar='int',
29          type=int,
30          default=50)
31
32      parser.add_argument(
33          '-m',
34          '--max',
35          help='Maximum number of reads',
36          metavar='int',
37          type=int,
38          default=0)
39
40      parser.add_argument(
41          '-f',
42          '--input_format',
43          help='Intput format',
44          metavar='IN_FMT',
45          type=str,
46          choices=['fastq', 'fasta'],
47          default='')
48
49      parser.add_argument(
50          '-F',
51          '--output_format',
52          help='Output format',
53          metavar='OUT_FMT',
54          type=str,
55          choices=['fastq', 'fasta'],
56          default='')
57
58      parser.add_argument(
59          '-o',
60          '--outfile',
61          help='Output file',
62          metavar='FILE',
63          type=str,
```

```
64              default='')
65
66       return parser.parse_args()
67
68
69   # --------------------------------------------------
70   def warn(msg):
71       """Print a message to STDERR"""
72       print(msg, file=sys.stderr)
73
74
75   # --------------------------------------------------
76   def die(msg='Something bad happened'):
77       """warn() and exit with error"""
78       warn(msg)
79       sys.exit(1)
80
81
82   # --------------------------------------------------
83   def main():
84       """main"""
85       args = get_args()
86       file = args.file
87       pct = args.pct
88       out_file = args.outfile
89       max_num_reads = args.max
90       min_num = 0
91       max_num = 100
92
93       if not os.path.isfile(file):
94           die('"{}" is not a file'.format(file))
95
96       in_fmt = args.input_format
97       if not in_fmt:
98           _, ext = os.path.splitext(file)
99           in_fmt = 'fastq' if re.match('\.f(ast)?q$', ext) else 'fasta'
100
101      out_fmt = args.output_format or in_fmt
102
103      if not min_num < pct < max_num:
104          msg = '--pct "{}" must be between {} and {}'
105          die(msg.format(pct, min_num, max_num))
106
107      if not out_file:
108          base, _ = os.path.splitext(file)
109          out_file = '{}.sub{}.{}'.format(base, pct, out_fmt)
```

156

```
110
111        out_fh = open(out_file, 'wt')
112        num_taken = 0
113        total_num = 0
114
115        with open(file) as fh:
116            for rec in SeqIO.parse(fh, in_fmt):
117                total_num += 1
118                if randint(min_num, max_num) <= pct:
119                    num_taken += 1
120                    SeqIO.write(rec, out_fh, out_fmt)
121                    if max_num_reads > 0 and num_taken == max_num_reads:
122                        break
123
124        out_fh.close()
125
126        print('Wrote {} of {} ({:.02f}%) to "{}"'.format(
127            num_taken, total_num, num_taken / total_num * 100, out_file))
128
129
130  # --------------------------------------------------
131  if __name__ == '__main__':
132      main()
```

FASTA splitter

I seem to have implemented my own FASTA splitter a few times in as many languages. Here is one that writes a maximum number of sequences to each output file. It would not be hard to instead write a maximum number of bytes, but, for the short reads I usually handle, this works fine. Again I will use the BioPython SeqIO module to parse the FASTA files.

```
$ cat -n fa_split.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:  Ken Youens-Clark
     4  Purpose: Split FASTA files
     5  NB:      If you have FASTQ files, maybe just use "split"?
     6  """
     7
     8  import argparse
     9  import os
    10  import sys
    11  from Bio import SeqIO
    12
```

```
13
14   # -------------------------------------------------
15   def get_args():
16       """get args"""
17       parser = argparse.ArgumentParser(
18           description='Split FASTA/Q files',
19           formatter_class=argparse.ArgumentDefaultsHelpFormatter)
20
21       parser.add_argument('file', help='FASTA input file(s)', nargs='+')
22
23       parser.add_argument(
24           '-f',
25           '--input_format',
26           help='Input file format',
27           type=str,
28           metavar='FORMAT',
29           choices=['fasta', 'fastq'],
30           default='fasta')
31
32       parser.add_argument(
33           '-F',
34           '--output_format',
35           help='Output file format',
36           type=str,
37           metavar='FORMAT',
38           choices=['fasta', 'fastq'],
39           default='fasta')
40
41       parser.add_argument(
42           '-n',
43           '--sequences_per_file',
44           help='Number of sequences per file',
45           type=int,
46           metavar='NUM',
47           default=50)
48
49       parser.add_argument(
50           '-o',
51           '--out_dir',
52           help='Output directory',
53           type=str,
54           metavar='DIR',
55           default='fasplit')
56
57       return parser.parse_args()
58
```

```
59
60   # --------------------------------------------------
61   def warn(msg):
62       """Print a message to STDERR"""
63       print(msg, file=sys.stderr)
64
65
66   # --------------------------------------------------
67   def die(msg='Something bad happened'):
68       """warn() and exit with error"""
69       warn(msg)
70       sys.exit(1)
71
72
73   # --------------------------------------------------
74   def main():
75       """main"""
76       args = get_args()
77       files = args.file
78       input_format = args.input_format
79       output_format = args.output_format
80       out_dir = args.out_dir
81       seqs_per_file = args.sequences_per_file
82
83       if not os.path.isdir(out_dir):
84           os.mkdir(out_dir)
85
86       if seqs_per_file < 1:
87           die('--sequences_per_file "{}" cannot be less than one'.format(
88               seqs_per_file))
89
90       num_files = 0
91       num_seqs_written = 0
92       for i, file in enumerate(files, start=1):
93           print('{:3d}: {}'.format(i, os.path.basename(file)))
94           num_files += 1
95           num_seqs_written += process(
96               file=file,
97               input_format=input_format,
98               output_format=output_format,
99               out_dir=out_dir,
100              seqs_per_file=seqs_per_file)
101
102      print('Done, processed {} sequence{} from {} file{} into "{}"'.format(
103          num_seqs_written, '' if num_seqs_written == 1 else 's', num_files, ''
104          if num_files == 1 else 's', out_dir))
```

```
105
106
107  # ----------------------------------------------------
108  def process(file, input_format, output_format, out_dir, seqs_per_file):
109      """
110      Spilt file into smaller files into out_dir
111      Optionally convert to output format
112      Return number of sequences written
113      """
114      if not os.path.isfile(file):
115          warn('"{}" is not valid'.format(file))
116          return 0
117
118      basename, ext = os.path.splitext(os.path.basename(file))
119      out_fh = None
120      i = 0
121      num_written = 0
122      nfile = 0
123      for record in SeqIO.parse(file, input_format):
124          if i == seqs_per_file:
125              i = 0
126              if out_fh is not None:
127                  out_fh.close()
128                  out_fh = None
129
130          i += 1
131          num_written += 1
132          if out_fh is None:
133              nfile += 1
134              path = os.path.join(out_dir,
135                                  basename + '.' + '{:04d}'.format(nfile) + ext)
136              out_fh = open(path, 'wt')
137
138          SeqIO.write(record, out_fh, output_format)
139
140      return num_written
141
142
143  # ----------------------------------------------------
144  if __name__ == '__main__':
145      main()
```

You can run this on the FASTA files in the `examples` directory to split them
into files of 50 sequences each:

```
$ ./fa_split.py *.fa
  1: CAM_SMPL_GS108.fa
```

```
  2: CAM_SMPL_GS112.fa
Done, processed 1000 sequences from 2 files into "fasplit"
$ ls -lh fasplit/
total 1088
-rw-r--r--  1 kyclark  staff     22K Feb 19 15:41 CAM_SMPL_GS108.0001.fa
-rw-r--r--  1 kyclark  staff     28K Feb 19 15:41 CAM_SMPL_GS108.0002.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS108.0003.fa
-rw-r--r--  1 kyclark  staff     23K Feb 19 15:41 CAM_SMPL_GS108.0004.fa
-rw-r--r--  1 kyclark  staff     22K Feb 19 15:41 CAM_SMPL_GS108.0005.fa
-rw-r--r--  1 kyclark  staff     26K Feb 19 15:41 CAM_SMPL_GS108.0006.fa
-rw-r--r--  1 kyclark  staff     29K Feb 19 15:41 CAM_SMPL_GS108.0007.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS108.0008.fa
-rw-r--r--  1 kyclark  staff     26K Feb 19 15:41 CAM_SMPL_GS108.0009.fa
-rw-r--r--  1 kyclark  staff     24K Feb 19 15:41 CAM_SMPL_GS108.0010.fa
-rw-r--r--  1 kyclark  staff     26K Feb 19 15:41 CAM_SMPL_GS112.0001.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS112.0002.fa
-rw-r--r--  1 kyclark  staff     28K Feb 19 15:41 CAM_SMPL_GS112.0003.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS112.0004.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS112.0005.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS112.0006.fa
-rw-r--r--  1 kyclark  staff     28K Feb 19 15:41 CAM_SMPL_GS112.0007.fa
-rw-r--r--  1 kyclark  staff     29K Feb 19 15:41 CAM_SMPL_GS112.0008.fa
-rw-r--r--  1 kyclark  staff     27K Feb 19 15:41 CAM_SMPL_GS112.0009.fa
-rw-r--r--  1 kyclark  staff     16K Feb 19 15:41 CAM_SMPL_GS112.0010.fa
```

We can verify that things worked:

```
$ for file in fasplit/*; do echo -n $file && grep '^>' $file | wc -l; done
fasplit/CAM_SMPL_GS108.0001.fa        50
fasplit/CAM_SMPL_GS108.0002.fa        50
fasplit/CAM_SMPL_GS108.0003.fa        50
fasplit/CAM_SMPL_GS108.0004.fa        50
fasplit/CAM_SMPL_GS108.0005.fa        50
fasplit/CAM_SMPL_GS108.0006.fa        50
fasplit/CAM_SMPL_GS108.0007.fa        50
fasplit/CAM_SMPL_GS108.0008.fa        50
fasplit/CAM_SMPL_GS108.0009.fa        50
fasplit/CAM_SMPL_GS108.0010.fa        50
fasplit/CAM_SMPL_GS112.0001.fa        50
fasplit/CAM_SMPL_GS112.0002.fa        50
fasplit/CAM_SMPL_GS112.0003.fa        50
fasplit/CAM_SMPL_GS112.0004.fa        50
fasplit/CAM_SMPL_GS112.0005.fa        50
fasplit/CAM_SMPL_GS112.0006.fa        50
fasplit/CAM_SMPL_GS112.0007.fa        50
fasplit/CAM_SMPL_GS112.0008.fa        50
fasplit/CAM_SMPL_GS112.0009.fa        50
```

```
fasplit/CAM_SMPL_GS112.0010.fa      50
```

FASTQ

FASTA (sequence) plus "quality" scores for each base call gives us "FASTQ."
Here is an example:

```
$ head -4 !$
head -4 input.fastq
@M00773:480:000000000-BLYPT:1:2106:12063:1841 1:N:0:AGGCGACCTTA
TTTCTGTGCCAGCAGCCGCGGTAAGACAGAGGTGGCGAGCGTTGTTCGGATTTACTGGGCGTAAAGCGCGGGTAGGCGGTTCGGCCAGTCAC
+
CCCCCGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGEFGGFEGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGGFGGG
```

Because of inherent logical flaws in this file format, the only sane representation
is for the record to consist of four lines:

1. header ('@', ID, desc, yadda yadda yadda)
2. sequence
3. spacer
4. quality scores (phred 33/64)

Here is what the record looks like:

```
>>> from Bio import SeqIO
>>> rec = list(SeqIO.parse('input.fastq', 'fastq'))[0]
>>> rec = list(SeqIO.parse('input.fastq', 'fastq'))[0]
>>> print(rec)
ID: M00773:480:000000000-BLYPT:1:2106:12063:1841
Name: M00773:480:000000000-BLYPT:1:2106:12063:1841
Description: M00773:480:000000000-BLYPT:1:2106:12063:1841 1:N:0:AGGCGACCTTA
Number of features: 0
Per letter annotation for: phred_quality
Seq('TTTCTGTGCCAGCAGCCGCGGTAAGACAGAGGTGGCGAGCGTTGTTCGGATTTA...CGC', SingleLetterAlphabet())
```

But this looks pretty much like a FASTA file, so where is the quality information?
We have to look here (http://biopython.org/DIST/docs/api/Bio.SeqIO.QualityIO-module.html):

```
>>> print(rec.format("qual"))
>M00773:480:000000000-BLYPT:1:2106:12063:1841 1:N:0:AGGCGACCTTA
34 34 34 34 34 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38
38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 36 37 38
38 37 36 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38
38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38 38
38 38 38 38 38 38 38 38 37 38 38 38 38 38 38 38 38 38 38 38
38 38 38 38 38 38 38 36 38 38 38 38 38 38 38 36 38 38 38 38
38 38 35 38 35 38 38 38 38 38 38 38 38 38 38 38 37 35 38 38
```

```
38 38 38 38 38 38 37 37 37 37 35 37 38 22 37 37 38 38 38 38
38 38 38 38 38 22 36 38 38 38 38 38 35 38 38 36 38 38 38 38
38 38 28 36 37 38 35 38 38 37 38 38 35 36 38 38 38 38 37 37
34 20 26 36 36 35 37 36 37 38 36 38 37 34 37 38 36 36 34 34
23 30 20 34 36 36 9 25 20 9 26 30 37 38 38 37 38 34 34 37
38 32 37 37 38 38 38 35 38 38 37 37 38 34 35 36 34 38 38 38
38 36 26 36 36 23 36 34 28 18 24 15 26 20 22 20 29 23 27 10
24 37 38 38 37 34 27 23 34 38 37 37 25 24 10 24 11 27 35 20
8
```

We can combine the bases and their quality scores into a list of tuples (which can naturally become a dictionary):

```
>>> list(zip(rec.seq, rec.format('qual')))
[('T', '>'), ('T', 'M'), ('T', 'O'), ('C', 'O'), ...
>>> for base, qual in zip(rec.seq, rec.format('qual')):
...    print('base = "{}" qual = "{}"'.format(base, qual))
...    break
...
base = "T" qual = ">"
```

The scores are based on the ordinal represenation of the quality characters' ASCII values. Cf:

- https://www.rapidtables.com/code/text/ascii-table.html
- https://www.drive5.com/usearch/manual/quality_score.html

We can convert FASTQ to FASTA by simply changing the leading "@" in the header to ">" and then removing lines 3 and 4 from each record. Here is an [g]awk one-liner to do that:

```
#!/bin/gawk -f

### fq2fa.awk
##
## Copyright Tomer Altman
##
### Desription:
##
## Given a FASTQ formatted file, transform it into a FASTA nucleotide file.

(FNR % 4) == 1 || (FNR % 4) == 2 { gsub("^@", ">"); print }
```

Can you write one in Python?

GFF

Two of the most common output files in bioinformatics, GFF (General Feature Format) and BLAST's tab/CSV files do not include headers, so it's up to you to merge in the headers. Additionally, some of the lines may be comments (they start with **#** just like bash and Python), so you should skip those. Further, the last field in GFF is basically a dumping ground for whatever else the data provider felt like putting there. Usually it's a bunch of "key=value" pairs, but there's no guarantee. Let's take a look at parsing the GFF output from Prodigal:

```
$ cat -n parse_prodigal_gff.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:  Ken Youens-Clark <kyclark@email.arizona.edu>
     4  Purpose: Parse the GFF output of Prodigal
     5  """
     6
     7  import argparse
     8  import os
     9  import sys
    10
    11
    12  # --------------------------------------------------
    13  def get_args():
    14      """get args"""
    15      parser = argparse.ArgumentParser(
    16          description='Prodigal GFF parser',
    17          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    18
    19      parser.add_argument('gff', metavar='FILE', help='Prodigal GFF file')
    20
    21      parser.add_argument(
    22          '-m',
    23          '--min',
    24          help='Min score',
    25          metavar='float',
    26          type=float,
    27          default=0)
    28
    29      return parser.parse_args()
    30
    31
    32  # --------------------------------------------------
    33  def warn(msg):
    34      """Print a message to STDERR"""
    35      print(msg, file=sys.stderr)
```

```python
36
37
38   # ----------------------------------------------------
39   def die(msg='Something bad happened'):
40       """warn() and exit with error"""
41       warn(msg)
42       sys.exit(1)
43
44
45   # ----------------------------------------------------
46   def main():
47       """main"""
48       args = get_args()
49       gff_file = args.gff
50       min_score = args.min
51
52       if not os.path.isfile(gff_file):
53           die('GFF "{}" is not a file'.format(gff_file))
54
55       flds = [
56           'seqname', 'source', 'feature', 'start', 'end', 'score', 'strand',
57           'frame', 'attribute'
58       ]
59
60       for line in open(gff_file):
61           if line.startswith('#'):
62               continue
63
64           vals = line.rstrip().split('\t')
65           rec = dict(zip(flds, vals))
66           attrs = {}
67
68           for x in rec['attribute'].split(';'):
69               if '=' in x:
70                   key, value = x.split('=')
71                   attrs[key] = value
72
73           score = attrs.get('score')
74           if score is not None and float(score) >= min_score:
75               print('{} {}'.format(rec['seqname'], score))
76
77
78   # ----------------------------------------------------
79   if __name__ == '__main__':
80       main()
```

XML

Here's an example that looks at XML from the NCBI taxonomy. Here is what the raw file looks like:

```
$ head ena-101.xml
<?xml version="1.0" encoding="UTF-8"?>
<SAMPLE alias="SAMD00024455" accession="DRS018892" broker_name="DDBJ">
    <IDENTIFIERS>
        <PRIMARY_ID>DRS018892</PRIMARY_ID>
        <EXTERNAL_ID namespace="BioSample">SAMD00024455</EXTERNAL_ID>
        <SUBMITTER_ID namespace="">SAMD00024455</SUBMITTER_ID>
    </IDENTIFIERS>
    <TITLE>Surface water bacterial community from the East China Sea Site 100</TITLE>
    <SAMPLE_NAME>
        <TAXON_ID>408172</TAXON_ID>
```

The whitespace in XML is not significant and simply bloats the size of the file, so often you will get something that is unreadable. I recommend you install the program `xmllint` to look at such files. If you inspect the file, you can see that XML gives us a way to represent hierarchical data unlike CSV files which are essentially "flat" (unless you start sticking things like lists and key/value pairs [dictionaries]). We need to use a specific XML parser and use accessors that look quite a bit like file paths. There is a "root" of the XML from which we can descend into the structure to find data. Here is a program that will extract various parts of the XML.

```
$ cat -n xml_ena.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date   : 2019-02-22
     5  Purpose: Rock the Casbah
     6  """
     7
     8  import argparse
     9  import os
    10  import sys
    11  from xml.etree.ElementTree import ElementTree
    12
    13
    14  # --------------------------------------------------
    15  def get_args():
    16      """get command-line arguments"""
    17      parser = argparse.ArgumentParser(
    18          description='Argparse Python script',
    19          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
```

```python
20
21      parser.add_argument('xml', metavar='XML', help='XML input', nargs='+')
22
23      parser.add_argument(
24          '-o',
25          '--outdir',
26          help='Output directory',
27          metavar='str',
28          type=str,
29          default='out')
30
31      return parser.parse_args()
32
33
34  # --------------------------------------------------
35  def warn(msg):
36      """Print a message to STDERR"""
37      print(msg, file=sys.stderr)
38
39
40  # --------------------------------------------------
41  def die(msg='Something bad happened'):
42      """warn() and exit with error"""
43      warn(msg)
44      sys.exit(1)
45
46
47  # --------------------------------------------------
48  def main():
49      """Make a jazz noise here"""
50      args = get_args()
51      xml_files = args.xml
52      out_dir = args.outdir
53
54      if not os.path.isdir(out_dir):
55          os.makedirs(out_dir)
56
57      for file in xml_files:
58          print('>>>>>>', file)
59          tree = ElementTree()
60          root = tree.parse(file)
61
62          d = []
63          for key, value in root.attrib.items():
64              d.append(('sample.' + key, value))
65
```

```
66              for id_ in root.find('IDENTIFIERS'):
67                  d.append(('id.' + id_.tag, id_.text))
68
69              for attr in root.findall('SAMPLE_ATTRIBUTES/SAMPLE_ATTRIBUTE'):
70                  d.append(('attr.' + attr.find('TAG').text, attr.find('VALUE').text))
71
72              for key, value in d:
73                  print('{:25}: {}'.format(key, value))
74
75              print()
76
77     # --------------------------------------------------
78     if __name__ == '__main__':
79         main()
$ ./xml_ena.py ena-101.xml
>>>>>> ena-101.xml
sample.alias             : SAMD00024455
sample.accession         : DRS018892
sample.broker_name       : DDBJ
id.PRIMARY_ID            : DRS018892
id.EXTERNAL_ID           : SAMD00024455
id.SUBMITTER_ID          : SAMD00024455
attr.sample_name         : 100A
attr.collection_date     : 2013-08-15/2013-08-28
attr.depth               : 0.5m
attr.env_biome           : coastal biome
attr.env_feature         : natural environment
attr.env_material        : water
attr.geo_loc_name        : China:the East China Sea
attr.lat_lon             : 29.3 N 122.08 E
attr.project_name        : seawater bacterioplankton
attr.BioSampleModel      : MIMARKS.survey.water
attr.ENA-SPOT-COUNT      : 54843
attr.ENA-BASE-COUNT      : 13886949
attr.ENA-FIRST-PUBLIC    : 2015-02-15
attr.ENA-LAST-UPDATE     : 2018-08-15
```

SwissProt

The SwissProt format is one, like GenBank and EMBL, that allows for detailed
annotation of a sequence whereas FASTA/Q are primarily devoted to the se-
quence/quality and sometimes metadata/annotations are crudely shoved into
the header line. Parsing SwissProt, however, is no more difficult thanks to the
SeqIO module. Most of the interesting non-sequence data is in the annotations
which is a dictionary where the keys are strings like "accessions" and "keywords"

and the values are ints, strings, and lists.

Here is an example program to print out the accessions, keywords, and taxonomy in a SwissProt record:

```
$ cat -n swissprot.py
     1  #!/usr/bin/env python3
     2
     3  import argparse
     4  import sys
     5  from Bio import SeqIO
     6
     7
     8  # --------------------------------------------------
     9  def get_args():
    10      """get args"""
    11      parser = argparse.ArgumentParser(
    12          description='Parse Swissprot file',
    13          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    14
    15      parser.add_argument('file', metavar='FILE', help='Swissprot file')
    16
    17      return parser.parse_args()
    18
    19
    20  # --------------------------------------------------
    21  def die(msg='Something bad happened'):
    22      """print message and exit with error"""
    23      print(msg)
    24      sys.exit(1)
    25
    26
    27  # --------------------------------------------------
    28  def main():
    29      """main"""
    30      args = get_args()
    31      file = args.file
    32
    33      for i, record in enumerate(SeqIO.parse(file, "swiss"), start=1):
    34          print('{:3}: {}'.format(i, record.id))
    35          annotations = record.annotations
    36
    37          for annot_type in ['accessions', 'keywords', 'taxonomy']:
    38              if annot_type in annotations:
    39                  print('\tANNOT {}:'.format(annot_type))
    40                  val = annotations[annot_type]
    41                  if type(val) is list:
```

```
   42                      for v in val:
   43                          print('\t\t{}'.format(v))
   44                  else:
   45                      print('\t\t{}'.format(val))
   46
   47
   48
   49   # -------------------------------------------------
   50   if __name__ == '__main__':
   51       main()
$ ./swissprot.py input.swiss
  1: G5EEM5
   ANNOT accessions://
       Nematoda
       Chromadorea
       Rhabditida
       Rhabditoidea
       Rhabditidae
       Peloderinae
       Caenorhabditis
```

You should look at the sample "input.swiss" file to get a greater understanding of what is contained.

JSON

JSON stands for JavaScript Object Notation, and it has become the lingua franca of data exchange on the Internet. For our example, I will use the JSON that is returned by https://www.imicrobe.us/api/v1/samples/578. We need to `import json` and use `json.load` to read from an open file handle (there is also `loads` – load string) to parse the data from JSON into a Python dictionary. We could `print` that, but it's not nearly as pretty as printing the JSON which we can do with `json.dumps` (dump string) and the keyword argument `indent=4` to get nice indentation.

```
$ cat -n json_parse.py
     1  #!/usr/bin/env python3
     2
     3  import json
     4
     5  file = '578.json'
     6  data = json.load(open(file))
     7  print(json.dumps(data, indent=4))
$ ./json_parse.py | head -12
{
    "sample_id": 578,
```

```
    "project_id": 26,
    "sample_acc": "CAM_SMPL_GS108",
    "sample_name": "GS108",
    "sample_type": "Metagenome",
    "sample_description": "GS108",
    "url": "",
    "creation_date": "2018-07-06T04:43:09.000Z",
    "project": {
        "project_id": 26,
        "project_code": "CAM_PROJ_GOS",
```

If you `head 578.json`, you will see there is no whitespace, so this is a nicer way to look at the data; however, if all we wanted was to look at pretty JSON, we could do this:

```
$ python -m json.tool 578.json
```

Chapter 8

Introduction to Regular Expressions in Python

The term "regular expression" is a formal, linguistic term you might be interested to read about (https://en.wikipedia.org/wiki/Regular_language). For our purposes, regular expressions (AKA "regexes" or a "regex") is a way to formally describe some string of characters that we want to find. Regexes are an entirely separate DSL (domain-specific language) that we use inside Python, just like in the previous chapter we use SQL statements to communite with SQLite. While it's a bit of a drag to have to learn yet another language, the bonus is that you can use regular expressions in many places besides Python including with command line tools like `grep` and `awk` as well as within other languages like Perl and Rust.

We can `import re` to use the Python regular expression module and use it to search text. For instance, in the tic-tac-toe exercise, we needed to see if the `--player` argument was exactly one character that was either an 'X' or an 'O'. Here's code that can do that:

```python
for player in ['X', 'A', 'O', '5']:
    if len(player) == 1 and (player == 'X' or player == 'O'):
        print('{} OK'.format(player))
    else:
        print('{} bad'.format(player))
```

```
X OK
A bad
O OK
5 bad
```

A shorter way to write this could be:

```python
for player in ['X', 'A', 'B', '5']:
    if len(player) == 1 and player in 'XO':
        print('{} OK'.format(player))
    else:
        print('{} bad'.format(player))
```

```
X OK
A bad
B bad
5 bad
```

It's not too onerous, but it quickly gets worse as we get more complicated requirements. In that same exercise, we needed to check if `--state` was exactly 9 characters composed entirely of ".", "X", "O":

```
for state in ['XXX...OOO', 'XXX...OOA']:
    #print([(x, x in 'XO.') for x in state])
    print(state, 'OK' if len(state) == 9 and
          all(map(lambda x: x in 'XO.', state)) else 'No')

XXX...OOO OK
XXX...OOA No
```

Can we make this simpler? Well, when we were starting out with the Unix command line, one exercise had us using grep to look for lines that start with vowels. One solution was:

```
$ grep -io '^[aeiou]' scarlet.txt | sort | uniq -c
  59 A
  10 E
  91 I
  20 O
   6 U
 651 a
 199 e
 356 i
 358 o
 106 u
```

We used square brackets [] to enumerate all the vowels [aeiou] and used the -i flag to grep to indicate it should match case **insensitively**. Additionally, the ^ indicated that the match should occur at the start of the string. Those were regular expressions we were using.

The regex allows us to **describe** what we want rather than **implement** the code to find what we want. We can create a class of allowed characters with [XO] and additionally constraint it to be exactly one character wide with {1} after the class. (Note that {} for match length can be in the format {exactly}, {min,max}, {min,}, or {,max}.)

To use regular expressions:

```
import re
```

Now let's describe our pattern using a character class [XO] and the length {1}:

```
for player in ['X', 'O', 'A']:
    print(player, re.match('[XO]{1}', player))

X <_sre.SRE_Match object; span=(0, 1), match='X'>
O <_sre.SRE_Match object; span=(0, 1), match='O'>
A None
```

We can extend this to our state problem:

```
state = 'XXX...OOO'
print(state, re.match('[XO.]{9}', state))
```

173

```
XXX...OOO <_sre.SRE_Match object; span=(0, 9), match='XXX...OOO'>
```

```
state = 'XXX...OOA'
print(state, re.match('[XO.]{9}', state))
```

```
XXX...OOA None
```

Building regular expressions

How do we match a number?

```
print(re.match('1', '1'))
```

```
<_sre.SRE_Match object; span=(0, 1), match='1'>
```

But that only works for just "1"

```
print(re.match('2', '1'))
```

```
None
```

How do we match all the numbers from 0 to 9? We can create a character class that contains that range:

```
print(re.match('[0-9]', '1'))
```

```
<_sre.SRE_Match object; span=(0, 1), match='1'>
```

There is a short-hand for the character class [0-9] that is \d (digit)

```
re.match('\d', '1')
```

```
<_sre.SRE_Match object; span=(0, 1), match='1'>
```

But this only matches the first number we see:

```
re.match('\d', '123')
```

```
<_sre.SRE_Match object; span=(0, 1), match='1'>
```

We can use {} to indicate {min,max}, {min,}, {,max}, or {exactly}:

```
print(re.match('\d{1,4}', '8005551212'))
```

```
<_sre.SRE_Match object; span=(0, 4), match='8005'>
```

```
print(re.match('\d{1,}', '8005551212'))
```

```
<_sre.SRE_Match object; span=(0, 10), match='8005551212'>
```

```
print(re.match('\d{,5}', '8005551212'))
```

```
<_sre.SRE_Match object; span=(0, 5), match='80055'>
```

```
print(re.match('\d{8}', '8005551212'))
```

```
<_sre.SRE_Match object; span=(0, 8), match='80055512'>
```

match vs search

Note that we are using `re.match` which requires the regex to match **at the beginning of the string**:

```
print(re.match('\d{10}', 'That number to call is 8005551212!'))
```

```
None
```

If you want to match anywhere in the string, use `re.search`:

```
for s in ['123', 'abc456', '789def']:
    print(s, re.search('\d{3}', s))
```

```
123 <_sre.SRE_Match object; span=(0, 3), match='123'>
abc456 <_sre.SRE_Match object; span=(3, 6), match='456'>
789def <_sre.SRE_Match object; span=(0, 3), match='789'>
```

To anchor your match to the beginning of the string, use the `^`:

```
for s in ['123', 'abc456', '789def']:
    print(s, re.search('^\d{3}', s))
```

```
123 <_sre.SRE_Match object; span=(0, 3), match='123'>
abc456 None
789def <_sre.SRE_Match object; span=(0, 3), match='789'>
```

Use `$` for the end of the string:

```
for s in ['123', 'abc456', '789def']:
    print(s, re.search('\d{3}$', s))
```

```
123 <_sre.SRE_Match object; span=(0, 3), match='123'>
abc456 <_sre.SRE_Match object; span=(3, 6), match='456'>
789def None
```

And use both to say that the entire string from beginning to end must match:

```
for s in ['123', 'abc456', '789def']:
    print(s, re.search('^\d{3}$', s))
```

```
123 <_sre.SRE_Match object; span=(0, 3), match='123'>
abc123 None
123def None
```

Returning to our previous problem of trying to see if we got *exactly* one "X" or "O" for our tic-tac-toe player:

```
for player in ['X', 'O', 'XX', 'OO']:
    print(player, re.match('[XO]{1}', player))
```

```
X <_sre.SRE_Match object; span=(0, 1), match='X'>
O <_sre.SRE_Match object; span=(0, 1), match='O'>
XX <_sre.SRE_Match object; span=(0, 1), match='X'>
```

```
OO <_sre.SRE_Match object; span=(0, 1), match='O'>
```

The problem is that there is a match of `[XO]{1}` in the strings "XX" and "OO" – there *is* exactly one X or O at the beginning of those strings. Since `re.match` already anchors the match to the beginning of the string, we could just add `$` to the end of our pattern:

```python
for player in ['X', 'O', 'XX', 'OO']:
    print(player, re.match('[XO]{1}$', player))
```

```
X <_sre.SRE_Match object; span=(0, 1), match='X'>
O <_sre.SRE_Match object; span=(0, 1), match='O'>
XX None
OO None
```

Or use `re.search` with `^$` to indicate a match over the entire string:

```python
for player in ['X', 'O', 'XX', 'OO']:
    print(player, re.search('^[XO]{1}$', player))
```

```
X <_sre.SRE_Match object; span=(0, 1), match='X'>
O <_sre.SRE_Match object; span=(0, 1), match='O'>
XX None
OO None
```

Matching SSNs and Dates

What if we wanted to recognize a US SSN (social security number)? We will use `re.compile` to create the regex and use it in a `for` loop:

```python
ssn_re = re.compile('\d{3}-\d{2}-\d{4}')
for s in ['123456789', '123-456-789', '123-45-6789']:
    print('{}: {}'.format(s, ssn_re.match(s)))
```

```
123456789: None
123-456-789: None
123-45-6789: <_sre.SRE_Match object; span=(0, 11), match='123-45-6789'>
```

SSNs always use a dash (-) as a number separator, but dates do not.

```python
date_re = re.compile('\d{4}-\d{2}-\d{2}')
dates = ['1999-01-02', '1999/01/02']
for d in dates:
    print('{}: {}'.format(d, date_re.match(d)))
```

```
1999-01-02: <_sre.SRE_Match object; span=(0, 10), match='1999-01-02'>
1999/01/02: None
```

Just as we created a character class with `[0-9]` to represent all the numbers from 0 to 9, we can create a class to represent the separators "/" and "-" with `[/-]`. As regular expressions get longer, it makes sense to break each unit onto

a different line and use Python's literal string expression to join them into a single string. As a bonus, we can comment on each unit of the regex.

```python
date_re = re.compile('\d{4}'    # year
                     '[/-]'     # separator
                     '\d{2}'    # month
                     '[/-]'     # separator
                     '\d{2}')   # day

dates = ['1999-01-02', '1999/01/02']
for d in dates:
    print('{}: {}'.format(d, date_re.match(d)))

1999-01-02: <_sre.SRE_Match object; span=(0, 10), match='1999-01-02'>
1999/01/02: <_sre.SRE_Match object; span=(0, 10), match='1999/01/02'>
```

You may notice that certain elements are repeated. If we followed DRY (Don't Repeat Yourself), we might want to make variables to hold each piece, but then we could not use the literal string joining trick above. In that case, just go back to using **+** to join strings:

```python
sep = '[/-]'
four_digits = '\d{4}'
two_digits = '\d{2}'

date_re = re.compile(four_digits + # year
                     sep         + # separator
                     two_digits  + # month
                     sep         + # separator
                     two_digits)   # day

dates = ['1999-01-02', '1999/01/02']
for d in dates:
    print('{}: {}'.format(d, date_re.match(d)))

1999-01-02: <_sre.SRE_Match object; span=(0, 10), match='1999-01-02'>
1999/01/02: <_sre.SRE_Match object; span=(0, 10), match='1999/01/02'>
```

Dates are not always written YYYY-MM-DD where the month/day are zero-padded left, e.g., "01" instead of "1". How could we handle that? Change our `two_digits` from `\d{2}` (exactly two) to `\d{1,2}` (one or two):

```python
sep = '[/-]'
four_digits = '\d{4}'
two_digits = '\d{1,2}'

date_re = re.compile(four_digits + # year
                     sep         + # separator
                     two_digits  + # month
```

177

```
                    sep          + # separator
                    two_digits)   # day

dates = ['1999-01-01', '1999/01/02', '1999/1/2']
for d in dates:
    print('{}: {}'.format(d, date_re.match(d)))

1999-01-01: <_sre.SRE_Match object; span=(0, 10), match='1999-01-01'>
1999/01/02: <_sre.SRE_Match object; span=(0, 10), match='1999/01/02'>
1999/1/2: <_sre.SRE_Match object; span=(0, 8), match='1999/1/2'>
```

If we wanted to extract each part of the date (year, month, day), we can use parentheses () around the parts we want to capture into groups. The group "0" is the whole string that was match, and they are numbered sequentially after that for each group.

Can you change the regex to match all three strings?

```
date_re = re.compile('(\d{4})'    # capture year (group 1)
                     '[/-]'       # separator
                     '(\d{1,2})'  # capture month (group 2)
                     '[/-]'       # separator
                     '(\d{1,2})') # capture day (group 3)

dates = ['1999-01-02', '1999/1/2', '1999.01.01']
for d in dates:
    match = date_re.match(d)
    print('{}: {}'.format(d, 'match' if match else 'miss'))
    if match:
        print(match.groups())
        print('year:', match.group(1))
    print()

1999-01-01: match
('1999', '01', '01')
year: 1999

1999/01/01: match
('1999', '01', '01')
year: 1999

1999.01.01: miss
```

As we add more groups, it can be confusing to remember them by their positions, so we can name them with ?P<name> just inside the opening paren.

```
date_re = re.compile('(?P<year>\d{4})'
                     '[/-]'
                     '(?P<month>\d{1,2})'
```

```
                    '[/-]'
                    '(?P<day>\d{1,2})')

dates = ['1999-1-2', '1999/01/02', '1999.01.01']

for d in dates:
    match = date_re.match(d)
    print('{}: {}'.format(d, 'match' if match else 'miss'))
    if match:
        print('{} = year "{}" month "{}" day "{}"'.format(d,
                                           match.group('year'),
                                           match.group('month'),
                                           match.group('day')))

    print()
1999-1-2: match
1999-1-2 = year "1999" month "1" day "2"

1999/01/02: match
1999/01/02 = year "1999" month "01" day "02"

1999.01.01: miss
```

Matching US Phone Numbers

What if we wanted to match a US phone number?

```
phone_re = re.compile('(\d{3})'   # area code
                      ' '         # a space
                      '\d{3}'     # prefix
                      '-'         # dash
                      '\d{4}')    # line number
print(phone_re.match('(800) 555-1212'))
```

```
None
```

Why didn't that work?

What do those parentheses do again? They group!

So we need to indicate that the parens are literal things to match by using backslashes \ to escape them.

```
phone_re = re.compile('\('       # left paren
                      '\d{3}'    # area code
                      '\)'       # right paren
                      ' '        # space
                      '\d{3}'    # prefix
```

```
                     '-'       # dash
                     '\d{4}') # line number
print(phone_re.match('(800) 555-1212'))
```

```
<_sre.SRE_Match object; span=(0, 14), match='(800) 555-1212'>
```

We could also use character classes to make this more readable:

```
phone_re = re.compile('[(]'    # left paren
                     '\d{3}'   # area code
                     '[)]'     # right paren
                     ' '       # space
                     '\d{3}'   # prefix
                     '-'       # dash
                     '\d{4}') # line number
```

```
print(phone_re.match('(800) 555-1212'))
```

```
<_sre.SRE_Match object; span=(0, 14), match='(800) 555-1212'>
```

There is not always a space after the area code, and it may sometimes it may be more than one space (or a tab?). We can use the **\s** to indicate any type of whitespace and ***** to indicate zero or more:

```
phone_re = re.compile('[(]'    # left paren
                     '\d{3}'   # area code
                     '[)]'     # right paren
                     '\s*'     # zero or more spaces
                     '\d{3}'   # prefix
                     '-'       # dash
                     '\d{4}') # line number
phones = ['(800)555-1212', '(800) 555-1212', '(800)  555-1212']
for phone in phones:
    print('{}\t{}'.format(phone, phone_re.match(phone)))
```

```
(800)555-1212   None
(800) 555-1212  None
(800)  555-1212 <_sre.SRE_Match object; span=(0, 15), match='(800)  555-1212'>
```

When the parens around the area code are optional, usually there is a dash to separate the area code:

```
phone_re = re.compile('[(]?'   # optional left paren
                     '\d{3}'   # area code
                     '[)]?'    # optional right paren
                     '[-]?'    # optional dash
                     '\s*'     # zero or more whitespace
                     '\d{3}'   # prefix
                     '-'       # dash
                     '\d{4}') # line number
```

```
phones = ['(800)555-1212', '(800) 555-1212', '800-555-1212']
for phone in phones:
    print('{}\t{}'.format(phone, phone_re.match(phone)))

(800)555-1212   <_sre.SRE_Match object; span=(0, 13), match='(800)555-1212'>
(800) 555-1212  <_sre.SRE_Match object; span=(0, 14), match='(800) 555-1212'>
800-555-1212    <_sre.SRE_Match object; span=(0, 12), match='800-555-1212'>
```

This has the affect of matching a dash after parens which is generally not a valid format:

```
phone_re = re.compile('[(]?'    # optional left paren
                      '\d{3}'   # three digits
                      '[)]?'    # optional right paren
                      '[-]?'    # optional dash
                      '\s*'     # zero or more spaces
                      '\d{3}'   # three digits
                      '-'       # dash
                      '\d{4}')  # four digits

phone_re.match('(800)-555-1212')

<_sre.SRE_Match object; span=(0, 14), match='(800)-555-1212'>
```

We really have to create two regexes to handle these cases:

```
phone_re1 = re.compile('[(]'
                       '\d{3}'
                       '[)]'
                       '\s*'
                       '\d{3}'
                       '-'
                       '\d{4}')

phone_re2 = re.compile('\d{3}'
                       '-'
                       '\d{3}'
                       '-'
                       '\d{4}')

phones = ['(800)555-1212', '(800) 555-1212', '800-555-1212', '(800)-555-1212']
for phone in phones:
    match1 = phone_re1.match(phone)
    match2 = phone_re2.match(phone)
    print('{}\t{}'.format(phone, 'match' if match1 or match2 else 'miss'))

(800)555-1212   match
(800) 555-1212  match
```

```
800-555-1212     match
(800)-555-1212   miss
```

I worked with a graphic artist who always insisted on using dots as the number separator, and sometimes there are no separators at all. The combination of these two regexes find the valid formats and skip the invalid one.

```python
phone_re1 = re.compile('[(]'
                       '\d{3}'
                       '[)]'
                       '\s*'
                       '\d{3}'
                       '[.-]'
                       '\d{4}')

phone_re2 = re.compile('\d{3}'
                       '[.-]?'
                       '\d{3}'
                       '[.-]?'
                       '\d{4}')

phones = ['8005551212', '(800)555-1212', '(800) 555-1212',
          '800-555-1212', '(800)-555-1212', '800.555.1212']

for phone in phones:
    match = phone_re1.match(phone) or phone_re2.match(phone)
    print('{}\t{}'.format(phone, 'match' if match else 'miss'))
```

```
8005551212       match
(800)555-1212    match
(800) 555-1212   match
800-555-1212     match
(800)-555-1212   miss
800.555.1212     match
```

OK, now let's normalize the numbers by using parens to capture the area code, prefix, and line number and then create a standard representation.

```python
phone_re1 = re.compile('[(]'
                       '(\d{3})'   # group 1
                       '[)]'
                       '\s*'
                       '(\d{3})'   # group 2
                       '[.-]'
                       '(\d{4})')  # group 3

phone_re2 = re.compile('(\d{3})'   # group 1
                       '[.-]?'
```

```
                    '(\d{3})'  # group 2
                    '[.-]?'
                    '(\d{4})') # group 3

phones = ['8005551212', '(800)555-1212', '(800) 555-1212',
          '800-555-1212', '(800)-555-1212', '800.555.1212']

for phone in phones:
    match = phone_re1.match(phone) or phone_re2.match(phone)
    standard = '{}-{}-{}'.format(match.group(1),
                                 match.group(2),
                                 match.group(3)) if match else 'miss'
    print('{}\t{}'.format(phone, standard))

8005551212   800-555-1212
(800)555-1212    800-555-1212
(800) 555-1212   800-555-1212
800-555-1212     800-555-1212
(800)-555-1212   miss
800.555.1212     800-555-1212
```

And if we add named capture groups...

```
phone_re1 = re.compile('[(]'
                       '(?P<area_code>\d{3})'
                       '[)]'
                       '\s*'
                       '(?P<prefix>\d{3})'
                       '[.-]'
                       '(?P<line_num>\d{4})')

phone_re2 = re.compile('(?P<area_code>\d{3})'
                       '[.-]?'
                       '(?P<prefix>\d{3})'
                       '[.-]?'
                       '(?P<line_num>\d{4})')

phones = ['8005551212', '(800)555-1212', '(800) 555-1212',
          '800-555-1212', '(800)-555-1212', '800.555.1212']

for phone in phones:
    match = phone_re1.match(phone) or phone_re2.match(phone)
    standard = '{}-{}-{}'.format(match.group('area_code'),
                                 match.group('prefix'),
                                 match.group('line_num')) if match else 'miss'
    print('{}\t{}'.format(phone, standard))

8005551212   800-555-1212
```

```
(800)555-1212    800-555-1212
(800) 555-1212   800-555-1212
800-555-1212     800-555-1212
(800)-555-1212   miss
800.555.1212     800-555-1212
```

And if we add named capture groups and named groups in `format`:

```python
phone_re1 = re.compile('[(]'
                       '(?P<area_code>\d{3})'
                       '[)]'
                       '\s*(?P<prefix>\d{3})'
                       '[.-]'
                       '(?P<line_num>\d{4})')

phone_re2 = re.compile('(?P<area_code>\d{3})'
                       '[.-]?'
                       '(?P<prefix>\d{3})'
                       '[.-]?'
                       '(?P<line_num>\d{4})')

phones = ['8005551212', '(800)555-1212', '(800) 555-1212',
          '800-555-1212', '(800)-555-1212', '800.555.1212']

for phone in phones:
    match = phone_re1.match(phone) or phone_re2.match(phone)
    tmpl = '{area_code}-{prefix}-{line_num}'
    standard = tmpl.format(prefix=match.group('prefix'),
                           area_code=match.group('area_code'),
                           line_num=match.group('line_num')) if match else 'miss'
    print('{}\t{}'.format(phone, standard))

8005551212  800-555-1212
(800)555-1212    800-555-1212
(800) 555-1212   800-555-1212
800-555-1212     800-555-1212
(800)-555-1212   miss
800.555.1212     800-555-1212
```

ENA Metadata

Let's examine the ENA metadata from the XML parsing example. We see there are many ways that latitude/longitude have been represented:

```
$ ./xml_ena.py *.xml | grep lat_lon
attr.lat_lon          : 27.83387,-65.4906
attr.lat_lon          : 29.3 N 122.08 E
```

```
attr.lat_lon              : 28.56_-88.70377
attr.lat_lon              : 39.283N 76.611 W
attr.lat_lon              : 78 N 5 E
attr.lat_lon              : missing
attr.lat_lon              : 0.00 N, 170.00 W
attr.lat_lon              : 11.46'45.7" 93.01'22.3"
```

How can we go about parsing all the various ways this data has been encoded? Regular expressions provide us a way to describe in very specific way what we want.

Let's start just with the idea of matching a number (where "number" is a string that could be parsed into a number) like "27.83387":

```
print(re.search('\d', '27.83387'))
```

```
<_sre.SRE_Match object; span=(0, 1), match='2'>
```

The \d pattern means "any number" which is the same as [0-9] where the [] creates a class of characters and 0-9 expands to all the numbers from zero to nine. The problem is that it only matches one number, 2. Change it to \d+ to indicate "one or more numbers":

```
re.search('\d+', '27.83387')
```

```
<_sre.SRE_Match object; span=(0, 2), match='27'>
```

Now let's capture the decimal point:

```
re.search('\d+.', '27.83387')
```

```
<_sre.SRE_Match object; span=(0, 3), match='27.'>
```

You might think that's perfect, but the . has a special meaning in regex. It means "one of anything", so it matches this, too:

```
re.search('\d+.', '27x83387')
```

```
<_sre.SRE_Match object; span=(0, 3), match='27x'>
```

To indicate we want a literal . we have to make it \. (backslash-escape):

```
print(re.search('\d+\.', '27.83387'))
print(re.search('\d+\.', '27x83387'))
```

```
<_sre.SRE_Match object; span=(0, 3), match='27.'>
None
```

Notice that the second try returns nothing.

To capture the bit after the ., add more numbers:

```
re.search('\d+\.\d+', '27.83387')
```

```
<_sre.SRE_Match object; span=(0, 8), match='27.83387'>
```

But we won't always see floats. Can we make this regex match integers, too? We can indicate that part of a pattern is optional by putting a `?` after it. Since we need more than one thing to be optional, we need to wrap it in parens:

```
print(re.search('\d+\.\d+', '27'))
print(re.search('\d+(\.\d+)?', '27'))
print(re.search('\d+(\.\d+)?', '27.83387'))
```

```
None
<_sre.SRE_Match object; span=(0, 2), match='27'>
<_sre.SRE_Match object; span=(0, 8), match='27.83387'>
```

What if there is a negative symbol in front? Add `-?` (an optional dash) at the beginning:

```
print(re.search('-?\d+(\.\d+)?', '-27.83387'))
print(re.search('-?\d+(\.\d+)?', '27.83387'))
print(re.search('-?\d+(\.\d+)?', '-27'))
print(re.search('-?\d+(\.\d+)?', '27'))
```

```
<_sre.SRE_Match object; span=(0, 9), match='-27.83387'>
<_sre.SRE_Match object; span=(0, 8), match='27.83387'>
<_sre.SRE_Match object; span=(0, 3), match='-27'>
<_sre.SRE_Match object; span=(0, 2), match='27'>
```

Sometimes we actually find a `+` at the beginning, so we can make an optional character class `[+-]?`:

```
print(re.search('[+-]?\d+(\.\d+)?', '-27.83387'))
print(re.search('[+-]?\d+(\.\d+)?', '+27.83387'))
print(re.search('[+-]?\d+(\.\d+)?', '27.83387'))
```

```
<_sre.SRE_Match object; span=(0, 9), match='-27.83387'>
<_sre.SRE_Match object; span=(0, 9), match='+27.83387'>
<_sre.SRE_Match object; span=(0, 8), match='27.83387'>
```

Now we can match things that basically look like a floating point number or an integer, both positive and negative.

Usually the data we want to find it part of a larger string, however, and the above fails to capture more than one thing, e.g.:

```
print(re.search('[+-]?\d+(\.\d+)?', 'Lat is "-27.83387" and lon is "+132.43."'))
```

```
<_sre.SRE_Match object; span=(8, 17), match='-27.83387'>
```

We really need to match more than once using our pattern matching to extract data. We saw earlier that we can use parens to group optional patterns, but the parens also end up creating a **capture group** that we can refer to by position:

```
re.findall('([+-]?\d+(\.\d+)?)','Lat is "-27.83387" and lon is "+132.43."')
```

```
[('-27.83387', '.83387'), ('+132.43', '.43')]
```

OK, it was a bit unexpected that we have matches for both the whole float and the decimal part. This is because of the dual nature of the parens, and in the case of using them to group the optional part we are also creating another capture. If we change () to (?:), we make this a non-capturing group:

```
re.findall('([+-]?\d+(?:\.\d+)?)', 'Lat is "-27.83387" and lon is "+132.43."')

['-27.83387', '+132.43']
```

There are many resources you can use to thoroughly learn regular expressions, so I won't try to cover them completely here. I will mostly try to introduce the general idea and show you some useful regexes you could steal.

Here is an example of how you can embed regexes in your Python code. This version can parse all the versions of latitude/longitude shown above. This code uses parens to create capture groups which it then uses `match.group(n)` to extract:

```
$ cat -n parse_lat_lon.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import re
     5  import sys
     6
     7  args = sys.argv[1:]
     8
     9  if len(args) != 1:
    10      print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  file = args[0]
    14
    15  float_re = r'[+-]?\d+\.*\d*'
    16  ll1 = re.compile('(' + float_re + ')\s*[,_]\s*(' + float_re + ')')
    17  ll2 = re.compile('(' + float_re + ')(?:\s*([NS]))?(?:\s*,)?\s+(' + float_re +
    18                   ')(?:\s*([EW])?)')
    19  loc_hms = r"""
    20  \d+\.\d+'\d+\.\d+"
    21  """.strip()
    22  ll3 = re.compile('(' + loc_hms + ')\s+(' + loc_hms + ')')
    23
    24  for line in open(file):
    25      line = line.rstrip()
    26      ll_match1 = ll1.search(line)
    27      ll_match2 = ll2.search(line)
    28      ll_match3 = ll3.search(line)
    29
```

```
30      if ll_match1:
31          lat, lon = ll_match1.group(1), ll_match1.group(2)
32          lat = float(lat)
33          lon = float(lon)
34          print('lat = {}, lon = {}'.format(lat, lon))
35      elif ll_match2:
36          lat, lat_dir, lon, lon_dir = ll_match2.group(
37              1), ll_match2.group(2), ll_match2.group(
38                  3), ll_match2.group(4)
39          lat = float(lat)
40          lon = float(lon)
41
42          if lat_dir == 'S':
43              lat *= -1
44
45          if lon_dir == 'W':
46              lon *= -1
47          print('lat = {}, lon = {}'.format(lat, lon))
48      elif ll_match3:
49          lat, lon = ll_match3.group(1), ll_match3.group(2)
50          print('lat = {}, lon = {}'.format(lat, lon))
51      else:
52          print('No match: "{}"'.format(line))
$ cat lat_lon.txt
attr.lat_lon            : 27.83387,-65.4906
attr.lat_lon            : 29.3 N 122.08 E
attr.lat_lon            : 28.56_-88.70377
This line will not be included
attr.lat_lon            : 39.283N 76.611 W
attr.lat_lon            : 78 N 5 E
attr.lat_lon            : missing
attr.lat_lon            : 0.00 N, 170.00 W
attr.lat_lon            : 11.46'45.7" 93.01'22.3"
$ ./parse_lat_lon.py lat_lon.txt
lat = 27.83387, lon = -65.4906
lat = 29.3, lon = 122.08
lat = 28.56, lon = -88.70377
No match: "This line will not be included"
lat = 39.283, lon = -76.611
lat = 78.0, lon = 5.0
No match: "attr.lat_lon            : missing"
lat = 0.0, lon = -170.0
lat = 11.46'45.7", lon = 93.01'22.3"
```

We see a similar problem with "collection_date":

```
$ ./xml_ena.py *.xml | grep collection
```

```
attr.collection_date      : March 24, 2014
attr.collection_date      : 2013-08-15/2013-08-28
attr.collection_date      : 20100910
attr.collection_date      : 02-May-2012
attr.collection_date      : Jul-2009
attr.collection_date      : missing
attr.collection_date      : 2013-12-23
attr.collection_date      : 5/04/2012
```

Imagine how you might go about parsing all these various representations of dates. Be aware that parsing date/time formats is so problematic and ubiquitous that many people have already written modules to assist you!

To run the code below, you will need to install the `dateparser` module:

```
$ python3 -m pip install dateparser
```

```python
import dateparser
for date in ['March 24, 2014',
             '2013-08-15',
             '20100910',
             '02-May-2012',
             'Jul-2009',
             '5/04/2012']:

    print('{:15}\t{}'.format(date, dateparser.parse(date)))
```

```
March 24, 2014  2014-03-24 00:00:00
2013-08-15      2013-08-15 00:00:00
20100910        2000-02-01 09:01:00
02-May-2012     2012-05-02 00:00:00
Jul-2009        2009-07-26 00:00:00
5/04/2012       2012-05-04 00:00:00
```

You can see it's not perfect, e.g., "20100910" should be "2010-09-10" and "Jul-2009" should not resolve to the 26th of July, but, honestly, what should it be? (Is the 1st any better?!) Still, this saves you writing a lot of code. And, trust me, **THIS IS REAL DATA**! While trying to parse latitude, longitude, collection date, and depth for 35K marine metagenomes from the ENA, I wrote a hundreds of lines of code and dozens of regular expressions!

Exercises

Write the regular expressions to parse the year, month, and day from the following date formats found in SRA metadata. When no day is present, e.g., "2/14," use "01" for the day.

```
d1 = "2012-03-09T08:59"
print(d1, re.match('', d1))

2012-03-09T08:59 <_sre.SRE_Match object; span=(0, 0), match=''>

d2 = "2012-03-09T08:59:03"

d3 = "2017-06-16Z"

d4 = "2015-01"

d5 = "2015-01/2015-02"

d6 = "2015-01-03/2015-02-14"

d7 = "20100910"

d8 = "12/06"

d9 = "2/14"

d10 = "2/14-12/15"

d11 = "2017-06-16Z"

# "Excel" format! What is that?! Look it up.
d12 = "34210"

d13 = "Dec-2015"

d14 = "March-2017"

d15 = "May, 2017"

d16 = "March-April 2017"

d17 = "July of 2011"

d18 = "2008 August"
```

Now combine all your code from the previous cell to normalize all the dates into the same format.

```
dates = ["2012-03-09T08:59", "2012-03-09T08:59:03", "2017-06-16Z",
         "2015-01", "2015-01/2015-02", "2015-01-03/2015-02-14",
         "20100910", "12/06", "2/14", "2/14-12/15", "2017-06-16Z",
         "34210", "Dec-2015", "March-2017", "May, 2017",
         "March-April 2017", "July of 2011", "2008 August"]


for date in dates:
    year = '1999'
    month = '01'
    day = '01'
    print('{}-{}-{}\t{}'.format(year, month, day, date))
```

```
1999-01-01  2012-03-09T08:59
1999-01-01  2012-03-09T08:59:03
1999-01-01  2017-06-16Z
1999-01-01  2015-01
1999-01-01  2015-01/2015-02
1999-01-01  2015-01-03/2015-02-14
1999-01-01  20100910
1999-01-01  12/06
1999-01-01  2/14
1999-01-01  2/14-12/15
1999-01-01  2017-06-16Z
1999-01-01  34210
1999-01-01  Dec-2015
1999-01-01  March-2017
1999-01-01  May, 2017
1999-01-01  March-April 2017
1999-01-01  July of 2011
1999-01-01  2008 August
```

Chapter 9

Writing Simple Games in Python

> "Treat your code like poetry and take it to the edge of the bare minimum." - ILYO

Games are a terrific way to learn. If you take something simple you know well, you have all the information you need to complete it. Something simple like tic-tac-toe – you know you need a board, some way for the user to select a cell, you need to keep track of who's playing (X or O), when they've made a bad move, and when someone has won. Games often need random values, interact with the user, employ infinite loops – in short, they are fascinating and fun to program and play.

Twelve Days of Christmas

Here is an implementation of the "Twelve Days of Christmas" song. It uses two loops to count up from 1 for each step and then a countdown from each step back to 1. Notice I use integers as the keys to the dictionaries.

```
$ cat -n twelve_days.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date   : 2019-03-19
     5  Purpose: Twelve Days of Christmas
     6  """
     7
     8  import argparse
     9  import sys
    10
    11
    12  # --------------------------------------------------
    13  def get_args():
    14      """get command-line arguments"""
    15      parser = argparse.ArgumentParser(
    16          description='Twelve Days of Christmas',
    17          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    18
    19      parser.add_argument(
    20          '-o',
    21          '--outfile',
    22          help='Outfile (STDOUT)',
```

```
23              metavar='str',
24              type=str,
25              default='')
26
27      parser.add_argument(
28          '-n',
29          '--number_days',
30          help='Number of days to sing',
31          metavar='int',
32          type=int,
33          default=12)
34
35      return parser.parse_args()
36
37
38  # --------------------------------------------------
39  def warn(msg):
40      """Print a message to STDERR"""
41      print(msg, file=sys.stderr)
42
43
44  # --------------------------------------------------
45  def die(msg='Something bad happened'):
46      """warn() and exit with error"""
47      warn(msg)
48      sys.exit(1)
49
50
51  # --------------------------------------------------
52  def main():
53      """Make a jazz noise here"""
54      args = get_args()
55      out_file = args.outfile
56      num_days = args.number_days
57
58      out_fh = open(out_file, 'wt') if out_file else sys.stdout
59
60      days = {
61          12: 'Twelve drummers drumming',
62          11: 'Eleven pipers piping',
63          10: 'Ten lords a leaping',
64          9: 'Nine ladies dancing',
65          8: 'Eight maids a milking',
66          7: 'Seven swans a swimming',
67          6: 'Six geese a laying',
68          5: 'Five gold rings',
```

```
69              4: 'Four calling birds',
70              3: 'Three French hens',
71              2: 'Two turtle doves',
72              1: 'a partridge in a pear tree',
73          }
74
75      cardinal = {
76              12: 'twelfth',
77              11: 'eleven',
78              10: 'tenth',
79              9: 'ninth',
80              8: 'eighth',
81              7: 'seventh',
82              6: 'sixth',
83              5: 'fifth',
84              4: 'fourth',
85              3: 'third',
86              2: 'second',
87              1: 'first',
88          }
89
90      if not num_days in days:
91          die('Cannot sing "{}" days'.format(num_days))
92
93      def ucfirst(s):
94          return s[0].upper() + s[1:]
95
96      for i in range(1, num_days + 1):
97          first = 'On the {} day of Christmas,\nMy true love gave to me,'
98          out_fh.write(first.format(cardinal[i]) + '\n')
99          for j in reversed(range(1, i + 1)):
100             if j == 1:
101                 if i == 1:
102                     out_fh.write('{}.\n'.format(ucfirst(days[j])))
103                 else:
104                     out_fh.write('And {}.\n'.format(days[j]))
105             else:
106                 out_fh.write('{},\n'.format(days[j]))
107
108         if i < max(days.keys()):
109             out_fh.write('\n')
110
111
112  # -------------------------------------------------
113  if __name__ == '__main__':
114      main()
```

Guessing Game

Let's write a simple program where the user has to guess a random number.
First, let's sketch out some pseudo-code:

```
Establish the range of numbers allowed and how many times the player can guess
Pick a random number in that range
Start a loop
    Ask the user for a guess
    Quit if the user asks
    Make sure the guess is a number and in the allowed range
    If the number was correctly guessed, stop and tell the user they won
    Let the user know if the number is high or low
    If the user has guessed too many times, stop and tell the user they lost (also insult th
```

Here's how it looks being played. Note that I use a binary search where I divide
the search space in half on each guess. I can usually guess the number correctly
in 5 guesses when the range is 1-50.

```
$ ./guess.py
Guess a number between 1 and 50 (q to quit): 25
"25" is too low.
Guess a number between 1 and 50 (q to quit): 37
"37" is too low.
Guess a number between 1 and 50 (q to quit): 45
"45" is too high.
Guess a number between 1 and 50 (q to quit): 40
"40" is too high.
Guess a number between 1 and 50 (q to quit): 38
"38" is correct. You win!
$ ./guess.py -x 100
Guess a number between 1 and 100 (q to quit): 50
"50" is too low.
Guess a number between 1 and 100 (q to quit): 75
"75" is too high.
Guess a number between 1 and 100 (q to quit): 62
"62" is too high.
Guess a number between 1 and 100 (q to quit): 55
"55" is too low.
Guess a number between 1 and 100 (q to quit): 58
"58" is too low.
Too many guesses, loser! The number was "59."
$ ./guess.py
Guess a number between 1 and 50 (q to quit): quit
Now you will never know the answer.
```

As usual, we'll start with `new_py.py`, and I'll use `argparse` to get the min/max

range with defaults of 1/50 and set the number of guesses to 5:

```
$ cat -n guess.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:  Ken Youens-Clark <kyclark@gmail.com>
     4  Purpose: Guess-the-number game
     5  """
     6
     7  import argparse
     8  import random
     9  import re
    10  import sys
    11
    12
    13  # --------------------------------------------------
    14  def get_args():
    15      """get args"""
    16      parser = argparse.ArgumentParser(
    17          description='Guessing game',
    18          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    19
    20      parser.add_argument(
    21          '-m',
    22          '--min',
    23          help='Minimum value',
    24          metavar='int',
    25          type=int,
    26          default=1)
    27
    28      parser.add_argument(
    29          '-x',
    30          '--max',
    31          help='Maximum value',
    32          metavar='int',
    33          type=int,
    34          default=50)
    35
    36      parser.add_argument(
    37          '-g',
    38          '--guesses',
    39          help='Number of guesses',
    40          metavar='int',
    41          type=int,
    42          default=5)
    43
```

```
44      return parser.parse_args()
45
46
47  # ----------------------------------------------------
48  def warn(msg):
49      """Print a message to STDERR"""
50      print(msg, file=sys.stderr)
51
52
53  # ----------------------------------------------------
54  def die(msg='Something bad happened'):
55      """warn() and exit with error"""
56      warn(msg)
57      sys.exit(1)
58
59
60  # ----------------------------------------------------
61  def main():
62      """main"""
63      args = get_args()
64      low = args.min
65      high = args.max
66      guesses_allowed = args.guesses
67      secret = random.randint(low, high)
68
69      if low < 1:
70          die('--min "{}" cannot be lower than 1'.format(low))
71
72      if guesses_allowed < 1:
73          die('--guesses "{}" cannot be lower than 1'.format(high))
74
75      if low > high:
76          die('--min "{}" is higher than --max "{}"'.format(low, high))
77
78      prompt = 'Guess a number between {} and {} (q to quit): '.format(low, high)
79      num_guesses = 0
80
81      while True:
82          guess = input(prompt)
83          num_guesses += 1
84
85          if re.match('q(uit)?', guess):
86              die('Now you will never know the answer.')
87
88          if not guess.isdigit():
89              warn('"{}" is not a number'.format(guess))
```

```
 90                    continue
 91
 92            num = int(guess)
 93
 94            if not low <= num <= high:
 95                print('Number "{}" is not in the allowed range'.format(num))
 96            elif num == secret:
 97                print('"{}" is correct. You win!'.format(num))
 98                break
 99            else:
100                print('"{}" is too {}.'.format(num, 'low'
101                                                 if num < secret else 'high'))
102
103            if num_guesses >= guesses_allowed:
104                die('Too many guesses, loser! The number was "{}."'.format(secret))
105
106
107  # --------------------------------------------------
108  if __name__ == '__main__':
109      main()
```

The `get_args` function will ensure we get some values for the range and number of guesses, but we should always assume garbage from the user, so we have to check them. To get a random number in our given range, we use the `random` module's `randint` function:

```
secret = random.randint(low, high)
```

If you ever need to "flip a coin" in your code, you can do this:

```
>>> import random
>>> random.randint(0,1)
1
>>> random.randint(0,1)
0
>>> random.randint(0,1)
0
```

The meat of the program will be an infinite loop where we keep asking the user:

```
prompt = 'Guess a number between {} and {} (q to quit): '.format(low, high)
```

Before we enter that loop, we'll need a variable to keep track of the number of guesses the user has made. This is the lone piece of "state" we need to track. Other games can have many pieces of information you need to track.

```
num_guesses = 0
```

The beginning of the play loop looks like this:

```
while True:
```

```
    guess = input('[{}] {}'.format(num_guesses, prompt))
    num_guesses += 1
```

Here I want the user to know how many guesses they've made so far. We want to give them a way out, so they can enter "q" (or "quit") to quit. I chose to use the `re` module for regular expressions so I can identify a string that is either "q" or "quit". The bit in `()?` is considered optional because the parens group it and the question mark makes it optional:

```
if re.match('q(uit)?', guess.lower()):
    die('Now you will never know the answer.')
```

The input from the user will be a string, and we are going to need to convert it to an integer to see if it is the secret number. Before we do that, we must check that it is a digit. We can use the `isdigit` method that all strings have. Look at `help(str)` in your Python REPL to see other useful methods like `isalnum`, `isalpha`, `islower`, etc.:

```
if not guess.isdigit():
    warn('"{}" is not a number'.format(guess))
    continue
```

If it's not a digit, we `continue` to go to the next iteration of the loop. If we move ahead, then it's OK to convert the guess by using the `int` method to coerce the string the user typed into an integer value:

```
>>> int('8')
8
>>> type(int('8'))
<type 'int'>
```

There's an alternate way to handle the conversion of the guess using a `try/except` block (see `guess-try.py`) where the call to `int` is wrapped in a `try` block that has an `except` block that catches the exception that Python throws when it fails. Cf https://docs.python.org/3/tutorial/errors.html:

```
num = 0
try:
    num = int(guess)
except:
    warn('"{}" is not an integer'.format(guess))
    continue
```

Now we need to determine if the user has guessed too many times, if the number if too high or low, or if they've won the game. Lastly we see if the user has exceeded the maximum number of guesses:

```
if not low <= num <= high:
    print('Number "{}" is not in the allowed range'.format(num))
elif num == secret:
    print('"{}" is correct. You win!'.format(num))
```

```
        break
else:
    print('"{}" is too {}.'.format(num, 'low'
                                    if num < secret else 'high'))

if num_guesses >= guesses_allowed:
    die('Too many guesses, loser! The number was "{}."'.format(secret))
```

Hangman

Here is an implementation of the game "Hangman" that uses dictionaries to maintain the "state" of the program – that is, all the information needed for each round of play such as the word being guessed, how many misses the user has made, which letters have been guessed, etc. The program uses the `argparse` module to gather options from the user while providing default values so that nothing needs to be provided. The `main` function is used just to gather the parameters and then run the `play` function which recursively calls itself, each time passing in the new "state" of the program. Inside `play`, we use the `get` method of `dict` to safely ask for keys that may not exist and use defaults. When the user finishes or quits, `play` will simply call `sys.exit` to stop.

Here is what it looks like being played:

```
$ ./hangman.py
_ _ _ _ _ _ _ _ _ (Misses: 0)
Your guess? ("?" for hint, "!" to quit) a
_ a _ _ _ _ a _ _ (Misses: 0)
Your guess? ("?" for hint, "!" to quit) e
e a _ _ _ _ a _ _ (Misses: 0)
Your guess? ("?" for hint, "!" to quit) i
e a _ _ _ _ a _ _ (Misses: 1)
Your guess? ("?" for hint, "!" to quit) o
e a _ _ _ _ a _ _ (Misses: 2)
Your guess? ("?" for hint, "!" to quit) u
e a _ _ _ _ a _ _ (Misses: 3)
Your guess? ("?" for hint, "!" to quit) ?
e a _ _ h _ a _ _ (Misses: 4)
Your guess? ("?" for hint, "!" to quit) t
e a _ t h _ a _ _ (Misses: 4)
Your guess? ("?" for hint, "!" to quit) r
e a r t h _ a r _ (Misses: 4)
Your guess? ("?" for hint, "!" to quit) w
e a r t h w a r _ (Misses: 4)
Your guess? ("?" for hint, "!" to quit) d
You win. You guessed "earthward" with "4" misses!
```

Here is the code:

```
$ cat -n hangman.py
     1  #!/usr/bin/env python3
     2  """
     3  Author:   Ken Youens-Clark <kyclark@gmail.com>
     4  Purpose: Hangman game
     5  """
     6
     7  import argparse
     8  import os
     9  import random
    10  import re
    11  import sys
    12
    13
    14  # --------------------------------------------------
    15  def get_args():
    16      """parse arguments"""
    17      parser = argparse.ArgumentParser(
    18          description='Hangman',
    19          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    20
    21      parser.add_argument(
    22          '-l', '--maxlen', help='Max word length', type=int, default=10)
    23
    24      parser.add_argument(
    25          '-n', '--minlen', help='Min word length', type=int, default=5)
    26
    27      parser.add_argument(
    28          '-m', '--misses', help='Max number of misses', type=int, default=10)
    29
    30      parser.add_argument(
    31          '-w',
    32          '--wordlist',
    33          help='Word list',
    34          type=str,
    35          default='/usr/share/dict/words')
    36
    37      return parser.parse_args()
    38
    39  # --------------------------------------------------
    40  def bail(msg):
    41      """Print a message to STDOUT and quit with no error"""
    42      print(msg)
    43      sys.exit(0)
```

```
44
45   # -------------------------------------------------
46   def warn(msg):
47       """Print a message to STDERR"""
48       print(msg, file=sys.stderr)
49
50
51   # -------------------------------------------------
52   def die(msg='Something bad happened'):
53       """warn() and exit with error"""
54       warn(msg)
55       sys.exit(1)
56
57   # -------------------------------------------------
58   def main():
59       """main"""
60       args = get_args()
61       max_len = args.maxlen
62       min_len = args.minlen
63       max_misses = args.misses
64       wordlist = args.wordlist
65
66       if not os.path.isfile(wordlist):
67           die('--wordlist "{}" is not a file.'.format(wordlist))
68
69       if min_len < 1:
70           die('--minlen must be positive')
71
72       if not 3 <= max_len <= 20:
73           die('--maxlen should be between 3 and 20')
74
75       if min_len > max_len:
76           die('--minlen ({}) is greater than --maxlen ({})'.format(
77               min_len, max_len))
78
79       good_word = re.compile('^[a-z]{' + str(min_len) + ',' + str(max_len) + '}$')
80       words = [w for w in open(wordlist).read().split() if good_word.match(w)]
81       word = random.choice(words)
82       play({'word': word, 'max_misses': max_misses})
83
84
85   # -------------------------------------------------
86   def play(state):
87       """Loop to play the game"""
88       word = state.get('word') or ''
89
```

```
90       if not word: die('No word!')
91
92       guessed = state.get('guessed') or list('_' * len(word))
93       prev_guesses = state.get('prev_guesses') or set()
94       num_misses = state.get('num_misses') or 0
95       max_misses = state.get('max_misses') or 0
96
97       if ''.join(guessed) == word:
98           msg = 'You win. You guessed "{}" with "{}" miss{}!'
99           bail(msg.format(word, num_misses, '' if num_misses == 1 else 'es'))
100
101      if num_misses >= max_misses:
102          bail('You lose, loser!  The word was "{}."'.format(word))
103
104      print('{} (Misses: {})'.format(' '.join(guessed), num_misses))
105      new_guess = input('Your guess? ("?" for hint, "!" to quit) ').lower()
106
107      if new_guess == '!':
108          bail('Better luck next time, loser.')
109      elif new_guess == '?':
110          new_guess = random.choice([x for x in word if x not in guessed])
111          num_misses += 1
112
113      if not re.match('^[a-z]$', new_guess):
114          print('"{}" is not a letter'.format(new_guess))
115          num_misses += 1
116      elif new_guess in prev_guesses:
117          print('You already guessed that')
118      elif new_guess in word:
119          prev_guesses.add(new_guess)
120          last_pos = 0
121          while True:
122              pos = word.find(new_guess, last_pos)
123              if pos < 0:
124                  break
125              elif pos >= 0:
126                  guessed[pos] = new_guess
127                  last_pos = pos + 1
128      else:
129          num_misses += 1
130
131      play({
132          'word': word,
133          'guessed': guessed,
134          'num_misses': num_misses,
135          'prev_guesses': prev_guesses,
```

```
136            'max_misses': max_misses
137        })
138
139
140  # ------------------------------------------------
141  if __name__ == '__main__':
142      main()
```

NB: I would mention that my approach to recursively calling the `play` function with a dictionary for state rather than creating an infinite `while` loop was influenced by my experience programming in the Elm language.

Some notes on code:

I don't want to guess at words that are too short or too long, so we set up `min_len` and `max_len` variables and then use those to build a regular expression that decribes a string that is composed of alphabet characters in that range of length:

```
regex = re.compile('^[a-z]{' + str(min_len) + ',' + str(max_len) + '}$')
```

To visualize this, pretend we set 5 and 10 for lower and upper bounds:

```
>>> min_len = 5
>>> max_len = 10
>>> '^[a-z]{' + str(min_len) + ',' + str(max_len) + '}$'
'^[a-z]{5,10}$'
```

Now perhaps you can see the regex we've created?

```
1 2     3      4
^ [a-z] {5,10} $
```

1. The beginning of the string
2. The character class composed of the letters from "a" to "z"
3. A length from 5 to 10
4. The end of the string

I chose to use `re.compile` to turn this into a variable containing the regex so I can use it:

```
>>> good_word = re.compile('^[a-z]{' + str(min_len) + ',' + str(max_len) + '}$')
>>> good_word.match('foo')
>>> good_word.match('foobar')
<_sre.SRE_Match object at 0x103f42ed0>
```

There's a lot packed into this line:

```
words = [w for w in open(wordlist).read().split() if good_word.match(w)]
```

So I `open` the words file and `read` it and them immediately `split` into words. Then I use a "list comprehension" to set up a little `for` loop over each word to take it `if` the word matches the regex. I could have written it more verbosely,

but this way is quite succinct and correct in a way that is actually harder to get right in a longer version like so:

```
words = []
for line in open(wordlist):
    for word in line.split():
        if good_word.match(word):
            words.append(word)
```

Similar to the guessing game, we need to randomly choose from our `words` which the `random.choice` function does exactly:

```
>>> import random
>>> random.choice(['foo', 'bar', 'baz'])
'bar'
>>> random.choice(['foo', 'bar', 'baz'])
'foo'
```

With that, we can launch into the `play` with a minimal state:

```
play({'word': word, 'max_misses': max_misses})
```

The `play` function is defined receiving a single `state` variable which is expected to be a dictionary. I could have passed in each part of the state individually as named variables, but this way seems cleaner to me:

```
def play(state):
```

The first time through `play`, there will be no previous guesses, so I use `dict.get` to ask for this so my code won't blow up. If nothing is available, I create a new string by multiplying the `_` (underscore) character by the length of the word where the underscore will indicate to the user where a letter has not been guessed. Since I want to store this as a list and not a string, I use `list` to convert the string:

```
guessed = state.get('guessed') or list('_' * len(word))
```

I'd like to keep track of all the letters the user has previously guessed so that I can tell them when they guess the same letter twice. The best data structure for this is a `set` which is essentially a dictionary with values of `1` – we only care is a key is present or absent so the value is irrelevant.

```
prev_guesses = state.get('prev_guesses') or set()
```

First I need to see if the user has guessed the correct word:

```
if ''.join(guessed) == word:
    msg = 'You win. You guessed "{}" with "{}" miss{}!'
    bail(msg.format(word, num_misses, '' if num_misses == 1 else 'es'))
```

The `bail` function is one I wrote just for this program as there are several places where I needed to `print` a message and `exit` *without* an error code.

To get a new guess from the user, I use `input` and chain it to the `lower` method of the returned string to lowercase the value:

```
new_guess = input('Your guess? ("?" for hint, "!" to quit) ').lower()
```

Because `q` is a valid input from the user, I can't use it to `quit` so I decided to use `!`. I also wanted to show mercy by allowing hints with the `?`. I this is implemented using another list comprehension to find all the letters in `word` that are *not* in the `guessed` set and then randomly select from that list:

```
if new_guess == '!':
    bail('Better luck next time, loser.')
elif new_guess == '?':
    new_guess = random.choice([x for x in word if x not in guessed])
    num_misses += 1
```

I use a regex similar to the `good_word` to see if we have exactly one character that is a letter:

```
>>> re.match('^[a-z]$', 'x')
<_sre.SRE_Match object at 0x103f76850>
>>> re.match('^[a-z]$', '4')
>>> re.match('^[a-z]$', '>')
```

Sets make it easy to check for membership:

```
elif new_guess in prev_guesses:
        print('You already guessed that')
```

Then we check if the guess is a character in the word; if so, add it to our previous guesses:

```
elif new_guess in word:
    prev_guesses.add(new_guess)
```

This next bit is tricky. We need to `find` the position(s) of the character in the word. Since the character may occur more than once, we need to keep track of the last position where we found it and use that as the second optional argument to `find`. E.g., in "foo" the "o" occurs twice. Keeping in mind zero-based counting:

```
>>> zip(range(3), 'foo')
[(0, 'f'), (1, 'o'), (2, 'o')]
```

`'foo'.find('o')` will always return `1` unless we tell it to start looking after that position:

```
>>> 'foo'.find('o')
1
>>> 'foo'.find('o')
1
>>> 'foo'.find('o', 2)
2
```

So we set up a variable `last_pos` to set to *after* whatever `find` returns. We need that in order to turn the underscore in `guessed` into the actual letter. Note that `find` will return `-1` to indicate no matches.

```
last_pos = 0
while True:
    pos = word.find(new_guess, last_pos)
    if pos < 0:
        break
    elif pos >= 0:
        guessed[pos] = new_guess
        last_pos = pos + 1
```

Finally we call `play` again but with a new state:

```
play({
    'word': word,
    'guessed': guessed,
    'num_misses': num_misses,
    'prev_guesses': prev_guesses,
    'max_misses': max_misses
})
```

This is an example of a recursive algorithm (https://readtheplaque.com/plaque/the-toronto-recursive-history-project), and they only work if you first handle the "base case." For Hangman, that is where the user guesses the word or exceeds the number of guesses, both of which will `bail` on the program; otherwise, the program continues to the next iteration.

Chapter 10

Logging in Python

So far we've use `print` statements that go to STDOUT and the `warn` function that makes is slightly more convenient to write to STDERR. The trouble with this approach to writing and debugging code is that you need to remove all the `print`/`warn` statements prior to releasing your code or running your tests. With the `logging` module (https://docs.python.org/3/library/logging.html), you can sprinkle messages to yourself liberally throughout your code and chose *at run time* which ones to see.

Like with `random.seed`, calls to the `logging` module affect the **global state** of how logging happens. First you need to set up how the logging will happen using the `basicConfig` (https://docs.python.org/3/library/logging.html#logging.basicConfig). Typically you will set log message to go to a `filename` (if you don't indicate a filename then messages go to STDERR) with the `filemode` of "w" (write, which will overwrite existing files; default is "a" for append) at some `level` like `logging.DEBUG` (default is `logging.NOTSET` so everything prints). Here is a script (in **examples**) that does that:

```
$ cat -n basic.py
     1  #!/usr/bin/env python3
     2
     3  import logging
     4  import os
     5  import sys
     6
     7  prg = sys.argv[0]
     8  prg_name, _ = os.path.splitext(os.path.basename(prg))
     9  logging.basicConfig(
    10      filename=prg_name + '.log',
    11      filemode='w',
    12      level=logging.DEBUG
    13  )
    14
    15  logging.debug('DEBUG!')
    16  logging.critical('CRITICAL!')
```

Before running the program, see that there is no log file:

```
$ ls
basic.py* long.py*
```

Run it, and see that `basic.log` has been created:

```
$ ls
```

```
basic.log  basic.py*  long.py*
$ cat basic.log
DEBUG:root:DEBUG!
CRITICAL:root:CRITICAL!
```

The key is to understand the hierarchy of the levels:

1. CRITICAL
2. ERROR
3. WARNING
4. INFO
5. DEBUG
6. NOTSET

The log level includes everything above the level you set. As in the above program, we set it to `logging.DEBUG` and so a call to `critical` was included. If you change the program to `logging.CRITICAL`, then `error` through `debug` calls are not emitted:

```
$ cat -n basic.py
     1  #!/usr/bin/env python3
     2
     3  import logging
     4  import os
     5  import sys
     6
     7  prg = sys.argv[0]
     8  prg_name, _ = os.path.splitext(os.path.basename(prg))
     9  logging.basicConfig(
    10      filename=prg_name + '.log',
    11      filemode='w',
    12      level=logging.CRITICAL
    13  )
    14
    15  logging.debug('DEBUG!')
    16  logging.critical('CRITICAL!')
$ ./basic.py
$ cat basic.log
CRITICAL:root:CRITICAL!
```

If you find yourself repeatedly debugging some program or just need to know information about how it is proceeding, then `logging` is for you. Maybe you have some functions or system calls that take a long time; sometimes you want to monitor how they are going and other times (e.g., running unattended on the HPC) you don't. Here is a program that logs random levels and then sleeps for one second. To see how this could be useful, open two terminals and navigate to the `examples` directory.

Here is the program:

```
$ cat -n long.py
     1  #!/usr/bin/env python3
     2
     3  import argparse
     4  import logging
     5  import os
     6  import random
     7  import sys
     8  import time
     9
    10
    11  # --------------------------------------------------
    12  def get_args():
    13      """get command-line arguments"""
    14      parser = argparse.ArgumentParser(
    15          description='Demonstrate logging',
    16          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    17
    18      parser.add_argument(
    19          '-d', '--debug', help='Debug mode', action='store_true')
    20
    21      return parser.parse_args()
    22
    23
    24  # --------------------------------------------------
    25  def main():
    26      """Make a jazz noise here"""
    27      args = get_args()
    28
    29      prg = sys.argv[0]
    30      prg_name, _ = os.path.splitext(os.path.basename(prg))
    31      logging.basicConfig(
    32          filename=prg_name + '.log',
    33          filemode='a',
    34          level=logging.DEBUG if args.debug else logging.CRITICAL)
    35
    36      logging.debug('Starting')
    37      for i in range(1, 11):
    38          method = random.choice([
    39              logging.info, logging.warning, logging.error, logging.critical,
    40              logging.debug
    41          ])
    42          method('{}: Hey!'.format(i))
    43          time.sleep(1)
    44
    45      logging.debug('Done')
```

```
46
47      print('Done.')
48
49
50  # ------------------------------------------------
51  if __name__ == '__main__':
52      main()
```

Start running `long.py` in one terminal, then execute `tail -f long.log` in the other where `tail` is the program to show you the end of a file and `-f` tells `tail` to stay running and "follow" the file as it grows. (Use CTRL-C to stop following.) Following is what I see when I run `long.py`. Note that, since I didn't set the `-d|--debug` flag, my program will only log *critical* errors:

```
CRITICAL:root:5: Hey!
CRITICAL:root:8: Hey!
```

And when I run `long.py -d`, everything from "debug" on up is displayed:

```
DEBUG:root:Starting
WARNING:root:1: Hey!
ERROR:root:2: Hey!
DEBUG:root:3: Hey!
DEBUG:root:4: Hey!
CRITICAL:root:5: Hey!
INFO:root:6: Hey!
ERROR:root:7: Hey!
INFO:root:8: Hey!
DEBUG:root:9: Hey!
CRITICAL:root:10: Hey!
DEBUG:root:Done
```

Chapter 11

SQLite in Python

SQLite (https://www.sqlite.org) is a lightweight, SQL/relational database that is available by default with Python (https://docs.python.org/3/library/sqlite3.html). By using `import sqlite3` you can interact with an SQLite database. So, let's create one, returning to our earlier Centrifuge output. Here is the file "tables.sql" containing the SQL statements needed to drop and create the tables:

```
drop table if exists tax;
create table tax (
    tax_id integer primary key,
    tax_name text not null,
    ncbi_id int not null,
    tax_rank text default '',
    genome_size int default 0,
    unique (ncbi_id)
);

drop table if exists sample;
create table sample (
    sample_id integer primary key,
    sample_name text not null,
    unique (sample_name)
);

drop table if exists sample_to_tax;
create table sample_to_tax (
    sample_to_tax_id integer primary key,
    sample_id int not null,
    tax_id int not null,
    num_reads int default 0,
    abundance real default 0,
    num_unique_reads integer default 0,
    unique (sample_id, tax_id),
    foreign key (sample_id) references sample (sample_id),
    foreign key (tax_id) references tax (tax_id)
);
```

Like Python, has data types of strings, integers, and floats (https://sqlite.org/datatype3.html). Primary keys are unique values defining a record in a table. You can place constraints on the allowed values of a field with conditions like `default` values or `not null` requirements as well as having the database enforce that some values are `unique` (such as NCBI taxonomy IDs). You can also require that

a particular combination of fields be unique, e.g., the sample/tax table has a unique constraint on the pairing of the sample/tax IDs. Additionally, this database uses foreign keys (https://sqlite.org/foreignkeys.html) to maintain relationships between tables. We will see in a moment how that prevents us from accidentally creating "orphan" records.

We are going to create a minimal database to track the abundance of species in various samples. The biggest rule of relational databases is to not repeat data. There should be one place to store each entity. For us, we have a "sample" (the Centrifuge ".tsv" file), a "taxonomy" (NCBI tax ID/name), and the relationship of the sample to the taxonomy. I have my own particular naming convention when it comes to relational tables/fields:

1. Name tables in the singular, e.g. "sample" not "samples"
2. Name the primary key [tablename] + underscore + "id", e.g., "sample_id"
3. Name linking tables [table1] + underscore + "to" + underscore + [table2]
4. Always have a primary key that is an auto-incremented integer

You can instantiate the database by calling `make db` in the "csv" directory to *first remove the existing database* and then recreate it by redirecting the "tables.sql" file into `sqlite3`:

```
$ make db
find . -name centrifuge.db -exec rm {} \;
sqlite3 centrifuge.db < tables.sql
```

You can then run `sqlite3 centrifuge.db` to use the CLI (command-line interface) to the database. Use `.help` inside SQLite to see all the "dot" commands (they begin with a ., cf. https://sqlite.org/cli.html):

```
$ sqlite3 centrifuge.db
SQLite version 3.13.0 2016-05-18 10:57:30
Enter ".help" for usage hints.
sqlite>
```

I often rely on the `.schema` command to look at the tables in an SQLite db. If you run that, you should see essentially the same thing as was in the "tables.sql" file. An alternate way to create the database is to use the `.read tables.sql` command from within SQLite to have it read and execute the SQL statements in that file.

We can manually insert a record into the `tax` table with an `insert` statement (https://sqlite.org/lang_insert.html). Note how SQLite treats strings and numbers exactly like Python – strings must be in quotes, numbers should be plain:

```
sqlite> insert into tax (tax_name, ncbi_id) values ('Homo sapiens', 3606);
```

We can add a dummy "sample" and link them like so:

```
sqlite> insert into sample (sample_name) values ('foo');
sqlite> insert into sample_to_tax (sample_id, tax_id, num_reads, abundance) values (1, 1, 1(
```

Verify that the data is there with a `select` statement (https://sqlite.org/lang_select.html):

```
sqlite> select count(*) from tax;
1
sqlite> select * from tax;
1|Homo sapiens|3606||0
```

Use `.headers on` to see the column names:

```
sqlite> .headers on
sqlite> select * from tax;
tax_id|tax_name|ncbi_id|tax_rank|genome_size
1|Homo sapiens|3606||0
sqlite> select * from sample;
sample_id|sample_name
1|foo
```

That's still a bit hard to read, so we can set `.mode column` to see a bit better:

```
sqlite> select * from sample;
sample_id   sample_name
----------  -----------
1           foo
sqlite> select * from tax;
tax_id      tax_name      ncbi_id     tax_rank    genome_size
----------  ------------  ----------  ----------  -----------
1           Homo sapiens  3606                    0
sqlite> select * from sample_to_tax;
sample_to_tax_id  sample_id   tax_id      num_reads   abundance   num_unique_reads
----------------  ----------  ----------  ----------  ----------  ----------------
1                 1           1           100         0.01        0
```

Often what we want is to `join` the tables so we can see just the data we want,
e.g., use this SQL:

```
select s.sample_name, t.tax_name, s2t.num_reads
from sample s, tax t, sample_to_tax s2t
where s.sample_id=s2t.sample_id
and s2t.tax_id=t.tax_id;
```

And you should see:

```
sample_name  tax_name      num_reads
-----------  ------------  ----------
foo          Homo sapiens  100
````
```

```
Now let's try to delete the `sample` record after we have turned on the enforcement of forei
```

````

```
sqlite> PRAGMA foreign_keys = ON;
sqlite> delete from sample where sample_id=1;
Error: FOREIGN KEY constraint failed
````
```

It would be bad to remove our sample and leave the sample/tax records in place. This is wha

Obviously we're not going to manually enter our data by hand, so let's write a script to imp

First we're going to need to get our data, so do `make data` to download some TSV files from

$ ./load_centrifuge.py *.tsv 1: Importing "YELLOWSTONE_SMPL_20717"
(2) Loading "Synechococcus sp. JA-3-3Ab" (321327) Loading "Synechococcus
sp. JA-2-3B'a(2-13)" (321332) 2: Importing "YELLOWSTONE_SMPL_20719"
(3) Loading "Streptococcus suis" (1307) Loading "synthetic construct" (32630)
3: Importing "YELLOWSTONE_SMPL_20721" (4) Loading "Staphylococcus
sp. AntiMn-1" (1715860) 4: Importing "YELLOWSTONE_SMPL_20723" (5)
5: Importing "YELLOWSTONE_SMPL_20725" (6) 6: Importing "YELLOW-
STONE_SMPL_20727" (7) Done """`

Here is the code that does that:

```
#!/usr/bin/env python3
"""Load Centrifuge into SQLite db"""

import argparse
import csv
import os
import re
import sqlite3
import sys


# --------------------------------------------------
def get_args():
    """get args"""
    parser = argparse.ArgumentParser(description='Load Centrifuge data')
    parser.add_argument('tsv_file', metavar='file',
                        help='Sample TSV file', nargs='+')
    parser.add_argument('-d', '--dbname', help='Centrifuge db name',
                        metavar='str', type=str, default='centrifuge.db')
    return parser.parse_args()
```

Our main is going to handle the arguments, ensuring the --dbname is a valid file,
then processing each of the tsv_file arguments (note the nargs declaration to
show that the program takes one or more TSV files). Note that in order to keep
this function short, I created two other functions, to import the samples and
TSV files:

```
# --------------------------------------------------
def main():
    """main"""
    args = get_args()
    tsv_files = args.tsv_file
    dbname = args.dbname

    if not os.path.isfile(dbname):
        print('Bad --dbname "{}"'.format(dbname))
        sys.exit(1)

    db = sqlite3.connect(dbname)

    for fnum, tsv_file in enumerate(tsv_files):
        if not os.path.isfile(tsv_file):
            print('Bad tsv_file "{}"'.format(tsv_file))
            sys.exit(1)

        sample_name, ext = os.path.splitext(tsv_file)

        if ext != '.tsv':
            print('"{}" does not end with ".tsv"'.format(tsv_file))
            sys.exit(1)

        if sample_name.endswith('.centrifuge'):
            sample_name = re.sub(r'\.centrifuge$', '', sample_name)

        sample_id = import_sample(sample_name, db)
        print('{:3}: Importing "{}" ({})'.format(fnum + 1,
                                                 sample_name, sample_id))
        import_tsv(db, tsv_file, sample_id)

    print('Done')
```

Here is the code to import a "sample." It needs a sample_name (which we assume to be unique) and a database handle (which is a bit like filehandles which we've been dealing with – it's the actual conduit from your code to the database). First we have to check if the sample already exists in our table, and this requires we use a `cursor` (https://docs.python.org/3/library/sqlite3.html) to issue our `select` statement. Rather than putting the sample name directly into the SQL (which is very insecure, see SQL injection/"Bobby Tables" XKCD https://xkcd.com/327), we use a `?` and pass the string as an argument to the `execute` function. If nothing (`None`) is returned, we can safely `insert` the new record and get the newly created sample ID from the `lastrowid` function of the cursor; otherwise, the sample ID is in the `res` result list as the first field:

```
# --------------------------------------------------
```

```
def import_sample(sample_name, db):
    """Import sample"""
    cur = db.cursor()
    cur.execute('select sample_id from sample where sample_name=?',
                (sample_name,))
    res = cur.fetchone()

    if res is None:
        cur.execute('insert into sample (sample_name) values (?)',
                    (sample_name,))
        sample_id = cur.lastrowid
    else:
        sample_id = res[0]

    return sample_id
```

The code to import the TSV file is similar. We establish SQL statements to
find/insert/update the sample/tax record, then we use the **csv** module to parse
the TSV file, creating dictionaries of each record (a product of merging the first
line/headers with each row of data). Again, to keep this function short enough
to fit on a "page," there is a separate function to find or create the taxonomy
record.

```
# --------------------------------------------------
def import_tsv(db, file, sample_id):
    """Import TSV file"""
    find_sql = """
        select sample_to_tax_id
        from   sample_to_tax
        where  sample_id=?
        and    tax_id=?
    """

    insert_sql = """
        insert
        into   sample_to_tax
               (sample_id, tax_id, num_reads, abundance, num_unique_reads)
        values (?, ?, ?, ?, ?)
    """

    update_sql = """
        update sample_to_tax
        set    sample_id=?, tax_id=?, num_reads=?,
               abundance=?, num_unique_reads=?
        where  sample_to_tax_id=?
    """
```

```
        cur = db.cursor()
        with open(file) as csvfile:
            reader = csv.DictReader(csvfile, delimiter='\t')
            for row in reader:
                tax_id = find_or_create_tax(db, row)
                if tax_id:
                    cur.execute(find_sql, (sample_id, tax_id))
                    res = cur.fetchone()
                    num_reads = row.get('numReads', 0)
                    abundance = row.get('abundance', 0)
                    num_uniq = row.get('numUniqueReads', 0)

                    if res is None:
                        cur.execute(insert_sql,
                                        (sample_id, tax_id, num_reads,
                                         abundance, num_uniq))
                    else:
                        s2t_id = res[0]
                        cur.execute(update_sql,
                                        (sample_id, tax_id, num_reads,
                                         abundance, num_uniq, s2t_id))
                else:
                    print('No tax id!')

            db.commit()

    return 1
```

The find/create tax function works just the same as that for the sample:

```
# --------------------------------------------------
def find_or_create_tax(db, rec):
    """find or create the tax"""
    find_sql = 'select tax_id from tax where ncbi_id=?'
    insert_sql = """
        insert into tax (tax_name, ncbi_id, tax_rank, genome_size)
        values (?, ?, ?, ?)
    """

    cur = db.cursor()
    ncbi_id = rec.get('taxID', '')
    if re.match('^\d+$', ncbi_id):
        cur.execute(find_sql, (ncbi_id,))
        res = cur.fetchone()

        if res is None:
            name = rec.get('name', '')
```

```
            if name:
                print('Loading "{}" ({})'.format(name, ncbi_id))
                cur.execute(insert_sql,
                            (name, ncbi_id, rec['taxRank'],
                             rec['genomeSize']))
                tax_id = cur.lastrowid
            else:
                print('No "name" in {}'.format(rec))
                return None
        else:
            tax_id = res[0]

        return tax_id
    else:
        print('"{}" does not look like an NCBI tax id'.format(ncbi_id))
        return None
```

If you use `make data`, several files will be downloaded from the iMicrobe FTP site for use by the `make load` step run the loader program:

```
$ make load
./load_centrifuge.py *.tsv
  1: Importing "YELLOWSTONE_SMPL_20717" (1)
Loading "Synechococcus sp. JA-3-3Ab" (321327)
Loading "Synechococcus sp. JA-2-3B'a(2-13)" (321332)
  2: Importing "YELLOWSTONE_SMPL_20719" (2)
Loading "Streptococcus suis" (1307)
Loading "synthetic construct" (32630)
  3: Importing "YELLOWSTONE_SMPL_20721" (3)
Loading "Staphylococcus sp. AntiMn-1" (1715860)
  4: Importing "YELLOWSTONE_SMPL_20723" (4)
  5: Importing "YELLOWSTONE_SMPL_20725" (5)
  6: Importing "YELLOWSTONE_SMPL_20727" (6)
Done
```

Now we can inspect how many records were loaded into the database:

```
$ sqlite3 centrifuge.db
SQLite version 3.13.0 2016-05-18 10:57:30
Enter ".help" for usage hints.
sqlite> select count(*) from tax;
5
sqlite> select count(*) from sample;
6
sqlite> select count(*) from sample_to_tax;
18
```

But, again, we're not going to just sit here and manually write SQL to

check out the data. Let's write a program that takes an NCBI tax id as an argument and reports the samples where it is found. You will need to make tabulate to run the command to install the "tabulate" module (https://pypi.python.org/pypi/tabulate) in order to run this program:

```python
1   #!/usr/bin/env python3
2   """Query centrifuge.db for NCBI tax id"""
3
4   import argparse
5   import os
6   import re
7   import sys
8   import sqlite3
9   from tabulate import tabulate
10
11  # --------------------------------------------------
12  def get_args():
13      """get args"""
14      parser = argparse.ArgumentParser(description='Argparse Python script')
15      parser.add_argument('-d', '--dbname', help='Centrifuge db name',
16                          metavar='str', type=str, default='centrifuge.db')
17      parser.add_argument('-o', '--orderby', help='Order by',
18                          metavar='str', type=str, default='abundance')
19      parser.add_argument('-s', '--sortorder', help='Sort order',
20                          metavar='str', type=str, default='desc')
21      parser.add_argument('-t', '--taxid', help='NCBI taxonomy id',
22                          metavar='str', type=str, required=True)
23      return parser.parse_args()
24
25  # --------------------------------------------------
26  def main():
27      """main"""
28      args = get_args()
29      dbname = args.dbname
30      order_by = args.orderby
31      sort_order = args.sortorder
32
33      if not os.path.isfile(dbname):
34          print('"{}" is not a valid file'.format(dbname))
35          sys.exit(1)
36
37      flds = set(['tax_name', 'num_reads', 'abundance', 'sample_name'])
38      if not order_by in flds:
39          print('"{}" not an allowed --orderby, choose from {}'.format(
40              order_by, ', '.join(flds)))
41          sys.exit(1)
```

```
42
43      sorting = set(['asc', 'desc'])
44      if not sort_order in sorting:
45          print('"{}" not an allowed --sortorder, choose from {}'.format(
46              order_by, ', '.join(sorting)))
47          sys.exit(1)
48
49      tax_ids = []
50      for tax_id in re.split(r'\s*,\s*', args.taxid):
51          if re.match(r'^\d+$', tax_id):
52              tax_ids.append(tax_id)
53          else:
54              print('"{}" does not look like an NCBI tax id'.format(tax_id))
55
56      if len(tax_ids) == 0:
57          print('No tax ids')
58          sys.exit(1)
59
60      db = sqlite3.connect(dbname)
61      cur = db.cursor()
62      sql = """
63          select    s.sample_name, t.tax_name, s2t.num_reads, s2t.abundance
64          from      sample s, tax t, sample_to_tax s2t
65          where     s.sample_id=s2t.sample_id
66          and       s2t.tax_id=t.tax_id
67          and       t.ncbi_id in ({})
68          order by {} {}
69      """.format(', '.join(tax_ids), order_by, sort_order)
70
71      cur.execute(sql)
72
73      samples = cur.fetchall()
74      if len(samples) > 0:
75          cols = [d[0] for d in cur.description]
76          print(tabulate(samples, headers=cols))
77      else:
78          print('No results')
79
80  # --------------------------------------------------
81  if __name__ == '__main__':
82      main()
```

It takes as arguments a required NCBI tax id that can be a single value or a comma-separated list. Options include the SQLite Centrifuge db, a column name to sort by, and whether to show in ascending or descending order. The output is formatted with the `tabulate` module to produce a simple text table. To query

by one tax ID:

```
$ ./query_centrifuge.py -t 321327
sample_name            tax_name                     num_reads    abundance
---------------------  --------------------------   -----------  -----------
YELLOWSTONE_SMPL_20721 Synechococcus sp. JA-3-3Ab          315         0.98
YELLOWSTONE_SMPL_20723 Synechococcus sp. JA-3-3Ab         6432         0.98
YELLOWSTONE_SMPL_20727 Synechococcus sp. JA-3-3Ab         1219         0.96
YELLOWSTONE_SMPL_20717 Synechococcus sp. JA-3-3Ab           19         0.53
YELLOWSTONE_SMPL_20719 Synechococcus sp. JA-3-3Ab          719         0.27
YELLOWSTONE_SMPL_20725 Synechococcus sp. JA-3-3Ab         3781         0.2
```

To query by more than one:

```
$ ./query_centrifuge.py -t 321327,1307
sample_name            tax_name                     num_reads    abundance
---------------------  --------------------------   -----------  -----------
YELLOWSTONE_SMPL_20721 Synechococcus sp. JA-3-3Ab          315         0.98
YELLOWSTONE_SMPL_20723 Synechococcus sp. JA-3-3Ab         6432         0.98
YELLOWSTONE_SMPL_20727 Synechococcus sp. JA-3-3Ab         1219         0.96
YELLOWSTONE_SMPL_20717 Synechococcus sp. JA-3-3Ab           19         0.53
YELLOWSTONE_SMPL_20719 Synechococcus sp. JA-3-3Ab          719         0.27
YELLOWSTONE_SMPL_20725 Synechococcus sp. JA-3-3Ab         3781         0.2
YELLOWSTONE_SMPL_20719 Streptococcus suis                    1         0
```

To order by "num_reads" instead of "abundance":

```
$ ./query_centrifuge.py -t 321327,1307 -o num_reads
sample_name            tax_name                     num_reads    abundance
---------------------  --------------------------   -----------  -----------
YELLOWSTONE_SMPL_20723 Synechococcus sp. JA-3-3Ab         6432         0.98
YELLOWSTONE_SMPL_20725 Synechococcus sp. JA-3-3Ab         3781         0.2
YELLOWSTONE_SMPL_20727 Synechococcus sp. JA-3-3Ab         1219         0.96
YELLOWSTONE_SMPL_20719 Synechococcus sp. JA-3-3Ab          719         0.27
YELLOWSTONE_SMPL_20721 Synechococcus sp. JA-3-3Ab          315         0.98
YELLOWSTONE_SMPL_20717 Synechococcus sp. JA-3-3Ab           19         0.53
YELLOWSTONE_SMPL_20719 Streptococcus suis                    1         0
```

To sort ascending:

```
$ ./query_centrifuge.py -t 321327,1307 -o num_reads -s asc
sample_name            tax_name                     num_reads    abundance
---------------------  --------------------------   -----------  -----------
YELLOWSTONE_SMPL_20719 Streptococcus suis                    1         0
YELLOWSTONE_SMPL_20717 Synechococcus sp. JA-3-3Ab           19         0.53
YELLOWSTONE_SMPL_20721 Synechococcus sp. JA-3-3Ab          315         0.98
YELLOWSTONE_SMPL_20719 Synechococcus sp. JA-3-3Ab          719         0.27
YELLOWSTONE_SMPL_20727 Synechococcus sp. JA-3-3Ab         1219         0.96
YELLOWSTONE_SMPL_20725 Synechococcus sp. JA-3-3Ab         3781         0.2
```

YELLOWSTONE_SMPL_20723 Synechococcus sp. JA-3-3Ab   6432   0.98

Chapter 12

Writing Tests For Your Python Programs

> Much of the essence of building a program is in fact the debugging
> of the specification. – Fred Brooks

Let's start with a simple example of a `hello` program that should say "Hello,
name!"

```
$ ./hello.py
Usage: hello.py NAME [NAME...]
$ ./hello.py Jan
Hello, Jan!
$ ./hello.py Bobby Peter Greg
Hello, Bobby!
Hello, Peter!
Hello, Greg!
```

Here is one way to write such a program *with an embedded test*:

```
 1  #!/usr/bin/env python3
 2
 3  import os
 4  import sys
 5
 6  def hello(name):
 7      return 'Hello, {}!'.format(name)
 8
 9  def test_hello():
10      assert hello('World') == 'Hello, World!'
11      assert hello('') == 'Hello, !'
12      assert hello('my name is Fred') == 'Hello, my name is Fred!'
13
14  def main():
15      args = sys.argv[1:]
16      if not args:
17          print('Usage: {} NAME [NAME...]'.format(os.path.basename(sys.argv[0])))
18          sys.exit(1)
19
20      for arg in args:
21          print(hello(arg))
22
23  if __name__ == '__main__':
24      main()
```

Specifically I've written this to use the PyTest (https://docs.pytest.org/en/latest/)

framework that will search for function names starting with `test_` and will run them.

```
$ pytest -v hello.py
============================== test session starts ===============================
platform darwin -- Python 3.6.8, pytest-4.2.0, py-1.7.0, pluggy-0.8.1 -- /anaconda3/bin/pyth
cachedir: .pytest_cache
rootdir: /Users/kyclark/work/biosys-analytics/lectures/18-writing-tests/examples, inifile:
plugins: remotedata-0.3.1, openfiles-0.3.2, doctestplus-0.2.0, arraydiff-0.3
collected 1 item

hello.py::test_hello PASSED                                              [100%]

=========================== 1 passed in 0.03 seconds ===========================
```

The `assert` function you see in the `test_hello` function is a built-in Python function that evaluates some predicate and will throw an error if the predicate is false. For instance, we `assert` that `hello("World")` should return the string `Hello, World!`. If this does not happen, the test will fail:

```
>>> def hello(name):
...     return 'Hello, {}!'.format(name)
...
>>> assert hello('World') == 'foo'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AssertionError
```

PyTest will find such errors and report them as failed tests:

```
$ cat -n hello_bad.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
     5
     6  def hello(name):
     7      return 'Hello, {}!'.format(name)
     8
     9  def test_hello():
    10      assert hello('World') == 'Hello, World.'
    11
    12  def main():
    13      args = sys.argv[1:]
    14      if not args:
    15          print('Usage: {} NAME [NAME...]'.format(os.path.basename(sys.argv[0])))
    16          sys.exit(1)
    17
```

```
    18       for arg in args:
    19           print(hello(arg))
    20
    21   if __name__ == '__main__':
    22       main()
$ pytest -v hello_bad.py
============================ test session starts ============================
platform darwin -- Python 3.6.8, pytest-4.2.0, py-1.7.0, pluggy-0.8.1 -- /anaconda3/bin/pytl
cachedir: .pytest_cache
rootdir: /Users/kyclark/work/biosys-analytics/lectures/18-writing-tests/examples, inifile:
plugins: remotedata-0.3.1, openfiles-0.3.2, doctestplus-0.2.0, arraydiff-0.3
collected 1 item

hello_bad.py::test_hello FAILED                                        [100%]

================================= FAILURES =================================
_____ test_hello _____

    def test_hello():
>       assert hello('World') == 'Hello, World.'
E       AssertionError: assert 'Hello, World!' == 'Hello, World.'
E         - Hello, World!
E         ?             ^
E         + Hello, World.
E         ?             ^

hello_bad.py:10: AssertionError
========================== 1 failed in 0.08 seconds ==========================
```

The error output highlights the differences between what was expected (`Hello,
World.` ending in a period) and what the `hello` function actually returned
(`Hello, World!` ending in an exclamation point).

I would recommend writing your tests for every function directly below the
function being tested and calling the test `test_function`. Try to make a
function do just one thing, then write tests to ensure it does that thing. Try
to write tests that probe the edge cases, e.g., passing an empty string or a
very long string. Here's a version where the `hello` function will only greet if
the argument is a `str`; otherwise it will return an admonishment. This is an
extremely contrived example because everything coming in via `sys.argv` is by
definition a string, so I will intentionally convert anything that looks like a digit
to an `int` so that we can see the error:

```
$ cat -n hello_fail.py
     1   #!/usr/bin/env python3
     2
     3   import os
```

```
 4  import sys
 5
 6  def hello(name):
 7      if type(name) is str:
 8          return 'Hello, {}!'.format(name)
 9      else:
10          return 'Can only say hello to a string'
11
12
13  def test_hello():
14      assert hello('World') == 'Hello, World!'
15      assert hello('') == 'Hello, !'
16      assert hello('my name is Fred') == 'Hello, my name is Fred!'
17
18      err = 'Can only say hello to a string'
19      assert hello(4) == err
20      assert hello(None) == err
21      assert hello(float) == err
22      assert hello(str) == err
23
24  def main():
25      args = sys.argv[1:]
26      if not args:
27          print('Usage: {} NAME [NAME...]'.format(os.path.basename(sys.argv[0])))
28          sys.exit(1)
29
30      for arg in args:
31          if arg.isdigit(): arg = int(arg)
32
33          print(hello(arg))
34
35  if __name__ == '__main__':
36      main()
```

```
$ ./hello_fail.py Bob 3 Sue
Hello, Bob!
Can only say hello to a string
Hello, Sue!
$ pytest -v hello_fail.py
=========================== test session starts ============================
platform darwin -- Python 3.6.8, pytest-4.2.0, py-1.7.0, pluggy-0.8.1 -- /anaconda3/bin/pyth
cachedir: .pytest_cache
rootdir: /Users/kyclark/work/biosys-analytics/lectures/18-writing-tests/examples, inifile:
plugins: remotedata-0.3.1, openfiles-0.3.2, doctestplus-0.2.0, arraydiff-0.3
collected 1 item

hello_fail.py::test_hello PASSED                                       [100%]
```

```
=========================== 1 passed in 0.04 seconds ===========================
```

These types of tests that live *inside* each of your source files and test invididual functions are known as "unit tests". As your software grows, you may find yourself breaking your functions into logically grouped files or modules. We can also write tests that live *outside* our program files to ensure the proper integration of modules as well as the user interface we present. All the `test.py` programs that have been included in your assignments have these types of tests – ensuring, for instance, that your program will create a "usage" statement if passed no arguments or `-h|--help`, will print message to STDERR and `sys.exit()` with a non-zero value when there is an error, or will run to completion given good input and produce the expected STDOUT and/or output files.

Here is a `test.py` that tests for the usage statement and error code on no input and then tests for one, argument, more than one argument, and an argument that is more than one word:

```
$ cat -n test.py
     1  #!/usr/bin/env python3
     2
     3  from subprocess import getstatusoutput
     4
     5  prg = './hello.py'
     6
     7  def test_usage():
     8      rv, out = getstatusoutput('{}'.format(prg))
     9      assert rv != 0
    10      assert out.lower().startswith('usage')
    11
    12  def test_runs_ok():
    13      rv1, out1 = getstatusoutput('{} Carl'.format(prg))
    14      assert rv1 == 0
    15      assert out1 == 'Hello, Carl!'
    16
    17      rv2, out2 = getstatusoutput('{} Barbara McClintock'.format(prg))
    18      assert rv2 == 0
    19      assert out2 == 'Hello, Barbara!\nHello, McClintock!'
    20
    21      rv3, out3 = getstatusoutput('{} "Barbara McClintock"'.format(prg))
    22      assert rv3 == 0
    23      assert out3 == 'Hello, Barbara McClintock!'
```

I typically create a `Makefile` with a `test` target to show users how to run the tests:

```
$ cat -n Makefile
     1  .PHONY: test
```

```
     2
     3  test:
     4      pytest -v test.py
$ make test
pytest -v test.py
============================ test session starts ===============================
platform darwin -- Python 3.6.8, pytest-4.2.0, py-1.7.0, pluggy-0.8.1 -- /anaconda3/bin/pyth
cachedir: .pytest_cache
rootdir: /Users/kyclark/work/biosys-analytics/lectures/18-writing-tests/examples, inifile:
plugins: remotedata-0.3.1, openfiles-0.3.2, doctestplus-0.2.0, arraydiff-0.3
collected 2 items

test.py::test_usage PASSED                                                [ 50%]
test.py::test_runs_ok PASSED                                              [100%]

=========================== 2 passed in 0.19 seconds ===========================
```

Chapter 13

Functional Ideas in Python

"Gematria" is a system for assigning a number to a word by summing the numeric values of each of the letters as defined by the Mispar godol (https://en.wikipedia.org/wiki/Gematria). For English characters, we can use the ASCII table (https://en.wikipedia.org/wiki/ASCII). It is not necessary, however, to encode this table in our program as Python provides the `ord` function to convert a character to its "ordinal" (order in the ASCII table) value as well as the `chr` function to convert a number to its "character."

```python
print('"{}" = "{}"'.format('A', ord('A')))
```

```
"A" = "65"
```

```python
print('"{}" = "{}"'.format('a', ord('a')))
```

```
"a" = "97"
```

```python
print('"{}" = "{}"'.format(88, chr(88)))
```

```
"88" = "X"
```

```python
print('"{}" = "{}"'.format(112, chr(112)))
```

```
"112" = "p"
```

To implement an ASCII version of gematria in Python, we need to turn each letter into a number and add them all together. So, to start, note that Python can use a `for` loop to cycle through all the members of a list (in order):

```python
for n in range(5):
    print(n)
```

```
0
1
2
3
4
```

```python
for char in ['p', 'y', 't', 'h', 'o', 'n']:
    print(char)
```

```
p
y
t
h
o
n
```

A "word" is simply a list of characters, so we can iterate over it just like a list of numbers:

```python
for char in "python":
    print(char)
```

```
p
y
t
h
o
n
```

Let's print the ordinal (ASCII) value instead:

```python
for char in "python":
    print('"{}" = "{}"'.format(char, ord(char)))
```

```
"p" = "112"
"y" = "121"
"t" = "116"
"h" = "104"
"o" = "111"
"n" = "110"
```

Now let's create a variable to hold the running sum of the values:

```python
word = "python"
total = 0
for char in word:
    total += ord(char)

print('"{}" = "{}"'.format(word, total))
```

```
"python" = "674"
```

Another way could be to create another list to hold the values and then use the sum function:

```python
word = "python"
all = []
for char in word:
    all.append(ord(char))

print(all)
print('"{}" = "{}"'.format(word, sum(all)))
```

```
[112, 121, 116, 104, 111, 110]
"python" = "674"
```

Map

We can use a `map` function to transform all the characters via the `ord` function. This is interesting because `map` is a function that takes another function as its first argument. The second is a list of items to feed into the function. The result is the transformed list. For instance, we can use the `str.upper` function to turn each letter (e.g., "p") into the upper-case version ("P"). NB: it's necessary to force the results into a `list`.

```python
list(map(str.upper, "python"))
```

```
['P', 'Y', 'T', 'H', 'O', 'N']
```

```python
list(map(ord, "python"))
```

```
[112, 121, 116, 104, 111, 110]
```

Now we can `sum` those numbers:

```python
sum(map(ord, "python"))
```

```
674
```

Now let's think about how we could apply this to all the words in a file. As above, we can use a `for` loop to iterate over all the lines in a file:

```python
for line in open('gettysburg.txt'):
    print(line)
```

```
Four score and seven years ago our fathers brought forth on this

continent, a new nation, conceived in Liberty, and dedicated to the

proposition that all men are created equal.
```

The original is single-spaced, so why is this printing double-spaced? The `for` loop reads each "line" which is a string of text up to and including a newline. The `print` by default adds a newline, so we either need to `print(line, end='')` to indicate we don't want anything at the end:

```python
for line in open('gettysburg.txt'):
    print(line, end='')
```

```
Four score and seven years ago our fathers brought forth on this
continent, a new nation, conceived in Liberty, and dedicated to the
proposition that all men are created equal.
```

Or we need to use the `rstrip` function to "strip" whitespace off the "r"ight side of the line:

```python
for line in open('gettysburg.txt'):
    print(line.rstrip())
```

```
Four score and seven years ago our fathers brought forth on this
continent, a new nation, conceived in Liberty, and dedicated to the
proposition that all men are created equal.
```

We can use the `split` function to get all the words for each line and a `for` loop
to iterate over those:

```python
for line in open('gettysburg.txt'):
    for word in line.rstrip().split():
        print(word)
```

```
Four
score
and
seven
years
ago
our
fathers
brought
forth
on
this
continent,
a
new
nation,
conceived
in
Liberty,
and
dedicated
to
the
proposition
that
all
men
are
created
equal.
```

We want to get rid of anything that is not character like the punctuation. There
is a function in the `str` library called `isalpha` that returns `True` or `False`:

```python
for char in "a8,X.b!G":
    print('"{}" = "{}"'.format(char, str.isalpha(char)))
```

```
"a" = "True"
```

```
"8" = "False"
"," = "False"
"X" = "True"
"." = "False"
"b" = "True"
"!" = "False"
"G" = "True"
```

Each `char` in the loop is itself a string, so we can call the method directly on the variable:

```python
for char in "a8,X.b!G":
    print('"{}" = "{}"'.format(char, char.isalpha()))
```

```
"a" = "True"
"8" = "False"
"," = "False"
"X" = "True"
"." = "False"
"b" = "True"
"!" = "False"
"G" = "True"
```

Filter

Similar to what we saw above with the `map` function, we can use `filter` to find all the characters in a string which are `True` for `isalpha`. `filter` is another "higher-order function" that takes another function for its first argument (called the "predicate") and a list as the second argument. Whereas `map` returns *all* the elements of the list transformed by the function, `filter` returns *only those for which the predicate is true*.

```python
list(filter(str.isalpha, "a8,X.b!G"))
```

```
['a', 'X', 'b', 'G']
```

The first argument for `map` and `filter` is called the "lambda," and sometimes you will see it written out explicitly like so:

```python
list(filter(lambda char: char.isalpha(), "a8,X.b!G"))
```

```
['a', 'X', 'b', 'G']
```

Here is a way to find only even numbers:

```python
list(filter(lambda x: x % 2 == 0, range(10)))
```

```
[0, 2, 4, 6, 8]
```

Let's turn that list of characters back into a word with the `join` function:

```
''.join(filter(str.isalpha, "a8,X.b!G"))
```

```
'aXbG'
```

Aside (Regular expressions)

NB: This is not the way I would actually remove punctuation in my own code. I'd be more likely to use regular expressions, e.g., "anything not A-Z, a-z, and 0-9":

```
import re
print(re.sub('[^A-Za-z0-9]', '', 'a8,X.b!G'))
```

```
a8XbG
```

The `string` class actually defines "punctuation":

```
import string
print(string.punctuation)
```

```
!"#$%&'()*+,-./:;<=>?@[\]^_`{|}~
```

So we could use that to create a character class of punctuation if that was the only thing we intended to remove:

```
import string
print(re.sub('[' + string.punctuation + ']', '', 'a8,X.b!G'))
```

```
a8XbG
```

Combining map and filter

So, going back to our Gettysburg example, here is a list of all the words without punctuation:

```
for line in open('gettysburg.txt'):
    for word in line.rstrip().split():
        print(''.join(filter(str.isalpha, word)))
```

```
Four
score
and
seven
years
ago
our
fathers
brought
```

```
forth
on
this
continent
a
new
nation
conceived
in
Liberty
and
dedicated
to
the
proposition
that
all
men
are
created
equal
```

Now, rather let's print the **sum** of the **chr** values for each cleaned up word:

```python
for line in open('gettysburg.txt'):
    for word in line.rstrip().split():
        clean = ''.join(filter(str.isalpha, word))
        print('"{}" = "{}"'.format(clean, sum(map(ord, clean))))
```

```
"Four" = "412"
"score" = "540"
"and" = "307"
"seven" = "545"
"years" = "548"
"ago" = "311"
"our" = "342"
"fathers" = "749"
"brought" = "763"
"forth" = "547"
"on" = "221"
"this" = "440"
"continent" = "978"
"a" = "97"
"new" = "330"
"nation" = "649"
"conceived" = "944"
"in" = "215"
```

```
"Liberty" = "731"
"and" = "307"
"dedicated" = "919"
"to" = "227"
"the" = "321"
"proposition" = "1222"
"that" = "433"
"all" = "313"
"men" = "320"
"are" = "312"
"created" = "728"
"equal" = "536"
```

Notice that we are calling `rstrip` for every line, so we could easily move that
into a `map`, and the "cleaning" code can likewise be moved into a `map`:

```python
for line in map(str.rstrip, open('gettysburg.txt')):
    for word in map(lambda w: ''.join(filter(str.isalpha, w)), line.split()):
        print('"{}" = "{}"'.format(word, sum(map(ord, word))))
```

```
"Four" = "412"
"score" = "540"
"and" = "307"
"seven" = "545"
"years" = "548"
"ago" = "311"
"our" = "342"
"fathers" = "749"
"brought" = "763"
"forth" = "547"
"on" = "221"
"this" = "440"
"continent" = "978"
"a" = "97"
"new" = "330"
"nation" = "649"
"conceived" = "944"
"in" = "215"
"Liberty" = "731"
"and" = "307"
"dedicated" = "919"
"to" = "227"
"the" = "321"
"proposition" = "1222"
"that" = "433"
"all" = "313"
"men" = "320"
```

```
"are" = "312"
"created" = "728"
"equal" = "536"
```

At this point, we have arguably sacrificed readability for the sake of using `map` and `filter` – another instance of "just because you can doesn't mean you should!"

We can improve readability, however, by creating our own functions with informative names. Also, since `onlychars` will get rid of the trailing newlines, we can remove the `line.rstrip()` call:

```python
def onlychars(word):
    return ''.join(filter(str.isalpha, word))

def word2num(word):
    return sum(map(ord, word))

for line in open('gettysburg.txt'):
    for word in map(onlychars, line.split()):
        print('"{}" = "{}"'.format(word, word2num(word)))
```

```
"Four" = "412"
"score" = "540"
"and" = "307"
"seven" = "545"
"years" = "548"
"ago" = "311"
"our" = "342"
"fathers" = "749"
"brought" = "763"
"forth" = "547"
"on" = "221"
"this" = "440"
"continent" = "978"
"a" = "97"
"new" = "330"
"nation" = "649"
"conceived" = "944"
"in" = "215"
"Liberty" = "731"
"and" = "307"
"dedicated" = "919"
"to" = "227"
"the" = "321"
"proposition" = "1222"
"that" = "433"
"all" = "313"
"men" = "320"
```

```
"are" = "312"
"created" = "728"
"equal" = "536"
```

Golfing

"Golfing" in code is when you try to express your code in fewer and fewer keystrokes. At some point you cross the line of cleanliness to absudity. Remember:

> It's such a fine line between stupid and clever. – David St. Hubbins (https://www.youtube.com/watch?v=wtXkD1BC564)

Here's a streamlined version that combines `open`, `read`, and `split` to read the entire file into a list of words which are `map`d into `word2num`.

NB: This version assumes you have enough memory to read an *entire file* and split it. The versions above which read and process each line consume only as much memory as any one line needs!

```
def onlychars(word):
    return ''.join(filter(str.isalpha, word))

def word2num(word):
    return str(sum(map(ord, word)))

print(' '.join(map(word2num,
                   map(onlychars,
                       open('gettysburg.txt').read().split())))))
```

412 540 307 545 548 311 342 749 763 547 221 440 978 97 330 649 944 215 731 307 919 227 321 1

To mimic the above output:

```
def onlychars(word):
    return ''.join(filter(str.isalpha, word))

def word2num(word):
    return str(sum(map(ord, onlychars(word))))

print('\n'.join(map(lambda word: '"{}" = "{}"'.format(word, word2num(word)),
                    map(onlychars,
                        open('gettysburg.txt').read().split())))))
```

```
"Four" = "412"
"score" = "540"
"and" = "307"
"seven" = "545"
"years" = "548"
```

```
"ago" = "311"
"our" = "342"
"fathers" = "749"
"brought" = "763"
"forth" = "547"
"on" = "221"
"this" = "440"
"continent" = "978"
"a" = "97"
"new" = "330"
"nation" = "649"
"conceived" = "944"
"in" = "215"
"Liberty" = "731"
"and" = "307"
"dedicated" = "919"
"to" = "227"
"the" = "321"
"proposition" = "1222"
"that" = "433"
"all" = "313"
"men" = "320"
"are" = "312"
"created" = "728"
"equal" = "536"
```

With this, I hope you're now understand what is meant by a "higher-order function" (functions that take other functions as arguments) and how they can streamline your code.

Exercise

Read your local dictionary (e.g., "/usr/share/dict/words") and find how many words share the same numeric representation. Which ones have the value "666"?

```python
from collections import defaultdict

def onlychars(word):
    return ''.join(filter(str.isalpha, word))

file = '/usr/share/dict/words'
num2word = defaultdict(list)
for line in map(str.rstrip, open(file)):
    for word in map(onlychars, line.split()):
        num = sum(map(ord, word))
        num2word[num].append(word)
```

```python
satan = '666'
if satan in num2word:
    print('Satan =', num2word[satan])
else:
    print('No Satan!')

count_per_n = []
for n, wordlist in num2word.items():
    count_per_n.append((len(wordlist), n))

top10 = list(reversed(sorted(count_per_n)))[:10]
for num_of_words, n in top10:
    print('{} ({} words) = {} ...'.format(n, len(num2word[n]), ', '.join(num2word[n][:3])))

No Satan!
973 (623 words) = Actaeaceae, activator, actorship ...
969 (621 words) = abrotanum, acclivous, acidulous ...
965 (611 words) = abhorrent, acoumeter, acronymic ...
855 (607 words) = abuseful, acanthus, acronych ...
861 (601 words) = Absyrtus, acaulous, adjuvant ...
856 (597 words) = abrastol, accismus, acervose ...
971 (596 words) = aburabozu, acropathy, acuteness ...
974 (594 words) = ablastous, absolvent, abysmally ...
972 (592 words) = accessory, acropolis, acutiator ...
1078 (587 words) = absentness, acrogenous, actinozoan ...
```

Chapter 14

Writing Pipelines in Python

> Falling in love with code means falling in love with problem solving
> and being a part of a forever ongoing conversation. – Kathryn Barrett

You might be surprised at how far you can push humble `make` to write analysis
pipelines. I'd encourage you to really explore Makefiles, reading the docs and
looking at other people's examples. You'll save yourself many hours if you
learn to use `make` well, even if you are just documenting how you ran your
Python program. Beyond `make`, there are many other frameworks for writing
pipelines such as Nextflow, Snakemake, Taverna, Pegasus and many more (cf
https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5429012/), many of which are
probably far superior to rolling your own in Python; however, we will do just
that as you will learn many valuable skills along the way. After all, hubris is one
of the three virtues of a great programmer:

> According to Larry Wall, the original author of the Perl programming
> language, there are three great virtues of a programmer:
>
> **Laziness**: The quality that makes you go to great effort to reduce
> overall energy expenditure. It makes you write labor-saving programs
> that other people will find useful and document what you wrote so
> you don't have to answer so many questions about it.
>
> **Impatience**: The anger you feel when the computer is being lazy.
> This makes you write programs that don't just react to your needs,
> but actually anticipate them. Or at least pretend to.
>
> **Hubris**: The quality that makes you write (and maintain) programs
> that other people won't want to say bad things about.
>
> *Programming Perl*, 2nd Edition, O'Reilly & Associates, 1996

A "pipeline" is chaining the output of one program or function as the input to
the next as many times as necessary to arrive at an end product. Sometimes the
whole pipeline can be written inside Python, but often in bioinformatics what
we have is one program written in Java/C/C++ we install from source that
creates some output that needs to be massaged by a program we write in bash
or Python that gets fed to a Perl script you found on BioStars that produces
some text file that we read into R to create some visualization. We're going to
focus on how to use Python to take input, call external programs, check on the
status, and feed the output to some other program.

Hello

In this first example, we'll pretend this "hello.sh" is something more interesting than it really is:

```
$ cat -n hello.sh
     1  #!/usr/bin/env bash
     2
     3  if [[ $# -lt 1 ]]; then
     4      printf "Usage: %s NAME\n" $(basename $0)
     5      exit 1
     6  fi
     7
     8  NAME=$1
     9
    10  if [[ $NAME == 'Lord Voldemort' ]]; then
    11      echo "Upon advice of my counsel, I respectfully refuse to say that name."
    12      exit 1
    13  fi
    14
    15  echo "Hello, $1!"
$ ./hello.sh
Usage: hello.sh NAME
$ ./hello.sh Jan
Hello, Jan!
$ ./hello.sh "Lord Voldemort"
Upon advice of my counsel, I respectfully refuse to say that name.
```

We'll write a Python program to feed names to the "hello.sh" program and monitor whether the program ran successfully.

```
$ cat -n run_hello.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date   : 2019-03-28
     5  Purpose: Run "hello.sh"
     6  """
     7
     8  import argparse
     9  import os
    10  import sys
    11  from subprocess import getstatusoutput
    12
    13
    14  # --------------------------------------------------
    15  def get_args():
```

```
16        """get command-line arguments"""
17        parser = argparse.ArgumentParser(
18            description='Simple pipeline',
19            formatter_class=argparse.ArgumentDefaultsHelpFormatter)
20
21        parser.add_argument(
22            'name', metavar='str', nargs='+', help='Names for hello.sh')
23
24        parser.add_argument(
25            '-p',
26            '--program',
27            help='Program to run',
28            metavar='str',
29            type=str,
30            default='./hello.sh')
31
32        return parser.parse_args()
33
34
35    # --------------------------------------------------
36    def warn(msg):
37        """Print a message to STDERR"""
38        print(msg, file=sys.stderr)
39
40
41    # --------------------------------------------------
42    def die(msg='Something bad happened'):
43        """warn() and exit with error"""
44        warn(msg)
45        sys.exit(1)
46
47
48    # --------------------------------------------------
49    def main():
50        """Make a jazz noise here"""
51        args = get_args()
52        prg = args.program
53
54        if not os.path.isfile(prg):
55            die('Missing expected program "{}"'.format(prg))
56
57        for name in args.name:
58            cmd = '{} "{}"'.format(prg, name)
59            rv, out = getstatusoutput(cmd)
60            if rv != 0:
61                warn('Failed to run: {}\nError: {}'.format(cmd, out))
```

```
62            else:
63                print('Success: "{}"'.format(out))
64
65        print('Done.')
66
67
68    # -------------------------------------------------
69    if __name__ == '__main__':
70        main()
```

In `get_args` we establish that we expect one or more positional arguments on the command line along with an optional `-p|--program` to run with those as arguments. One of the first items to check is if the `program` exists (we are expecting a full path with `./hello.sh` being the default), so line 54 checks this and calls `die` if it does not exist.

The main event starts on line 57 where we loop through the name arguments. On line 58, we create a command by making a string with the name of the program and the argument. Then we use `subprocess.getstatusoutput` to run this command and give us the return value (`rv`) and the output from the command (both STDERR and STDOUT get combined). If the return value is not zero ("zero errors"), then we use `warn` to report on STDERR that there was a failure, else we print "Success" along with the output from `hello.sh`.

If we run this, we see it stops when given a bad `program`:

```
$ ./run_hello.py -p foo Ken
Missing expected program "foo"
```

And we see it correctly reports the results for our inputs:

```
$ ./run_hello.py Jan Marcia "Lord Voldemort" Cindy
Success: "Hello, Jan!"
Success: "Hello, Marcia!"
Failed to run: ./hello.sh "Lord Voldemort"
Error: Upon advice of my counsel, I respectfully refuse to say that name.
Success: "Hello, Cindy!"
Done.
```

If you were submitting this job to run on an HPC, it would be launched by the job scheduler sometime later than when you submit it and would be run in an automated fashion. You would quickly learn that it's better to capture errors to an error file rather than let them comingle with STDOUT.

```
$ ./run_hello.py Jan Marcia "Lord Voldemort" Cindy 2>err
Success: "Hello, Jan!"
Success: "Hello, Marcia!"
Success: "Hello, Cindy!"
Done.
$ cat err
```

```
Failed to run: ./hello.sh "Lord Voldemort"
Error: Upon advice of my counsel, I respectfully refuse to say that name.
```

Parallel Hello

This works fairly well, but what if there are potentially dozens, hundreds, or thousands of names to greet? We are processing these in a serial fashion, but it's common that even laptops have more than one CPU that could we could use. Even with just 2 CPUs, we'd accomplish the task 2X faster than using just one. It's common to have 60-90 CPUs (or "cores") on HPC machines. If you aren't using them, you're wasting time!

The GNU `parallel` program (https://www.gnu.org/software/parallel/) provides a simple way to use more than one CPU to complete a batch of jobs. It takes as input the commands that need to be run and spins them out to all available CPUs (or as many as you limit it to), watching for jobs that fail, starting up new jobs when older ones finish.

To see it in action, let's compare these two programs in the "examples/gnu_parallel" directory. The first one simply prints the number 1-30 in order:

```
$ cat -n run.sh
     1  #!/usr/bin/env bash
     2
     3  for i in $(seq 1 30); do
     4      echo $i
     5  done
     6
     7  echo "Done."
```

The second one uses `parallel` to print them. While this is a trivial case, imagine something more intense like BLAST jobs.

```
$ cat -n run_parallel.sh
     1  #!/usr/bin/env bash
     2
     3  JOBS=$(mktemp)
     4
     5  for i in $(seq 1 30); do
     6      echo "echo $i" >> "$JOBS"
     7  done
     8
     9  NUM_JOBS=$(wc -l "$JOBS" | awk '{print $1}')
    10
    11  if [[ $NUM_JOBS -gt 0 ]]; then
    12      echo "Running $NUM_JOBS jobs"
```

```
13       parallel -j 8 --halt soon,fail=1 < "$JOBS"
14   fi
15
16   [[ -f "$JOBS" ]] && rm "$JOBS"
17
18   echo "Done."
```

And here is they look like when they are run:

```
$ ./run.sh
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
Done.
$ ./run_parallel.sh
Running 30 jobs
7
9
8
10
```

```
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
6
5
30
4
3
2
1
Done.
```

The `parallel` version looks out of order because the jobs are run as quickly as possible in whatever order that happens.

CD-HIT

Let's take the `cd-hit` cluster exercise and extend it to where we take the proteins FASTA, run cd-hit, and find the unclustered proteins all in one go. First things first, we need to ensure `cd-hit` is on our system. It's highly unlikely that it is, so let's figure out how to install it.

If you search on the Internet for `cd-hit`, you might end up at http://weizhongli-lab.org/cd-hit/ from which you go to the download page (http://weizhongli-lab.org/cd-hit/download.php) which directs you to the GitHub releases for the `cd-hit` repository (https://github.com/weizhongli/cdhit/releases). From there, we can download the source code tarball (`.tar.gz` file). For instance, I right-click on the link to copy the line address, then go to my HPC into my "downloads" directory and then use `wget` to retrieve the tarball. Next use `tar xvf` to "extract"

in a "verbose" fashion the "file" (followed by the tarball). Finally you should have a directory like `cd-hit-v4.8.1-2019-0228` into which you should `cd`.

If you look at the README, you'll see the way to compile this is to just type `make`. On my Mac laptop, I needed to compile without multi-threading support, so I used `make openmp=no`. That will run for a few seconds and look something like this:

```
$ make openmp=no
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-common.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-utility.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit.o cdhit-common.o cdhit-utility.o -lz -o cd-hit
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-est.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-est.o cdhit-common.o cdhit-utility.o -lz -o cd-hit-e
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-2d.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-2d.o cdhit-common.o cdhit-utility.o -lz -o cd-hit-2d
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-est-2d.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-est-2d.o cdhit-common.o cdhit-utility.o -lz -o cd-hi
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-div.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-div.o cdhit-common.o cdhit-utility.o -lz -o cd-hit-d
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-454.c++ -c
g++  -DNO_OPENMP -DWITH_ZLIB -O2  cdhit-454.o cdhit-common.o cdhit-utility.o -lz -o cd-hit-4
```

Often Makefiles will include an `install` target that will copy the new programs into a directory like `/usr/local/bin`. This one does not, so you'll have to manually copy the programs (e.g., `cd-hit`, `cd-hit-2d`, etc.) to whatever location you like. On an HPC (like Ocelote), you will not have permissions to copy to `/usr/local/bin`, so I'd recommend you create a directory like `$HOME/.local/bin` which you add to your `$PATH` and copy the binaries to that location.

Ensure you have a `cd-hit` binary you can use:

```
$ which cd-hit
/Users/kyclark/.local/bin/cd-hit
$ cd-hit -h | head
        ====== CD-HIT version 4.8.1 (built on Apr  9 2019) ======

Usage: cd-hit [Options]

Options

   -i    input filename in fasta format, required, can be in .gz format
   -o    output filename, required
   -c    sequence identity threshold, default 0.9
    this is the default cd-hit's "global sequence identity" calculated as:
```

Now we can try out our new code:

```
$ cat -n cdhit_unclustered.py
     1  #!/usr/bin/env python3
     2  """
     3  Author : kyclark
     4  Date    : 2019-02-20
     5  Purpose: Run cd-hit, find unclustered proteins
     6  """
     7
     8  import argparse
     9  import datetime
    10  import logging
    11  import os
    12  import re
    13  import signal
    14  import sys
    15  from subprocess import getstatusoutput
    16  from shutil import which
    17  from Bio import SeqIO
    18
    19
    20  # --------------------------------------------------
    21  def get_args():
    22      """get command-line arguments"""
    23      parser = argparse.ArgumentParser(
    24          description='Run cd-hit, find unclustered proteins',
    25          formatter_class=argparse.ArgumentDefaultsHelpFormatter)
    26
    27      parser.add_argument(
    28          '-p',
    29          '--proteins',
    30          help='Proteins FASTA',
    31          metavar='str',
    32          type=str,
    33          required=True)
    34
    35      parser.add_argument(
    36          '-c',
    37          '--seq_id_threshold',
    38          help='cd-hit Sequence identity threshold',
    39          metavar='float',
    40          type=float,
    41          default=0.9)
    42
    43      parser.add_argument(
    44          '-o',
    45          '--outfile',
```

```
46          help='Output file',
47          metavar='str',
48          type=str,
49          default='unclustered.fa')
50
51      parser.add_argument(
52          '-l',
53          '--logfile',
54          help='Log file',
55          metavar='str',
56          type=str,
57          default='.log')
58
59      parser.add_argument('-d', '--debug', help='Debug', action='store_true')
60
61      return parser.parse_args()
62
63
64  # --------------------------------------------------
65  def die(msg='Something bad happened'):
66      """log a critical message() and exit with error"""
67      logging.critical(msg)
68      sys.exit(1)
69
70
71  # --------------------------------------------------
72  def run_cdhit(proteins_file, seq_id_threshold):
73      """Run cd-hit"""
74      cdhit = which('cd-hit')
75
76      if not cdhit:
77          die('Cannot find "cd-hit"')
78
79      out_file = os.path.basename(proteins_file) + '.cdhit'
80      out_path = os.path.join(os.path.dirname(proteins_file), out_file)
81
82      logging.debug('Found cd-hit "{}"'.format(cdhit))
83      cmd = '{} -c {} -i {} -o {} -d 0'.format(cdhit, seq_id_threshold,
84                                               proteins_file, out_file)
85      logging.debug('Running "{}"'.format(cmd))
86      rv, out = getstatusoutput(cmd)
87
88      if rv != 0:
89          die('Non-zero ({}) return from "{}"\n{}\n'.format(rv, cmd, out))
90
91      if not os.path.isfile(out_file):
```

```
 92          die('Failed to create "{}"'.format(out_file))
 93
 94      logging.debug('Finished cd-hit, found cluster file "{}"'.format(out_file))
 95
 96      return out_file
 97
 98
 99  # --------------------------------------------------
100  def get_unclustered(cluster_file, proteins_file, out_file):
101      """Find the unclustered proteins in the cd-hit output"""
102
103      if not os.path.isfile(cluster_file):
104          die('cdhit "{}" is not a file'.format(cluster_file))
105
106      logging.debug('Parsing "{}"'.format(cluster_file))
107
108      clustered = set([rec.id for rec in SeqIO.parse(cluster_file, 'fasta')])
109
110      # Alternate (longer) way:
111      # clustered = set()
112      # for rec in SeqIO.parse(cluster_file, 'fasta'):
113      #     clustered.add(rec.id)
114
115      logging.debug('Will write to "{}"'.format(out_file))
116      out_fh = open(out_file, 'wt')
117      num_total = 0
118      num_unclustered = 0
119
120      for rec in SeqIO.parse(proteins_file, 'fasta'):
121          num_total += 1
122          prot_id = re.sub(r'\|.*', '', rec.id)
123          if not prot_id in clustered:
124              num_unclustered += 1
125              SeqIO.write(rec, out_fh, 'fasta')
126
127      logging.debug(
128          'Finished writing unclustered proteins'.format(num_unclustered))
129
130      return (num_unclustered, num_total)
131
132
133  # --------------------------------------------------
134  def main():
135      """Make a jazz noise here"""
136      args = get_args()
137      proteins_file = args.proteins
```

```
138      out_file = args.outfile
139      log_file = args.logfile
140
141      if not os.path.isfile(proteins_file):
142          die('--proteins "{}" is not a file'.format(arg_name, proteins_file))
143
144      logging.basicConfig(
145          filename=log_file,
146          filemode='a',
147          level=logging.DEBUG if args.debug else logging.CRITICAL)
148
149      def sigint(sig, frame):
150          logging.critical('INT: Exiting early!')
151          sys.exit(0)
152
153      signal.signal(signal.SIGINT, sigint)
154
155      banner = '#' * 50
156      logging.debug(banner)
157      logging.debug('BEGAN {}'.format(str(datetime.datetime.today())))
158
159      cluster_file = run_cdhit(proteins_file, args.seq_id_threshold)
160      num_unclustered, num_total = get_unclustered(cluster_file, proteins_file,
161                                                   out_file)
162
163      msg = 'Wrote {:,d} of {:,d} unclustered proteins to "{}"'.format(
164          num_unclustered, num_total, out_file)
165
166      print(msg)
167      logging.debug(msg)
168      logging.debug('FINISHED {}'.format(str(datetime.datetime.today())))
169      logging.debug(banner)
170
171
172  # --------------------------------------------------
173  if __name__ == '__main__':
174      main()
```

Chapter 15

Containers

Docker (https://www.docker.com/) is one of the most popular technologies for putting all your code and dependecies a "container" that can then be executed on any system that can execute Docker containers. For example, if you needed a program that only works with Python 2 and you don't want to actually install Python2 on your system (things can get really confused with `pip` and `conda` and such), it might be better to build the app into a container.

One down side to Docker is that it requires "root" permissions to run. Most HPC systems will not allow non-privedged users to run as root, so the wonderul folks at Lawrence Berkeley National Labs created Singularity to run containers as a normal user. Singularity containers can be built directly from Docker containers. Because of Singularity's flexibility and abilty to run on HPC platforms, I will focus only on this container.

NB: You must have root access on a Linux box (or an appropriate Linux VM) in order to build Docker and Singularity containers.

Simple Example

In the `examples/simple` directory is a Makefile that will show you the Singularity commands to build an image and a sample `image.def` definition file to instruct Singularity to build a container from a base Ubuntu image, add `git`, use `git` to clone this repository, add the necessary directory to the `$PATH`, and execute the `hello.sh` script automatically when the container is run.

Let's look at the Makefile first:

```
$ cat -n Makefile
     1  .PHONY = img img_old shell run clean
     2
     3  SINGULARITY = /usr/local/bin/singularity
     4  IMG = hello-0.0.1.img
     5  DEF = image.def
     6
     7  img: clean
     8      sudo $(SINGULARITY) build $(IMG) $(DEF)
     9
    10  img_old: clean
    11      sudo $(SINGULARITY) create --size 512 $(IMG)
    12      sudo $(SINGULARITY) bootstrap $(IMG) $(DEF)
    13
    14  shell:
```

```
    15        sudo $(SINGULARITY) shell --writable -B $(shell pwd):/tmp $(IMG)
    16
    17  run:
    18        sudo $(SINGULARITY) exec $(IMG) hello.sh
    19
    20  clean:
    21        rm -f $(IMG)
```

You just need to type `make` or `make img` to execute the `img` target. Notice this
target has a dependency on the `clean` target which will get rid of any existing
image.

Here is the `image.def` that tells Singularity how to build the image:

```
$ cat -n image.def
     1  BootStrap: docker
     2  From: ubuntu:latest
     3
     4  %environment
     5      PATH=/app/biosys-analytics/lectures/15-containers/examples/simple:$PATH
     6
     7  %runscript
     8      exec hello.sh
     9
    10  %post
    11      apt-get update
    12      apt-get install -y locales git
    13      locale-gen en_US.UTF-8
    14
    15      mkdir -p /app
    16      cd /app
    17
    18      git clone https://github.com/hurwitzlab/biosys-analytics.git
```

Notice how I set the `%environment` to include the `$PATH` to the directory con-
taining the `hello.sh` program. After you have an image, you can execute the
image itself like it were a program and Singularity with automatically use the
defined `runscript`:

```
$ make img
<...output ellided...>
Finalizing Singularity container
Calculating final size for metadata...
Skipping checks
Building Singularity image...
Singularity container built: hello-0.0.1.img
Cleaning up...
$ ./hello-0.0.1.img
```

```
Hello from Singularity!
$ make run
sudo /usr/local/bin/singularity exec hello-0.0.1.img hello.sh
Hello from Singularity!
```

I've included in the `img_old` targe the older syntax to `create` and `bootstrap`
a new container in case you are restricted to that version. One problem with
the older syntax is that the image is statically sized at build time. You may
need to allocate more disk space than you actually need for the purpose of
building, but after the build you might like to shrink the image's size. This is
not possible. The newer `build` command will create a base image and expand
and contract the size as needed leaving you with an image that is compressed
into the smallest possible size. The problem is that you cannot `shell` into the
image to test and debug; therefore, it is necessary to build with the older syntax.
Notice that I cannot `make shell` until I `make img_old`. With a shell, I can
`cd` into my directory and execute the script in the environment that it will be
deployed! (This is truly a big deal. You can finally deliver a complete executable
environment for your users!!!)

```
$ make shell
sudo /usr/local/bin/singularity shell --writable -B /home/u20/kyclark/work/biosys-analytics/
ERROR  : Unable to open squashfs image in read-write mode: Read-only file system
ABORT  : Retval = 255
make: *** [shell] Error 255
$ make img_old
$ make shell
sudo /usr/local/bin/singularity shell --writable -B /home/u20/kyclark/work/biosys-analytics/
WARNING: Non existent bind point (file) in container: '/etc/localtime'
Singularity: Invoking an interactive shell within container...

Singularity hello-0.0.1.img:~> cd /app/biosys-analytics/lectures/15-containers/examples/simp
Singularity hello-0.0.1.img:/app/biosys-analytics/lectures/15-containers/examples/simple> ./
Hello from Singularity!
```

Chapter 16

Common Patterns in Python

> "To me programming is more than an important practical art. It is also a gigantic undertaking in the foundations of knowledge." - Grace Hopper

Get positional command-line arguments

You can get the command-line arguments using `sys.argv` (argument vector), but it's annoying that the name of the Python program itself is in the first position (`sys.argv[0]`). To skip over this, take a slice of the argument vector starting at the second position (index `1`) which will succeed even if there are no arguments – you'll get an empty list, which is safe.

```
$ cat -n args.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
     5
     6  args = sys.argv[1:]
     7  num = len(args)
     8
     9  print('There are {} arg{}'.format(num, '' if num == 1 else 's'))
$ ./args.py
There are 0 args
$ ./args.py foo
There are 1 arg
$ ./args.py foo bar
There are 2 args
```

Put positional arguments into named variables

If you use `sys.argv[1]` and `sys.argv[2]` throughout your program, it degrades readability. It's better to copy the values into variables that have meaningful names like "file" or "num_lines".

```
$ cat -n name_args.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
```

257

```
     5
     6  args = sys.argv[1:]
     7
     8  if len(args) != 2:
     9      print('Usage: {} FILE NUM'.format(os.path.basename(sys.argv[0])))
    10      sys.exit(1)
    11
    12  file, num = args
    13
    14  file = args[0]
    15  num = args[1]
    16
    17  print('FILE is "{}", NUM is "{}"'.format(file, num))
$ ./name_args.py
Usage: name_args.py FILE NUM
$ ./name_args.py nobody.txt 10
FILE is "nobody.txt", NUM is "10"
```

Set defaults for optional arguments

```
$ cat -n default_arg.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
     5
     6  args = sys.argv[1:]
     7  num_args = len(args)
     8
     9  if not 1 <= num_args <= 2:
    10      print('Usage: {} FILE [NUM]'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  file = args[0]
    14  num = args[1] if num_args == 2 else 10
    15
    16  print('FILE is "{}", NUM is "{}"'.format(file, num))
$ ./default_arg.py
Usage: default_arg.py FILE [NUM]
$ ./default_arg.py nobody.txt
FILE is "nobody.txt", NUM is "10"
$ ./default_arg.py nobody.txt 5
FILE is "nobody.txt", NUM is "5"
```

Test argument is file and read

This program takes an argument, tests that it is a file, and then reads it. It's basically `cat`.

```
$ cat -n read_file.py
     1  #!/usr/bin/env python3
     2  """Read a file argument"""
     3
     4  import os
     5  import sys
     6
     7  args = sys.argv[1:]
     8
     9  if len(args) != 1:
    10      print('Usage: {} ARG'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  filename = args[0]
    14
    15  if not os.path.isfile(filename):
    16      print('"{}" is not a file'.format(filename), file=sys.stderr)
    17      sys.exit(1)
    18
    19  for line in open(filename):
    20      print(line, end='')
$ ./read_file.py foo
"foo" is not a file
$ ./read_file.py nobody.txt
I'm Nobody! Who are you?
Are you - Nobody - too?
Then there's a pair of us!
Don't tell! they'd advertise - you know!

How dreary - to be - Somebody!
How public - like a Frog -
To tell one's name - the livelong June -
To an admiring Bog!

Emily Dickinson
```

Write data to a file

To write a file, you need to **open** some filename with a second argument of the "mode" where

- **r**: read (default)
- **w**: write
- **t**: text mode (default)
- **b**: binary

You can combine the flags so that **wt** means "write a text file" which is what is done here.

If you **open** a file for writing and the file already exists, it will be overwritten, so it may behoove you to check if the file exists first!

```
$ cat -n write_file.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
     5
     6
     7  args = sys.argv[1:]
     8
     9  if len(args) < 1:
    10      print('Usage: {} ARG1 [ARG2...]'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  outfile = 'out.txt'
    14  out_fh = open(outfile, 'wt')
    15
    16  for arg in args:
    17      out_fh.write(arg + '\n')
    18
    19  out_fh.close()
    20  print('Done, see "{}"'.format(outfile))
$ ./write_file.py foo bar baz
Done, see "out.txt"
$ cat out.txt
foo
bar
baz
```

Test if an argument is a directory and list the contents

```
$ cat -n list_dir.py
     1  #!/usr/bin/env python3
     2  """Show contents of directory argument"""
     3
     4  import os
```

260

```
     5  import sys
     6
     7  args = sys.argv[1:]
     8
     9  if len(args) != 1:
    10      print('Usage: {} DIR'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  dirname = args[0]
    14
    15  if not os.path.isdir(dirname):
    16      print('"{}" is not a directory'.format(dirname), file=sys.stderr)
    17      sys.exit(1)
    18
    19  for entry in os.listdir(dirname):
    20      print(entry)
$ ./list_dir.py
Usage: list_dir.py DIR
$ ./list_dir.py .
list_dir.py
kmers.py
skip_loop.py
unpack_dict2.py
nobody.txt
name_args.py
create_dir.py
sort_dict_by_values.py
foo
args.py
sort_dict_by_keys.py
read_file.py
sort_dict_by_keys2.py
unpack_dict.py
codons.py
default_arg.py
```

Skip an iteration of a loop

Sometimes in a loop (`for` or `while`) you want to skip immediately to the top of the loop. You can use `continue` to do this. In this example, we skip the even-numbered lines by using the modulus `%` operator to find those line numbers which have a remainder of 0 after dividing by 2. We can use the `enumerate` function to provide both the array index and value of any list.

```
$ cat -n skip_loop.py
```

```
 1  #!/usr/bin/env python3
 2
 3  import os
 4  import sys
 5
 6
 7  args = sys.argv[1:]
 8
 9  if len(args) != 1:
10      print('Usage: {} FILE'.format(os.path.basename(sys.argv[0])))
11      sys.exit(1)
12
13  file = args[0]
14
15  if not os.path.isfile(file):
16      print('"{}" is not a file'.format(file), file=sys.stderr)
17      sys.exit(1)
18
19  for i, line in enumerate(open(file)):
20      if (i + 1) % 2 == 0:
21          continue
22
23      print(i + 1, line, end='')
```
```
$ ./skip_loop.py
Usage: skip_loop.py FILE
$ ./skip_loop.py nobody.txt
1 I'm Nobody! Who are you?
3 Then there's a pair of us!
5
7 How public - like a Frog -
9 To an admiring Bog!
11 Emily Dickinson
```

Create a directory if it does not exist

This program takes a directory name and looks to see if it already exists or needs
to be created.

```
$ cat -n create_dir.py
 1  #!/usr/bin/env python3
 2  """Test for a directory and create if needed"""
 3
 4  import os
 5  import sys
 6
```

```
     7  args = sys.argv[1:]
     8
     9  if len(args) != 1:
    10      print('Usage: {} DIR'.format(os.path.basename(sys.argv[0])))
    11      sys.exit(1)
    12
    13  dirname = args[0]
    14
    15  if os.path.isdir(dirname):
    16      print('"{}" exists'.format(dirname))
    17  else:
    18      print('Creating "{}"'.format(dirname))
    19      os.makedirs(dirname)
$ ./create_dir.py
Usage: create_dir.py DIR
$ ./create_dir.py foo
Creating "foo"
$ ./create_dir.py foo
"foo" exists
```

Unpack a dictionary's key/values pairs

The `.items()` method on a dictionary will return a list of tuples:

```
$ cat -n unpack_dict.py
     1  #!/usr/bin/env python3
     2  """Unpack dict"""
     3
     4  import os
     5  import sys
     6
     7  albums = {
     8      "2112": 1976,
     9      "A Farewell To Kings": 1977,
    10      "All the World's a Stage": 1976,
    11      "Caress of Steel": 1975,
    12      "Exit, Stage Left": 1981,
    13      "Fly By Night": 1975,
    14      "Grace Under Pressure": 1984,
    15      "Hemispheres": 1978,
    16      "Hold Your Fire": 1987,
    17      "Moving Pictures": 1981,
    18      "Permanent Waves": 1980,
    19      "Power Windows": 1985,
```

```
    20      "Signals": 1982,
    21  }
    22
    23  for tup in albums.items():
    24      album = tup[0]
    25      year = tup[1]
    26      print('{:4} {}'.format(year, album))
$ ./unpack_dict.py
1976 2112
1977 A Farewell To Kings
1976 All the World's a Stage
1975 Caress of Steel
1981 Exit, Stage Left
1975 Fly By Night
1984 Grace Under Pressure
1978 Hemispheres
1987 Hold Your Fire
1981 Moving Pictures
1980 Permanent Waves
1985 Power Windows
1982 Signals
```

But the `for` loop could unpack the tuple directly. Compare line 23 in the above and below programs.

```
$ cat -n unpack_dict2.py
     1  #!/usr/bin/env python3
     2  """Unpack dict"""
     3
     4  import os
     5  import sys
     6
     7  albums = {
     8      "2112": 1976,
     9      "A Farewell To Kings": 1977,
    10      "All the World's a Stage": 1976,
    11      "Caress of Steel": 1975,
    12      "Exit, Stage Left": 1981,
    13      "Fly By Night": 1975,
    14      "Grace Under Pressure": 1984,
    15      "Hemispheres": 1978,
    16      "Hold Your Fire": 1987,
    17      "Moving Pictures": 1981,
    18      "Permanent Waves": 1980,
    19      "Power Windows": 1985,
    20      "Signals": 1982,
    21  }
```

```
    22
    23  for album, year in albums.items():
    24      print('{:4} {}'.format(year, album))
$ ./unpack_dict2.py
1976 2112
1977 A Farewell To Kings
1976 All the World's a Stage
1975 Caress of Steel
1981 Exit, Stage Left
1975 Fly By Night
1984 Grace Under Pressure
1978 Hemispheres
1987 Hold Your Fire
1981 Moving Pictures
1980 Permanent Waves
1985 Power Windows
1982 Signals
```

Sort a dictionary by keys

To sort a dictionary by the keys, you have to understand that the `.sort()` method of an list mutates the list *in-place*. We get the keys of a dictionary with the `.keys()` method which does not support the `.sort()` method:

```
>>> d = dict(foo=1, bar=2)
>>> d.keys()
dict_keys(['foo', 'bar'])
>>> d.keys().sort()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'dict_keys' object has no attribute 'sort'
```

We could copy the keys into a list to sort like so:

```
>>> k = list(d.keys())
>>> k
['foo', 'bar']
>>> k.sort()
>>> k
['bar', 'foo']
```

Or we can use the `sorted()` function that accepts a list and *returns a sorted list*:

```
>>> d.keys()
dict_keys(['foo', 'bar'])
>>> sorted(d.keys())
```

```
['bar', 'foo']
```

Either way, once we have the sorted keys, we can get the associated values:

```
$ cat -n sort_dict_by_keys.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
     5
     6  albums = {
     7      "2112": 1976,
     8      "A Farewell To Kings": 1977,
     9      "All the World's a Stage": 1976,
    10      "Caress of Steel": 1975,
    11      "Exit, Stage Left": 1981,
    12      "Fly By Night": 1975,
    13      "Grace Under Pressure": 1984,
    14      "Hemispheres": 1978,
    15      "Hold Your Fire": 1987,
    16      "Moving Pictures": 1981,
    17      "Permanent Waves": 1980,
    18      "Power Windows": 1985,
    19      "Signals": 1982,
    20  }
    21
    22  for album in sorted(albums.keys()):
    23      print('{:25} {}'.format(album, albums[album]))
$ ./sort_dict_by_keys.py
2112                      1976
A Farewell To Kings       1977
All the World's a Stage   1976
Caress of Steel           1975
Exit, Stage Left          1981
Fly By Night              1975
Grace Under Pressure      1984
Hemispheres               1978
Hold Your Fire            1987
Moving Pictures           1981
Permanent Waves           1980
Power Windows             1985
Signals                   1982
```

Or we could unpack the tuples directly like above:

```
$ cat -n sort_dict_by_keys2.py
     1  #!/usr/bin/env python3
     2
```

```
    3  import os
    4  import sys
    5
    6  albums = {
    7      "2112": 1976,
    8      "A Farewell To Kings": 1977,
    9      "All the World's a Stage": 1976,
   10      "Caress of Steel": 1975,
   11      "Exit, Stage Left": 1981,
   12      "Fly By Night": 1975,
   13      "Grace Under Pressure": 1984,
   14      "Hemispheres": 1978,
   15      "Hold Your Fire": 1987,
   16      "Moving Pictures": 1981,
   17      "Permanent Waves": 1980,
   18      "Power Windows": 1985,
   19      "Signals": 1982,
   20  }
   21
   22  for album, year in sorted(albums.items()):
   23      print('{:25} {}'.format(album, year))
$ ./sort_dict_by_keys2.py
2112                      1976
A Farewell To Kings       1977
All the World's a Stage   1976
Caress of Steel           1975
Exit, Stage Left          1981
Fly By Night              1975
Grace Under Pressure      1984
Hemispheres               1978
Hold Your Fire            1987
Moving Pictures           1981
Permanent Waves           1980
Power Windows             1985
Signals                   1982
```

Sort a dictionary by values

To sort a dictionary by the values rather than the keys, we need to reverse the tuples which is what happens on line 24. Notice that in years when two albums were released, the **sorted** first sorts by the first tuple member (the year) and then the second (album name):

```
$ cat -n sort_dict_by_values.py
    1  #!/usr/bin/env python3
```

```
 2
 3  import os
 4  import sys
 5
 6  albums = {
 7      "2112": 1976,
 8      "A Farewell To Kings": 1977,
 9      "All the World's a Stage": 1976,
10      "Caress of Steel": 1975,
11      "Exit, Stage Left": 1981,
12      "Fly By Night": 1975,
13      "Grace Under Pressure": 1984,
14      "Hemispheres": 1978,
15      "Hold Your Fire": 1987,
16      "Moving Pictures": 1981,
17      "Permanent Waves": 1980,
18      "Power Windows": 1985,
19      "Signals": 1982,
20  }
21
22  # Create a list of (value, key) tuples
23  # sorted in descending order by the values
24  pairs = sorted([(x[1], x[0]) for x in albums.items()])
25
26  for year, album in pairs:
27      print('{} {}'.format(year, album))
$ ./sort_dict_by_values.py
1975 Caress of Steel
1975 Fly By Night
1976 2112
1976 All the World's a Stage
1977 A Farewell To Kings
1978 Hemispheres
1980 Permanent Waves
1981 Exit, Stage Left
1981 Moving Pictures
1982 Signals
1984 Grace Under Pressure
1985 Power Windows
1987 Hold Your Fire
```

Extract codons from DNA

This example assumes a codon length (`k`) of 3 and uses a handy third argument
to `range` that indicates the distance to skip in each iteration. The goal is to start
at position 0, then jump to position 3, then 6, etc., to extract all the codons.
Imagine how you could expand this to get all the codons in all the frames (this
one starts at "1" which is really "0" in the string):

```
$ cat -n codons.py
     1  #!/usr/bin/env python3
     2  """Extract codons from DNA"""
     3
     4  import os
     5  import sys
     6
     7  args = sys.argv[1:]
     8  num_args = len(args)
     9
    10  if not 1 <= num_args <= 2:
    11      print('Usage: {} DNA'.format(os.path.basename(sys.argv[0])))
    12      sys.exit(1)
    13
    14  string = args[0]
    15  k = 3
    16  n = len(string) - k + 1
    17
    18  for i in range(0, n, k):
    19      print(string[i:i+k])
$ ./codons.py
Usage: codons.py DNA
$ ./codons.py AAACCCGGGTTT
AAA
CCC
GGG
TTT
```

Extract k-mers from a string

K-mers are `k`-length contiguous sub-sequences from a string. They are similar
to codons (which are 3-mers), but we tend to move across the string by one
character than than the codon length (3). Notice this script guards against a
2nd argument that should be a number but is not:

```
$ cat -n kmers.py
     1  #!/usr/bin/env python3
```

```
 2  """Extract k-mers from string"""
 3
 4  import os
 5  import sys
 6
 7  args = sys.argv[1:]
 8  num_args = len(args)
 9
10  if not 1 <= num_args <= 2:
11      print('Usage: {} STR [K]'.format(os.path.basename(sys.argv[0])))
12      sys.exit(1)
13
14  string = args[0]
15  k = args[1] if num_args == 2 else '3'
16
17  # Guard against a string like "foo"
18  if not k.isdigit():
19      print('k "{}" is not a digit'.format(k))
20      sys.exit(1)
21
22  # Safe to convert now
23  k = int(k)
24
25  if len(string) < k:
26      print('There are no {}-length substrings in "{}"'.format(k, string))
27  else:
28      n = len(string) - k + 1
29      for i in range(0, n):
30          print(string[i:i+k])
$ ./kmers.py
Usage: kmers.py STR [K]
$ ./kmers.py foobar 10
There are no 10-length substrings in "foobar"
$ ./kmers.py AAACCCGGGTTT 3
AAA
AAC
ACC
CCC
CCG
CGG
GGG
GGT
GTT
TTT
```

Make All Items in a List Uppercase

If you need to check all the strings in a list in a case-insensitive fashion, one way would be to upper- or lower-case all the strings. A very Pythonic way is to use a list comprehension, but we can also borrow an idea from the purely functional programming world where we use a "higher-order function," which is a function that takes one or more other functions as arguments. In this case, the `map` function expects as it's first argument some other function, here `str.upper`. Notice it's not `str.upper()` with parens! That is the syntax for **calling** the `str.upper` function. We want to pass **the function itself**, so we leave off the parens. The function is applied to each item in the list and returns a new list. The original list remains unchanged.

```
>>> a = ['foo', 'bar', 'baz']
>>> [s.upper() for s in a]
['FOO', 'BAR', 'BAZ']
>>> list(map(str.upper, a))
['FOO', 'BAR', 'BAZ']
>>> a
['foo', 'bar', 'baz']
```

Another way to write the `map` is to use a `lambda` expression which is just a very short, anonymouse (unnamed) function:

```
>>> list(map(lambda s: s.upper(), a))
['FOO', 'BAR', 'BAZ']
```

Here is the code in action:

```
$ cat -n upper_list.py
     1  #!/usr/bin/env python3
     2
     3  import os
     4  import sys
     5
     6  args = sys.argv[1:]
     7
     8  if len(args) < 1:
     9      print('Usage: {} ARG [ARG...]'.format(os.path.basename(sys.argv[0])))
    10      sys.exit(1)
    11
    12  print('List comprehension')
    13  print(', '.join([x.upper() for x in args]))
    14
    15  print('Map')
    16  print(', '.join(map(str.upper, args)))
$ ./upper_list.py foo bar baz
List comprehension
```

```
FOO, BAR, BAZ
Map
FOO, BAR, BAZ
```